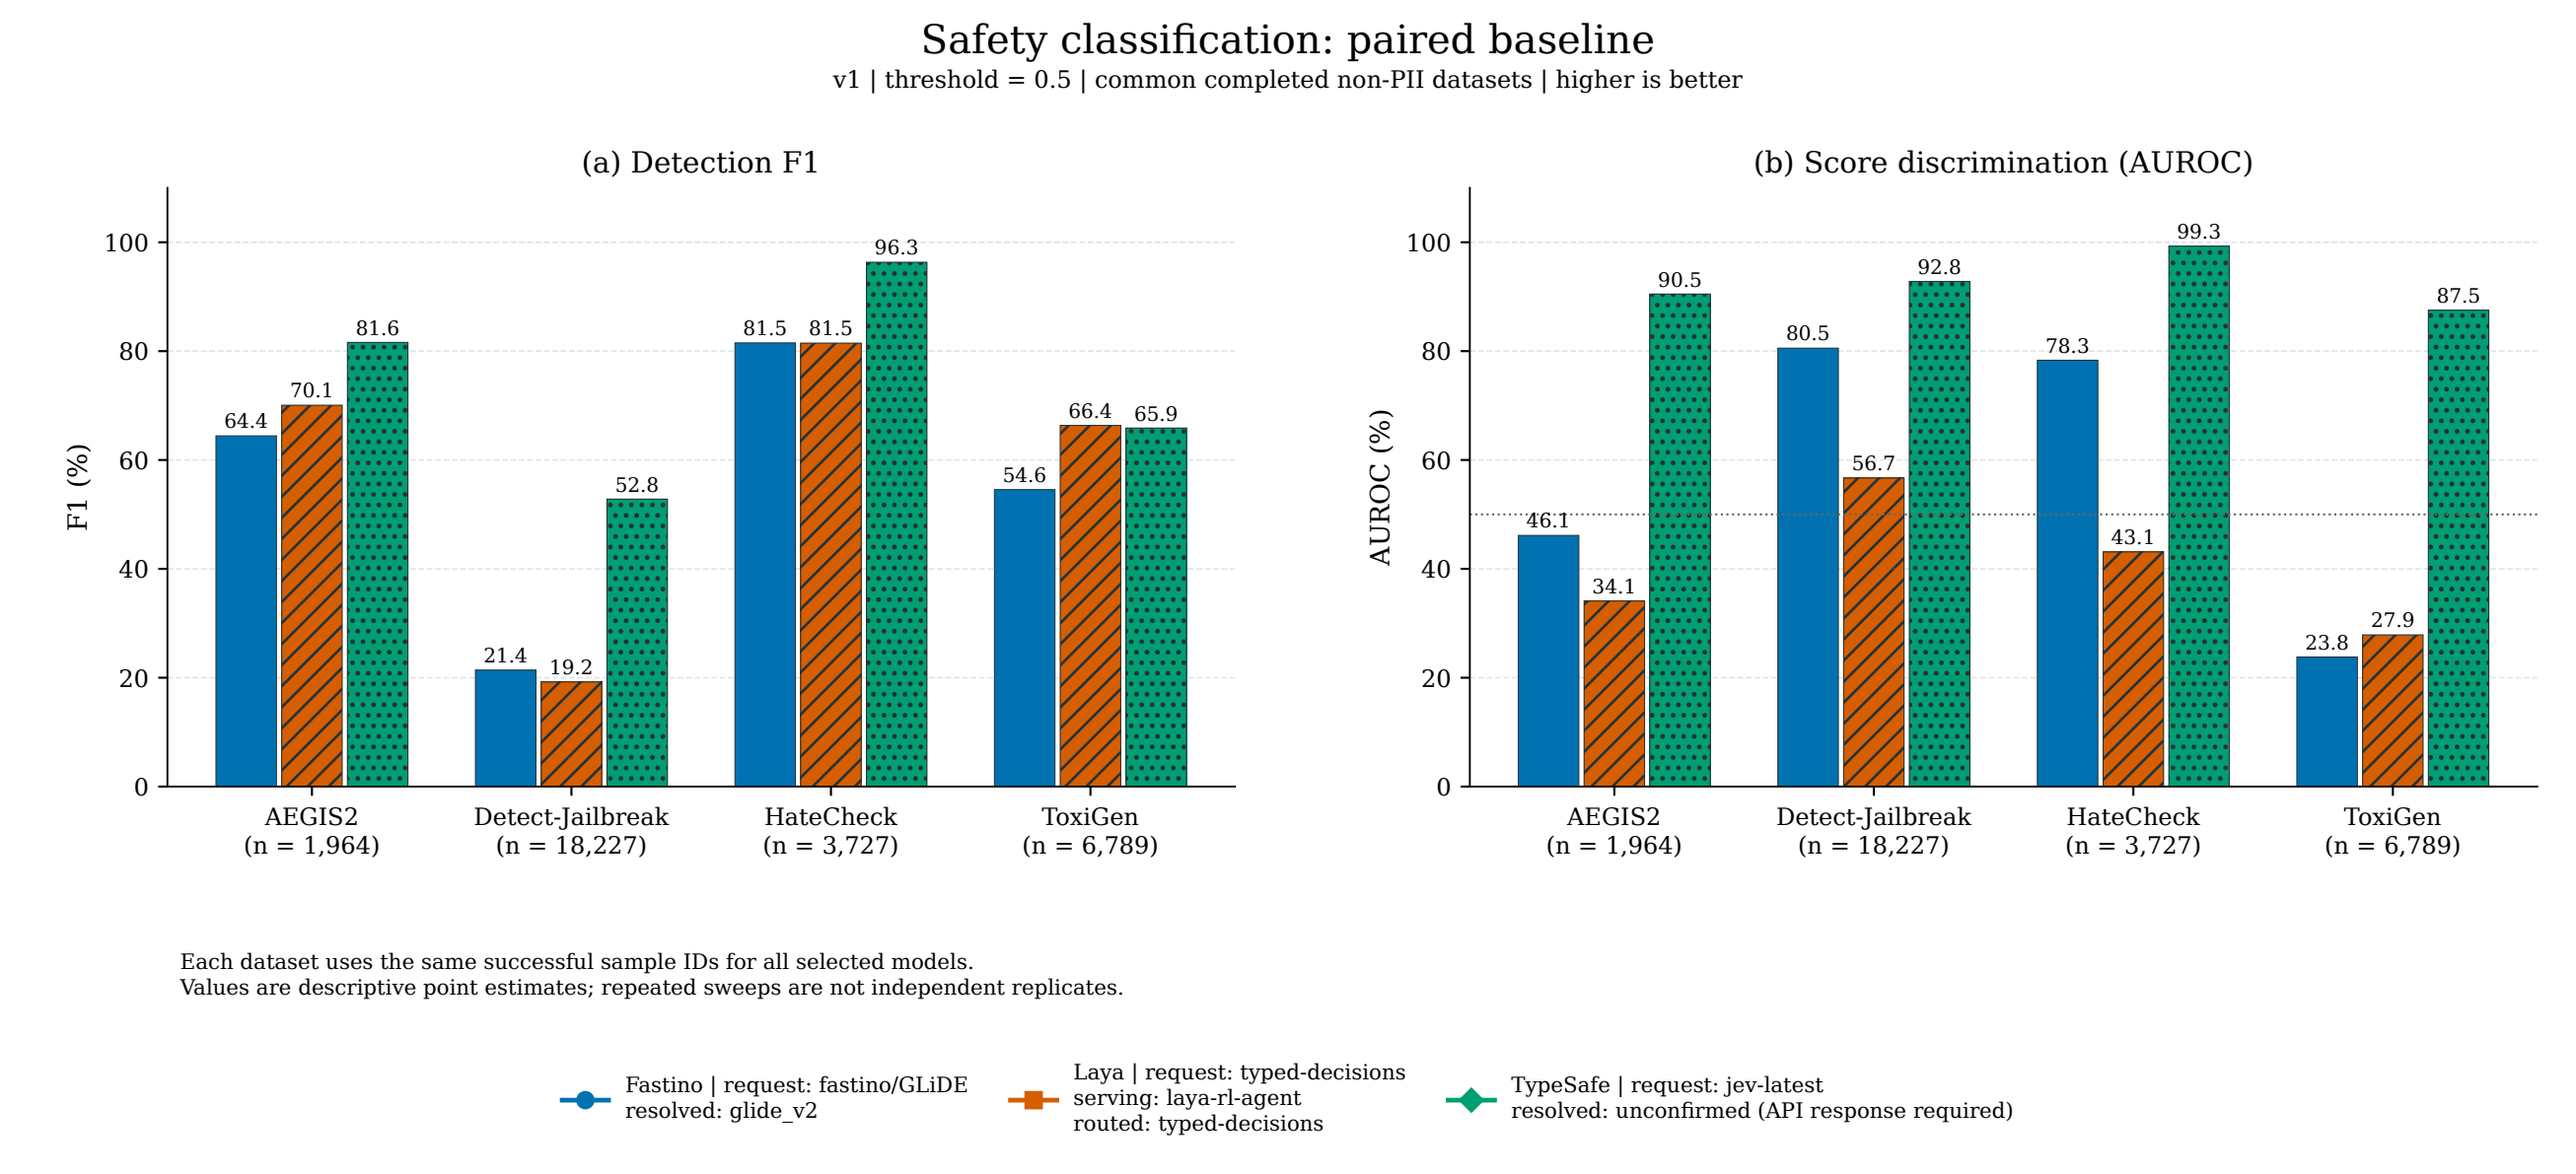
Safety classification: paired baseline
v1 | threshold = 0.5 | common completed non-PII datasets | higher is better
(a) Detection F1
F1 (%)
0
20
40
60
80
100
64.4
70.1
81.6
21.4
19.2
52.8
81.5
81.5
96.3
54.6
66.4
65.9
AEGIS2
(n = 1,964)
Detect-Jailbreak
(n = 18,227)
HateCheck
(n = 3,727)
ToxiGen
(n = 6,789)
(b) Score discrimination (AUROC)
AUROC (%)
0
20
40
60
80
100
46.1
34.1
90.5
80.5
56.7
92.8
78.3
43.1
99.3
23.8
27.9
87.5
AEGIS2
(n = 1,964)
Detect-Jailbreak
(n = 18,227)
HateCheck
(n = 3,727)
ToxiGen
(n = 6,789)
Each dataset uses the same successful sample IDs for all selected models.
Values are descriptive point estimates; repeated sweeps are not independent replicates.
Fastino | request: fastino/GLiDE
resolved: glide_v2
Laya | request: typed-decisions
serving: laya-rl-agent
routed: typed-decisions
TypeSafe | request: jev-latest
resolved: unconfirmed (API response required)