Error Rates and Latency Trade-offs

v1 | threshold = 0.5 | (a) 5 shared datasets | (b) 5 shared fully timed datasets

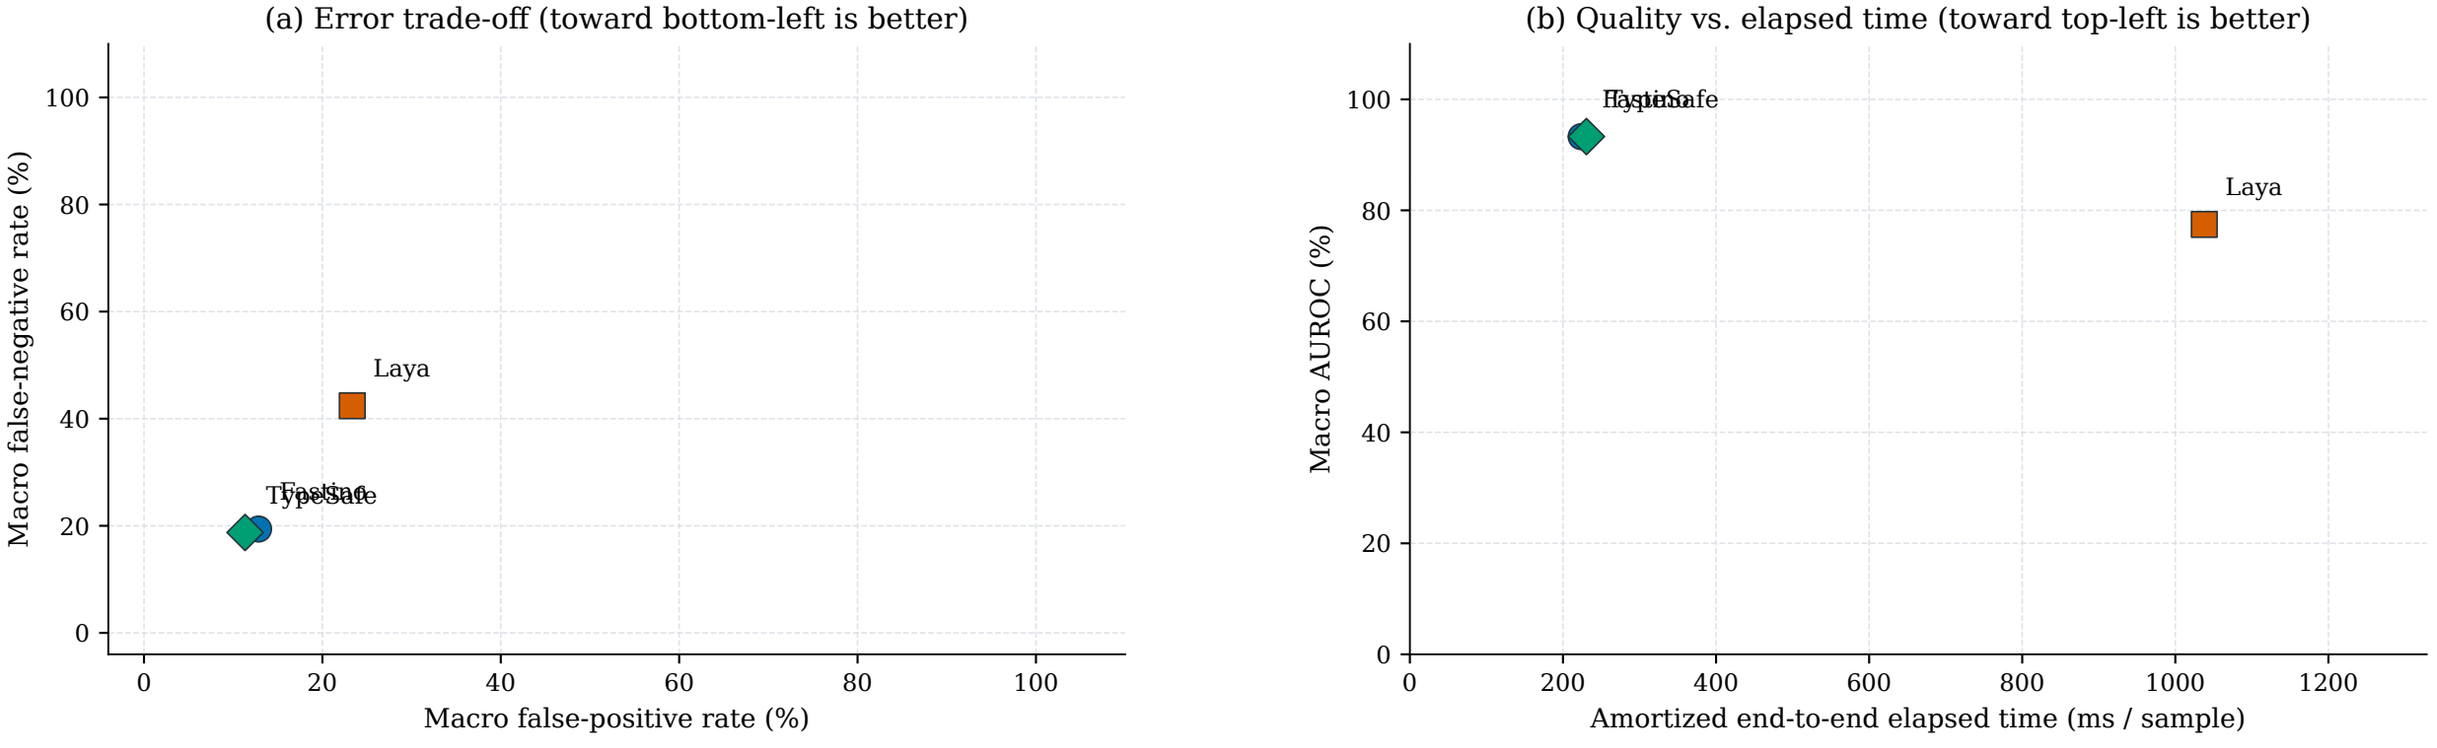

Panel (b) recomputes BOTH macro AUROC and sample-weighted elapsed on the same shared timed subset:
AEGIS2, HateCheck, Jailbreak-Classification, JailbreakHub, ToxiGen. No four-run timing averages substitute for missing v1 / 0.5 timing.
Elapsed time = total logged elapsed / all processed subset rows, including failed requests.
Runs occurred at different times; this is not pure inference latency or a request-level percentile.

Fastino | request: fastino/GLiDE
resolved: glide_v2

Laya | request: typed-decisions
serving: laya-rl-agent
routed: typed-decisions

TypeSafe | request: jev-latest
resolved: unconfirmed (API response required)