Model averages on the shared baseline

v1 | threshold = 0.5 | 4 shared non-PII datasets | equal dataset weights

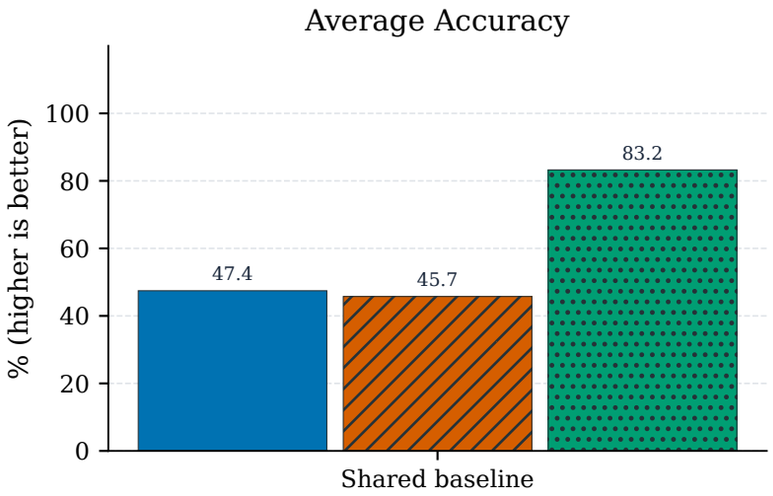

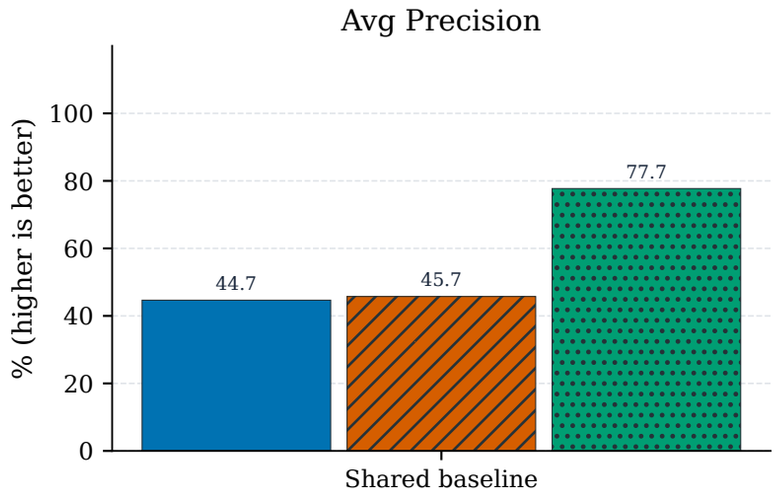

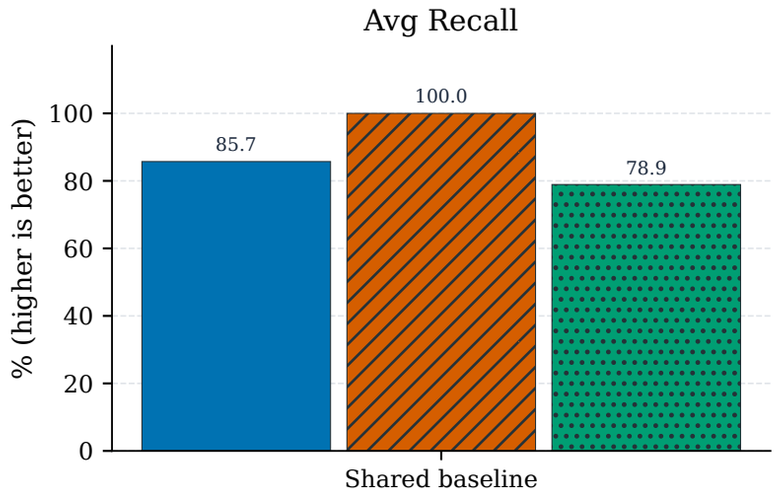

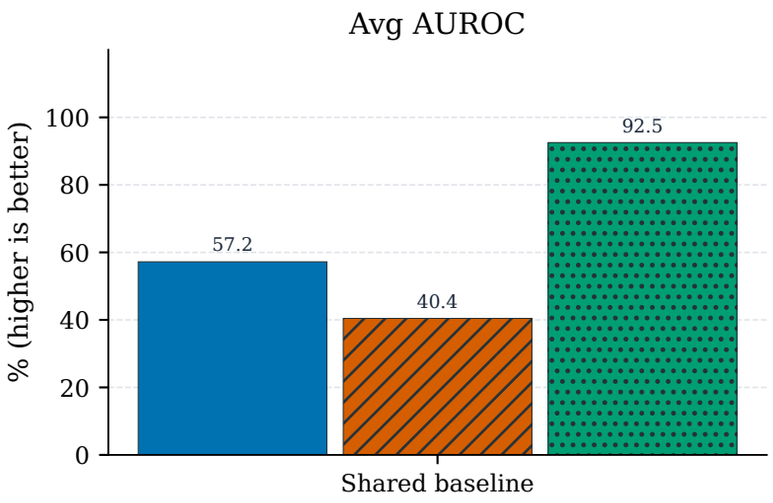

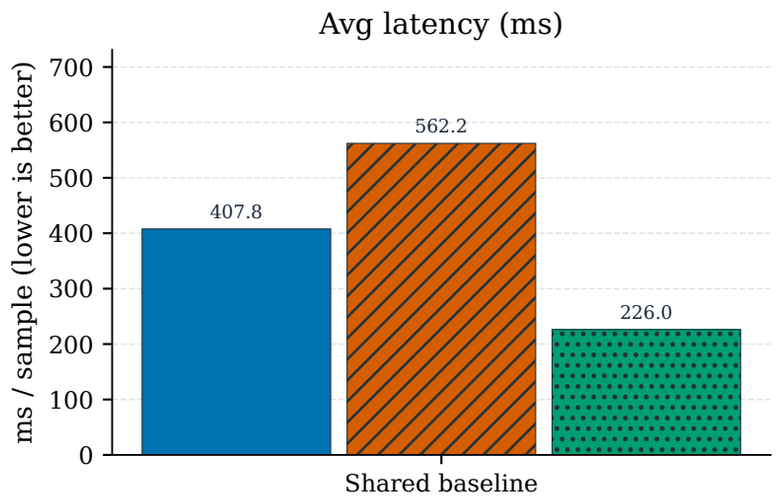

Quality: common successful sample IDs within each dataset, then an unweighted mean across datasets.
Latency: unweighted mean of per-dataset log elapsed / processed rows, including failures (not inference time).
N/A timing means at least one baseline dataset has no aligned complete timing log; missing is never zero.

Fastino | request: fastino/GLiDE
resolved: glide_v2

Laya | request: typed-decisions
serving: laya-rl-agent
routed: typed-decisions

TypeSafe | request: jev-latest
resolved: unconfirmed (API response required)