Sample-Aligned Model Averages (Fastino, Laya, TypeSafe)

v1 | threshold = 0.5 | quality: 5 shared non-PII datasets | latency: 5 shared fully timed datasets | equal dataset weights

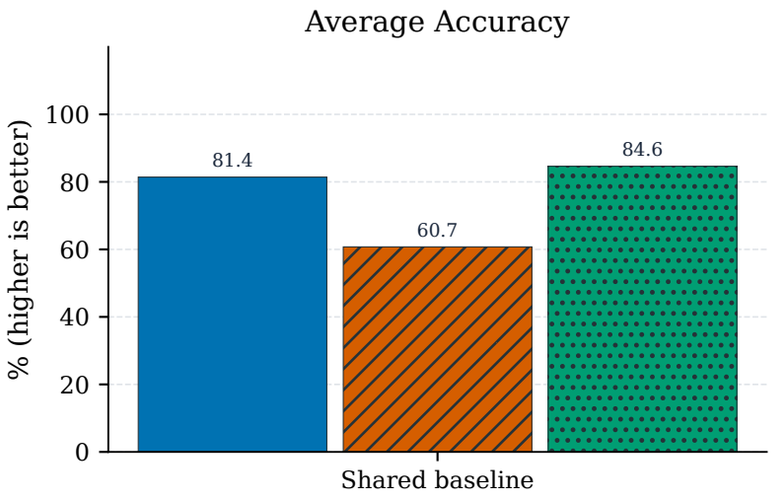

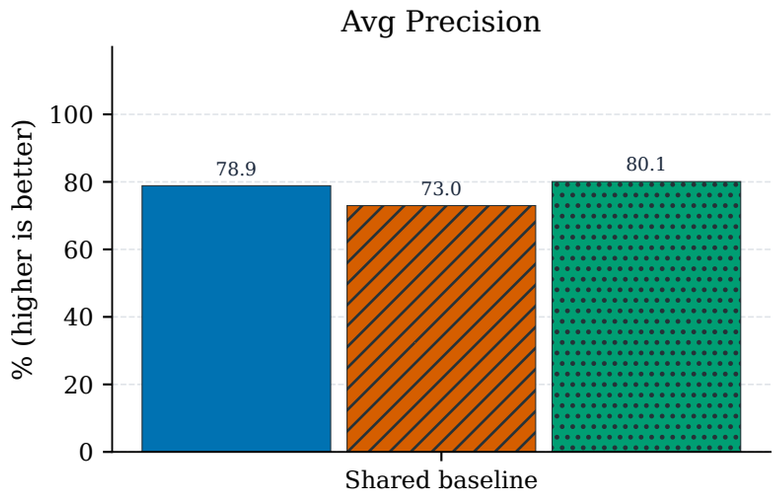

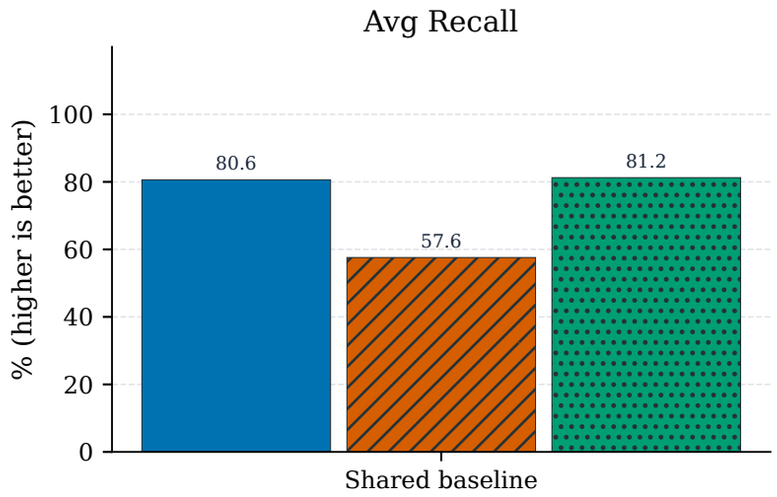

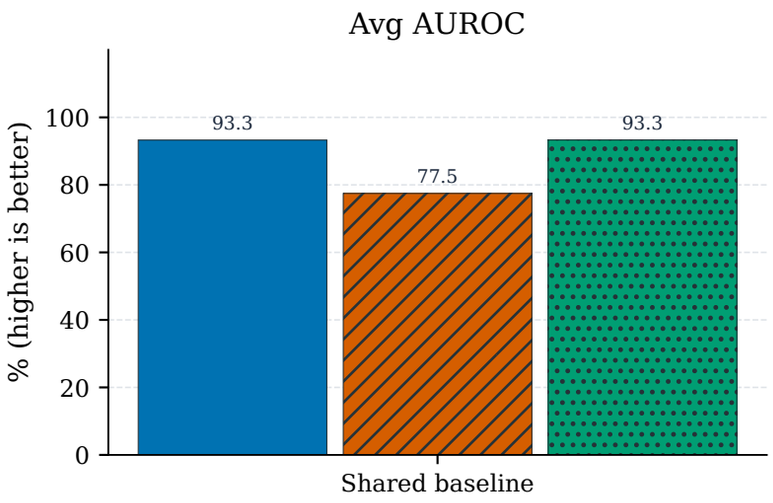

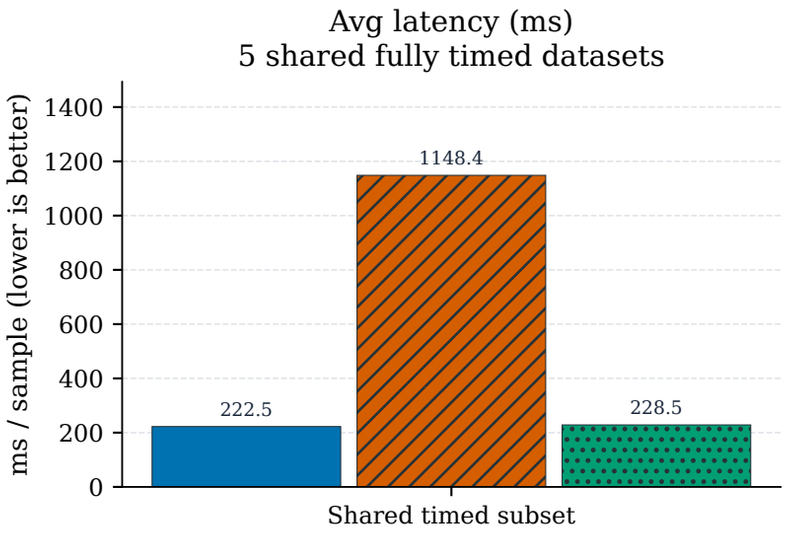

Quality: common successful sample IDs within each dataset, then an unweighted mean across datasets.
Latency: unweighted mean of per-dataset log elapsed / processed rows on the SAME timed subset for all models: AEGIS2, HateCheck, Jailbreak-Classification, JailbreakHub, ToxiGen. Quality and latency dataset scopes can differ.
Elapsed includes failures and network overhead, not pure inference time. No four-run timing substitutes.
N/A timing means no shared fully timed baseline datasets; missing is never zero.

Fastino | request: fastino/GLiDE
resolved: glide_v2

Laya | request: typed-decisions
serving: laya-rl-agent
routed: typed-decisions

TypeSafe | request: jev-latest
resolved: unconfirmed (API response required)