Repeat-call consistency: prompt v2

3 shared datasets | four calls per model | identical successful sample IDs across all models

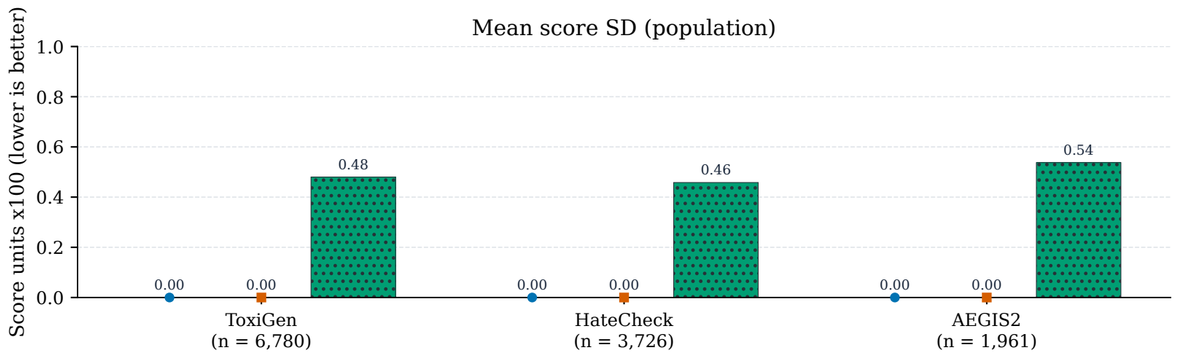

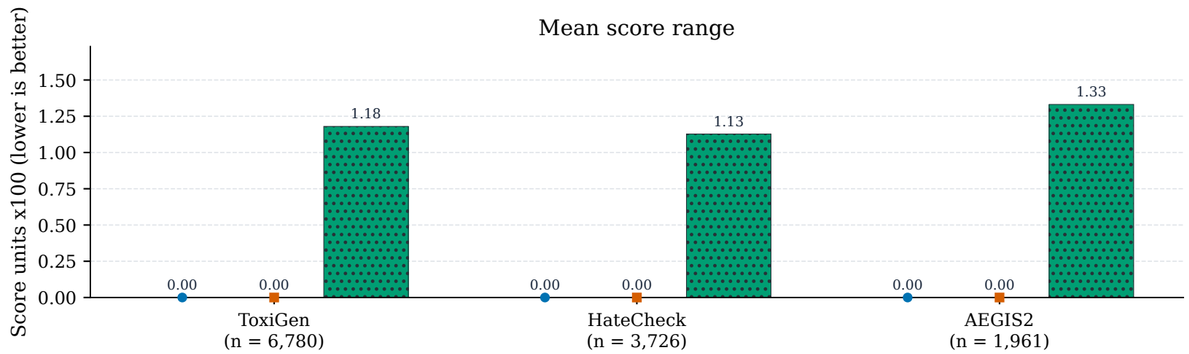

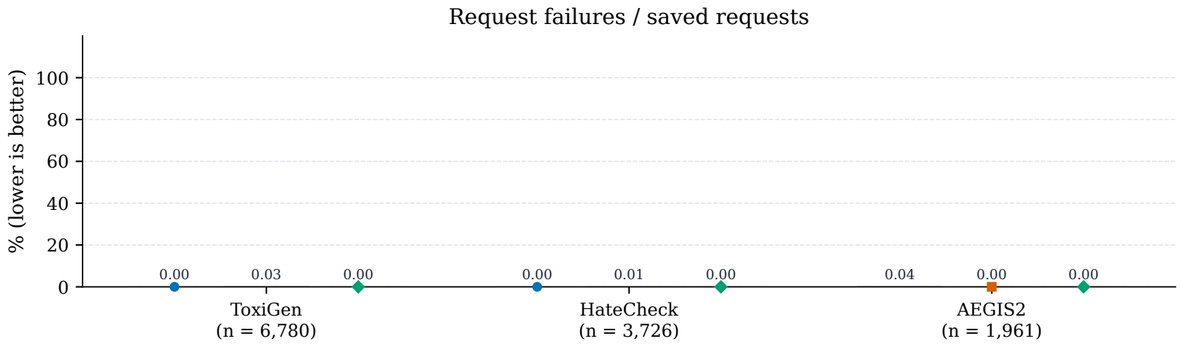

Score SD and range are displayed x100; flip rate uses score >= 0.5 for ALL four calls (not saved predictions).
Only datasets with complete four-call CSVs for every model and nonempty common successful IDs are shown.
Baseline markers show measured zeros, not missing data. Coverage uses all-model common IDs; failures use all saved requests.
Historical backend / routing / time changes may contribute. Consistency is not accuracy; this is exploratory.

Fastino | request: fastino/GLiDE
resolved: glide_v2

Laya | request: typed-decisions
serving: laya-rl-agent
routed: typed-decisions

TypeSafe | request: jev-latest
resolved: unconfirmed (API response required)