Model repeat-call consistency: paired dataset-macro averages

Prompt v1 and v2 are separate | common four-call successful IDs | shared non-PII datasets only

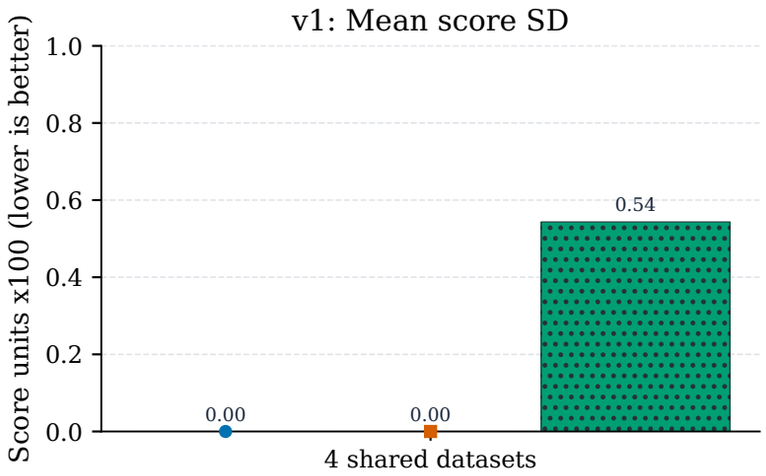

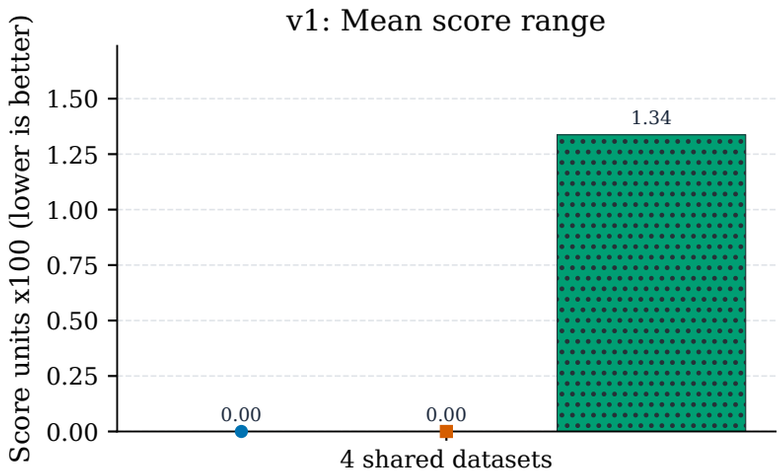

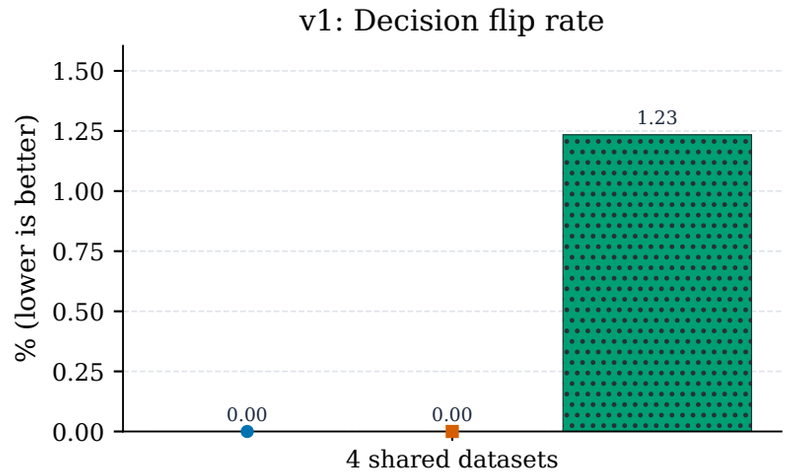

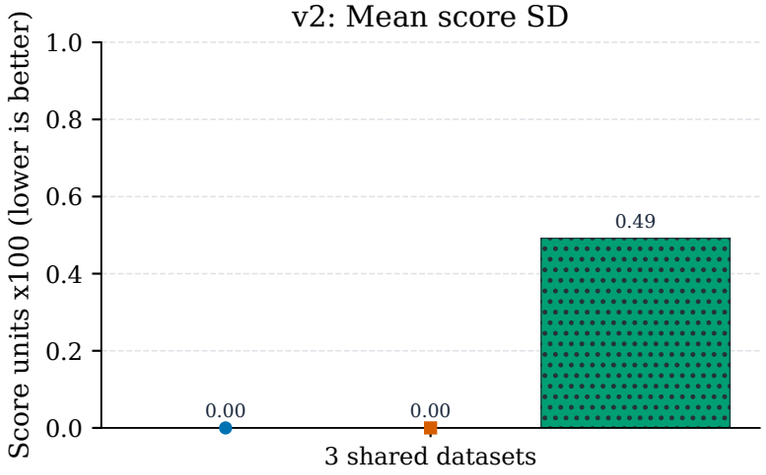

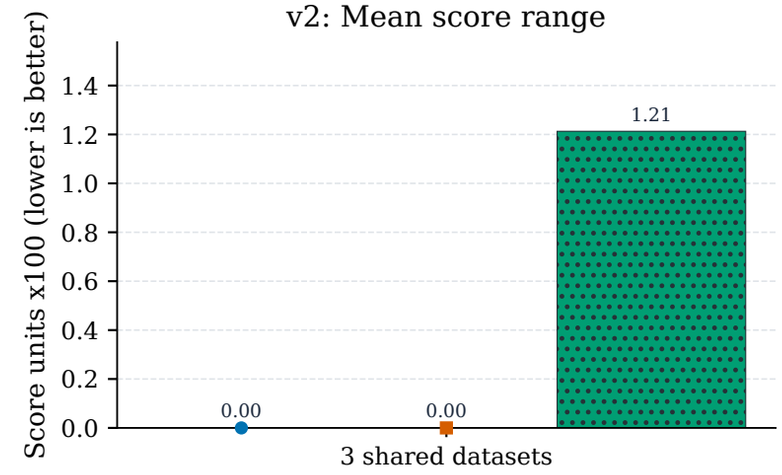

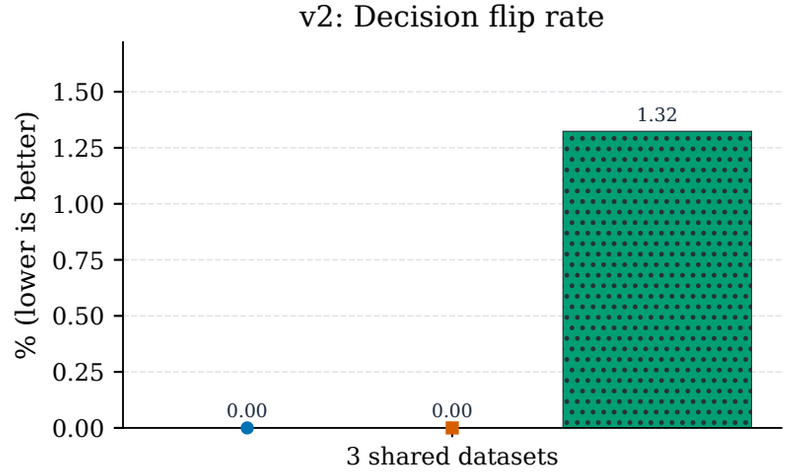

Equal dataset weights. Score SD / range displayed x100; decision flips use the fixed threshold 0.5.
Baseline markers show measured zeros. No eligible shared datasets = N/A; datasets and IDs may differ between prompts.
See the report for shared sample counts, coverage and failure rates; low variation does not imply correctness.

Fastino | request: fastino/GLiDE
resolved: glide_v2

Laya | request: typed-decisions
serving: laya-rl-agent
routed: typed-decisions

TypeSafe | request: jev-latest
resolved: unconfirmed (API response required)