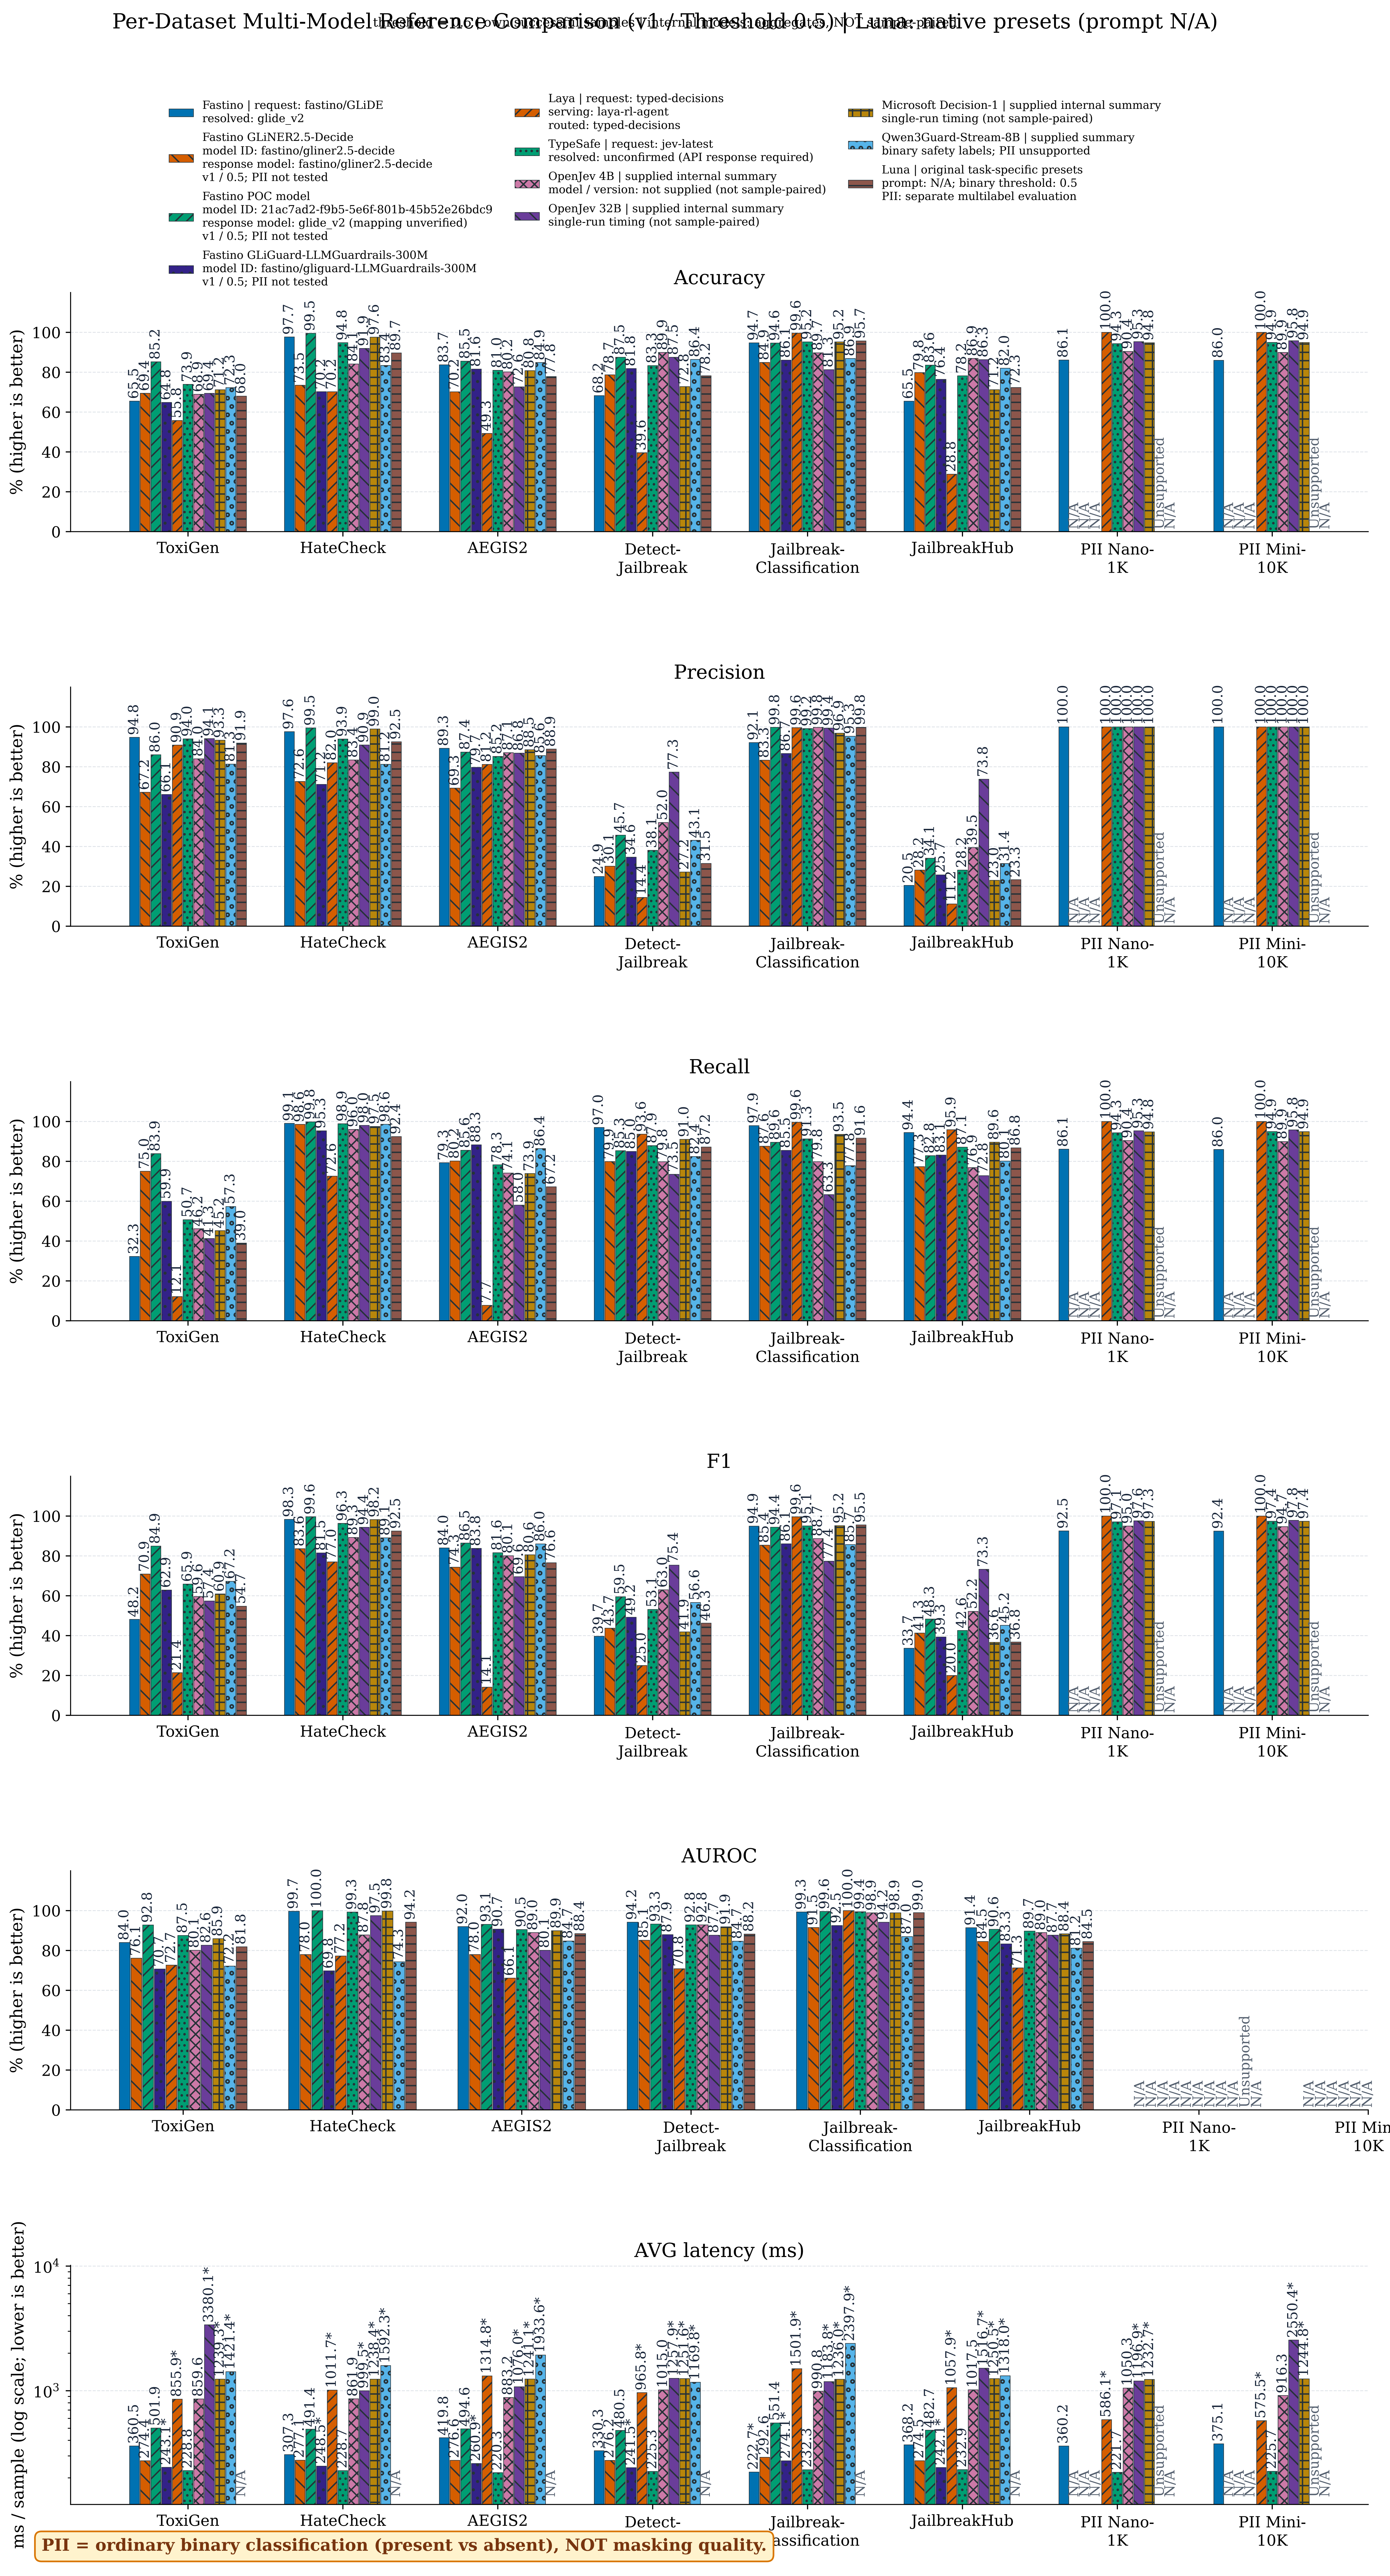

Per-Dataset Multi-Model Reference Comparison (v1 / Threshold 0.5) | Luna: native presets (prompt N/A)

threshold = 0.5 | own successful samples | internal models: aggregates, NOT sample-paired

PII = ordinary binary classification (present vs absent), NOT masking quality.

OpenJev 4B: four-run summary timing; local models prefer four-run log averages. 4B valid/error counts are unknown.
* = single-run v1 / 0.5 timing; local models require complete, matched logs; internal models use supplied summaries.
OpenJev 32B / Microsoft Decision-1: supplied single-run timing, including failures; protocols differ.
32B has 27,460 errors / 56,823 rows; quality excludes errors.
N/A = missing quality, undefined metric, or incomplete timing; missing values are not zero.
Descriptive only, not a fair ranking. PII is ordinary binary classification (present vs absent), not masking quality.
Qwen3Guard-Stream-8B: binary safety labels; PII unsupported (historical proxies are audit-only).
Qwen mean latency: log elapsed/processed (failures included); CSV P95: successful requests. No ranking.
Luna: original task-specific presets, no custom prompts (prompt N/A); binary threshold 0.5.
Other models use dataset-specific v1/v2 prompts. Task-definition differences may affect results.