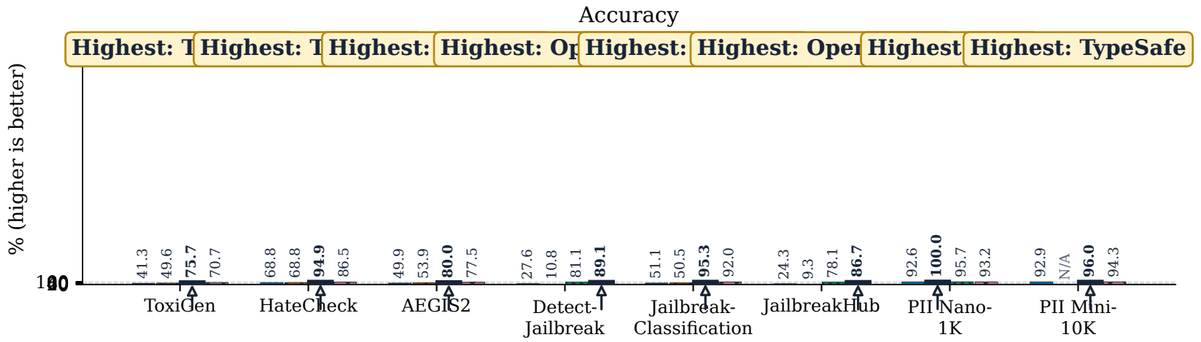
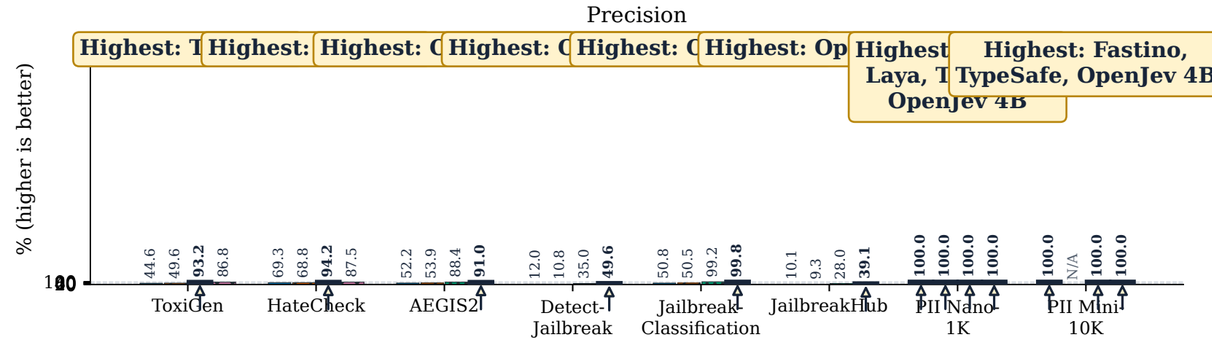
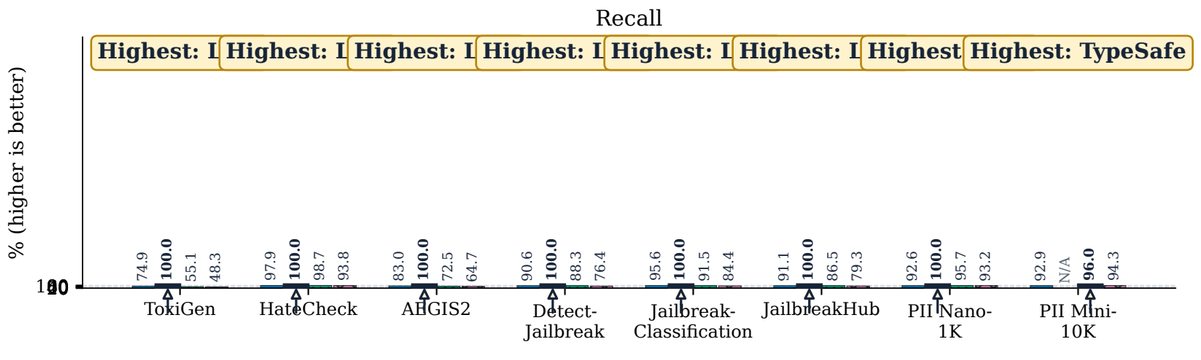
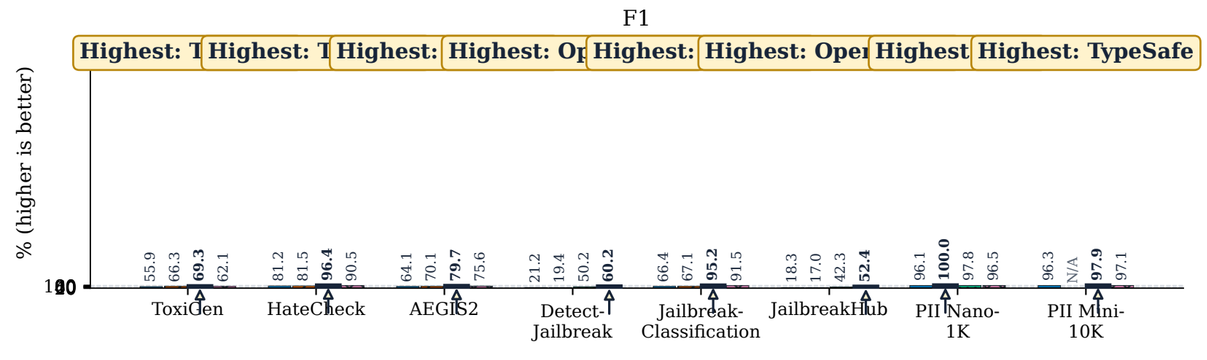
Per-Dataset Reference Comparison: v2 / Threshold 0.5 (Fastino, Laya, TypeSafe, OpenJev 4B)

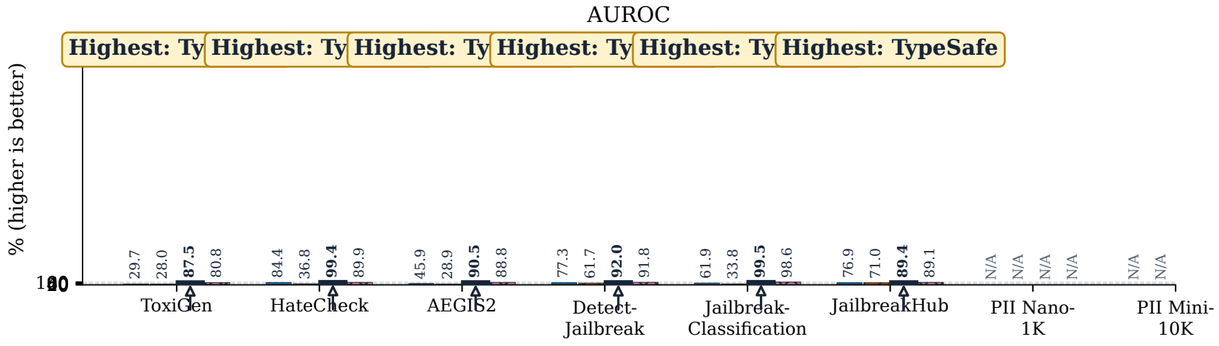
threshold = 0.5 | own successful samples | internal models: aggregates, NOT sample-paired



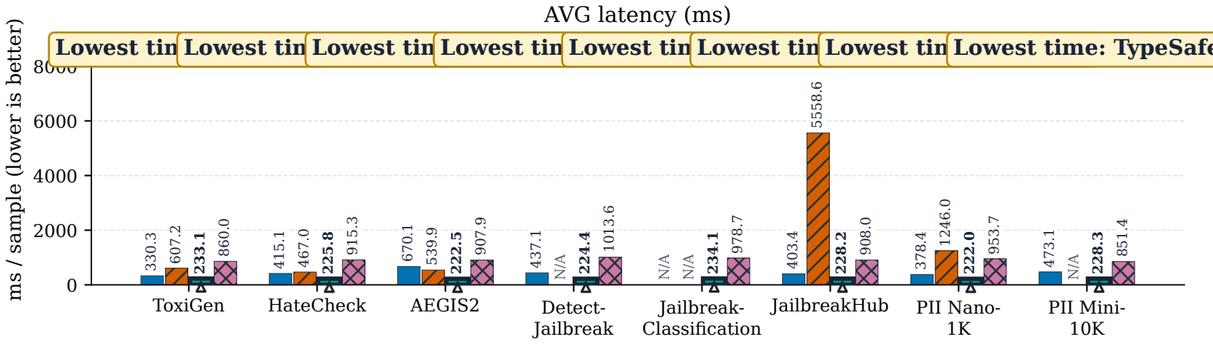








PII = ordinary binary classification (present vs absent), NOT masking quality.



OpenJev 4B: four-run summary timing; local models prefer four-run log averages. 4B valid/error counts are unknown.
N/A = missing quality, undefined metric, or incomplete timing; missing values are not zero.
Descriptive only, not a fair ranking. PII is ordinary binary classification (present vs absent), not masking quality.

Fastino | request: fastino/GLiDE
resolved: glide_v2

Laya | request: typed-decisions
serving: laya-rl-agent
routed: typed-decisions

TypeSafe | request: jev-latest
resolved: unconfirmed (API response required)

OpenJev 4B | supplied internal summary
model / version: not supplied (not sample-paired)