Four-model reference averages (not sample-paired)

threshold = 0.5 | shared non-PII dataset names | equal dataset weights | prompts kept separate

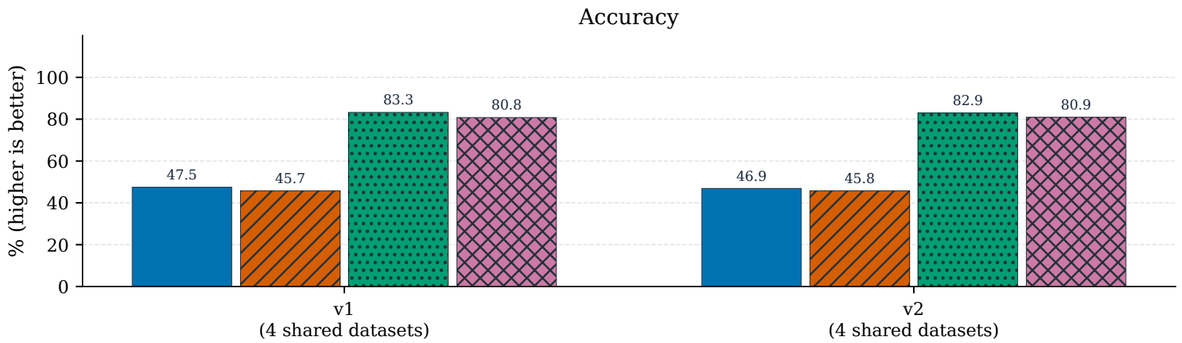

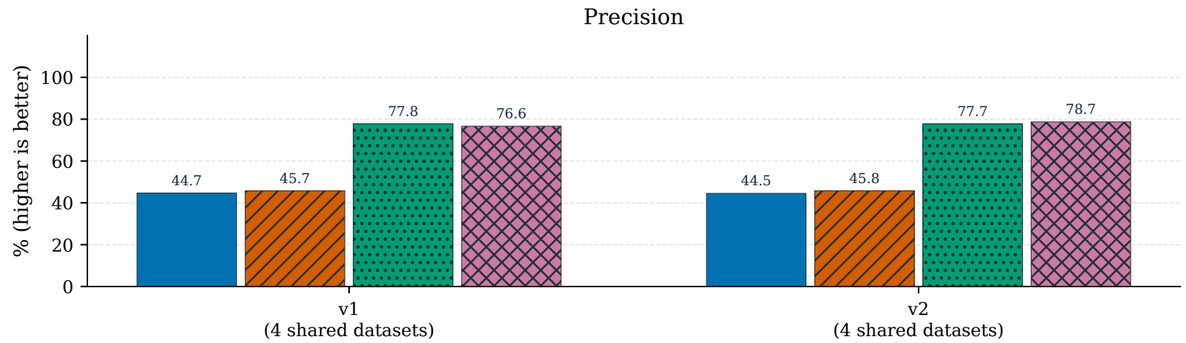

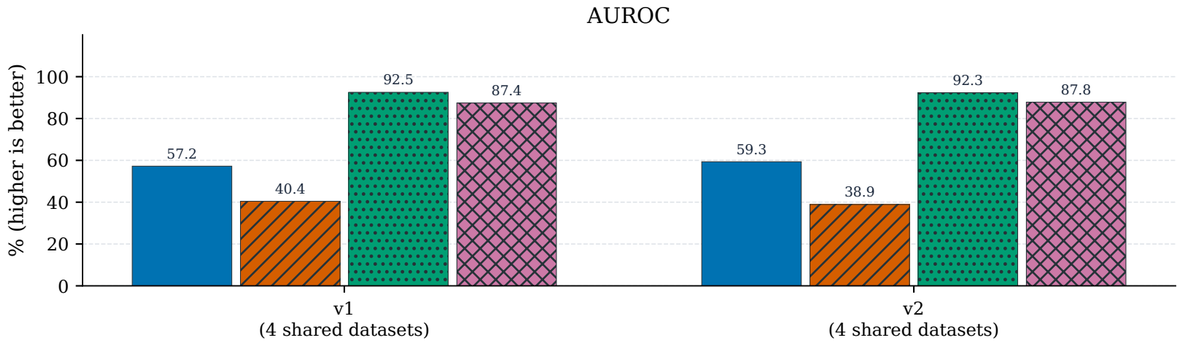

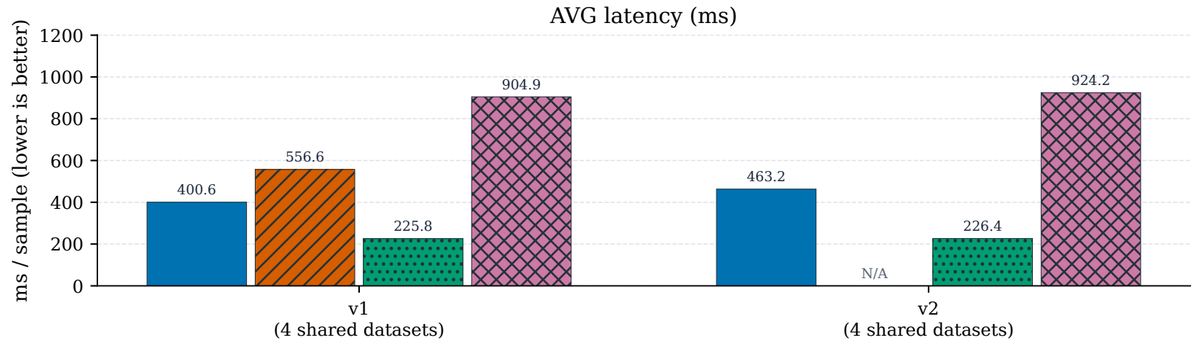

Each model uses its own successful samples; no common OpenJev sample IDs are available. Not a fair ranking.
Latency = mean per-run elapsed / processed across four complete runs, then an equal-weight dataset mean.
Missing any selected dataset's four-run timing makes average latency N/A; failure waiting is included.
OpenJev summary precision is retained as supplied; raw predictions and repeat-call stability are unavailable.

Fastino | request: fastino/GLiDE
resolved: glide_v2

Laya | request: typed-decisions
serving: laya-rl-agent
routed: typed-decisions

TypeSafe | request: jev-latest
resolved: unconfirmed (API response required)

OpenJev | supplied internal summary
model / version: not supplied (not sample-paired)