Model reference averages (not sample-paired)

v1 | threshold = 0.5 | shared non-PII dataset names | equal dataset weights | PII reported separately

Fastino | request: fastino/GLiDE
resolved: glide_v2

gliguard-LLMGuardrails-300M-fine-tuned
Fine-tuned · 300M
model ID: c94233af-98e6-4942-8e26-98ad204851bf

Laya | request: typed-decisions
serving: laya-rl-agent
routed: typed-decisions

TypeSafe | request: jev-latest
resolved: unconfirmed (API response required)

OpenJev 4B | supplied internal summary
model / version: not supplied (not sample-paired)

OpenJev 32B | supplied internal summary
single-run timing (not sample-paired)

Microsoft Decision-1 | supplied internal summary
single-run timing (not sample-paired)

Accuracy

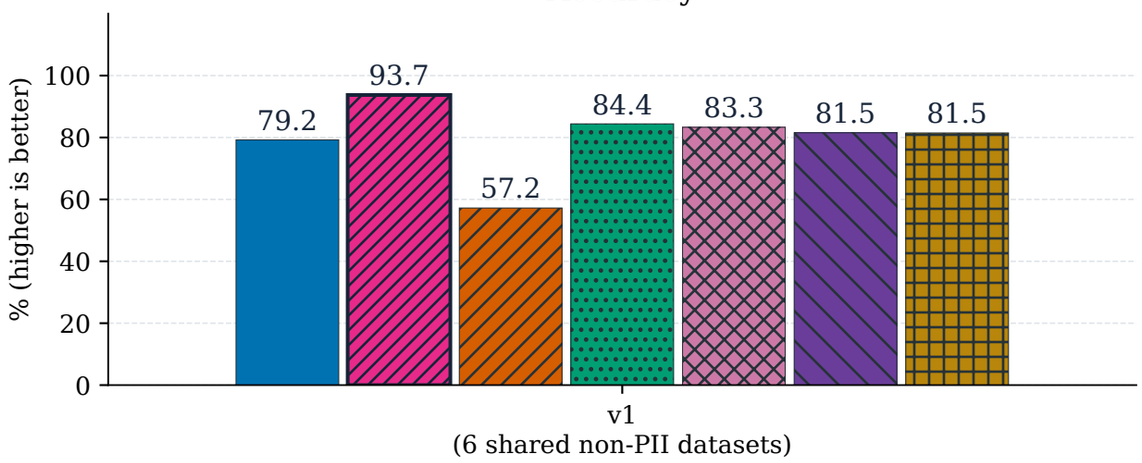

F1

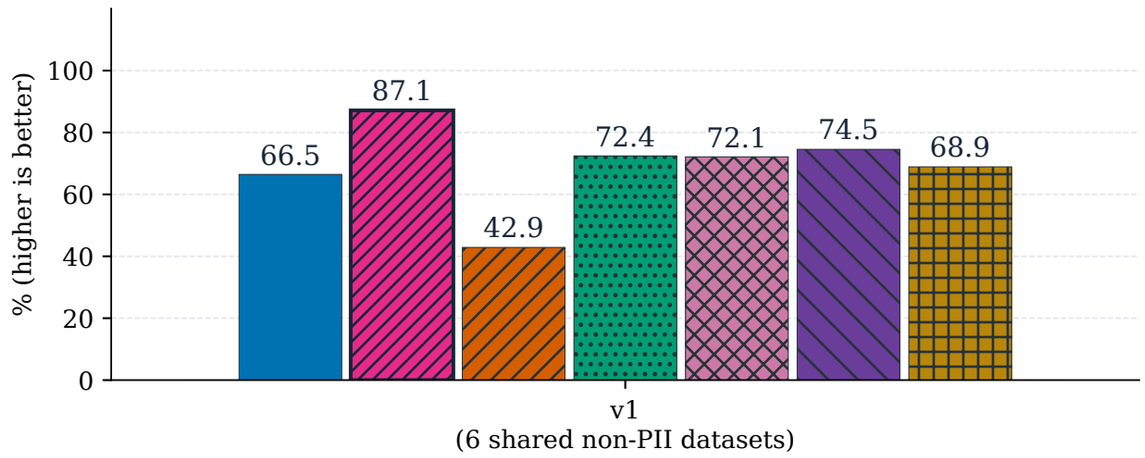

Precision

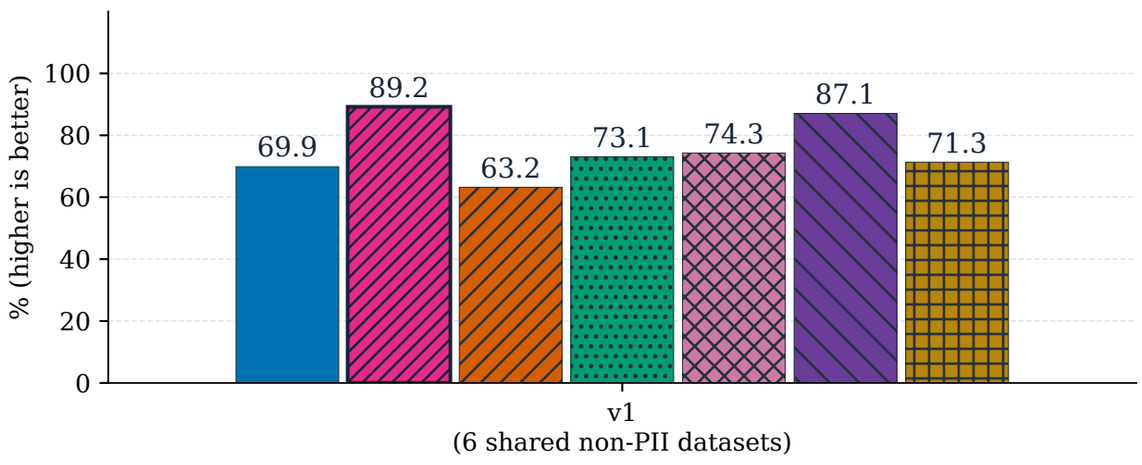

Recall

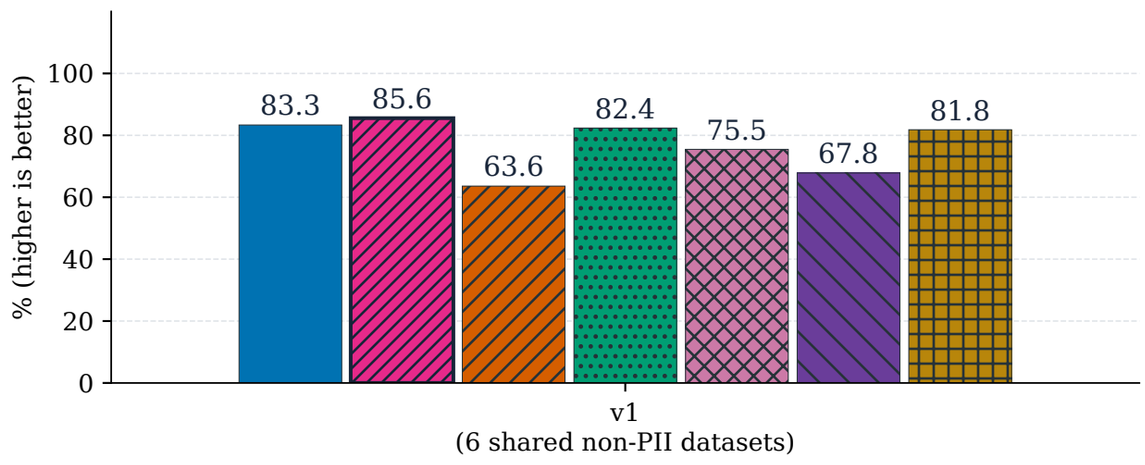

AUROC

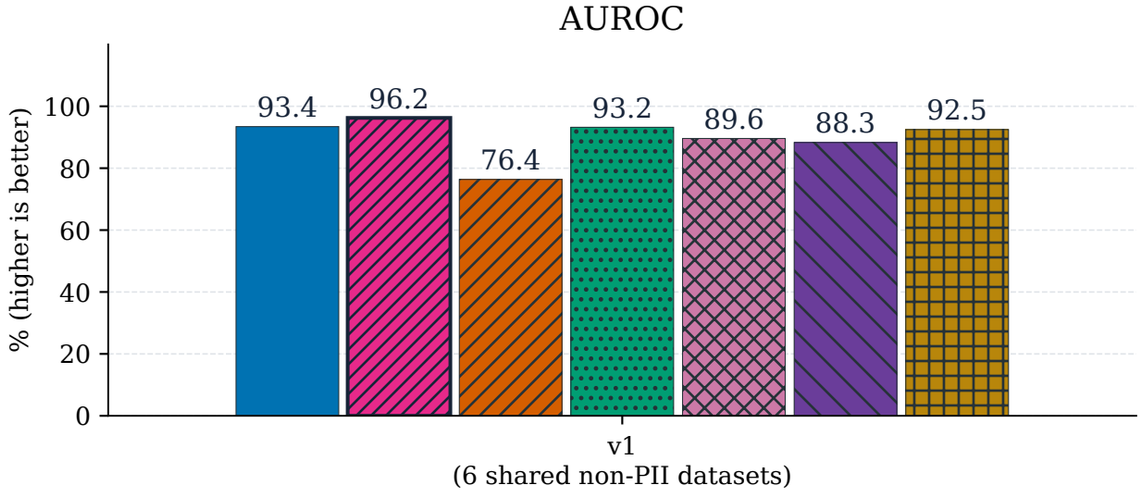

AVG latency (ms)

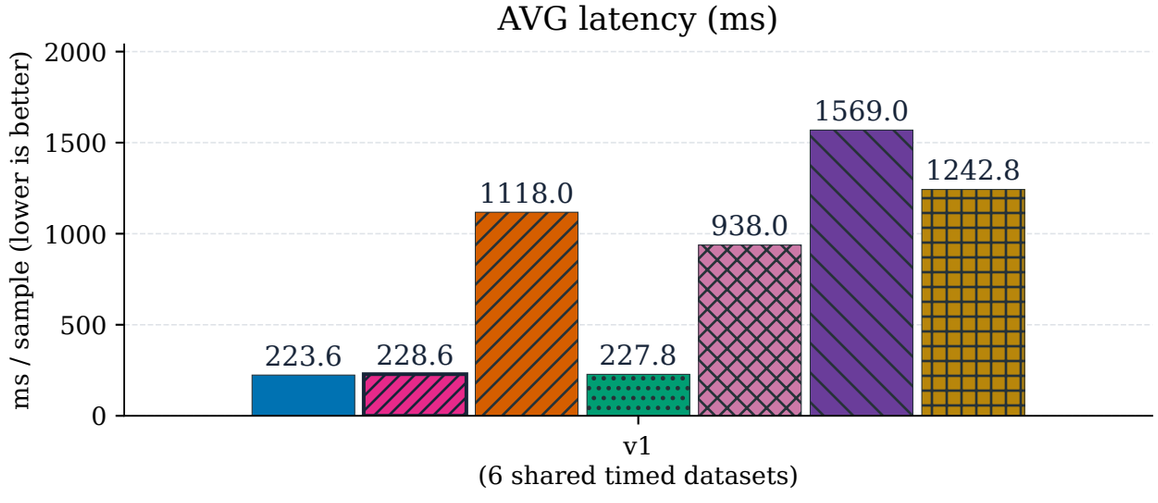

Each model uses its own successful samples; aggregate models have no paired sample IDs. Not a fair ranking.
Latency: equal dataset weights on the SAME fully timed v1 / 0.5 subset for all models:
ToxiGen, HateCheck, AEGIS2, Detect-Jailbreak, Jailbreak-Classification, JailbreakHub. Quality and latency dataset scopes can differ.
APIs: v1 / 0.5 single-run timing. OpenJev 4B: supplied summary; 32B / Decision-1: supplied single-run timing.
PII is ordinary binary classification (present vs absent), not masking quality; positives-only results are separate.
Missing quality stays N/A; no shared timing subset = N/A latency. Timing protocols differ; not inference time.