ToxiGen, HateCheck, AEGIS2: Multi-model Reference Averages | Luna: native presets (prompt N/A)

Computed averages: ToxiGen / HateCheck / AEGIS2 | supplied-model dataset scopes unverified

Accuracy

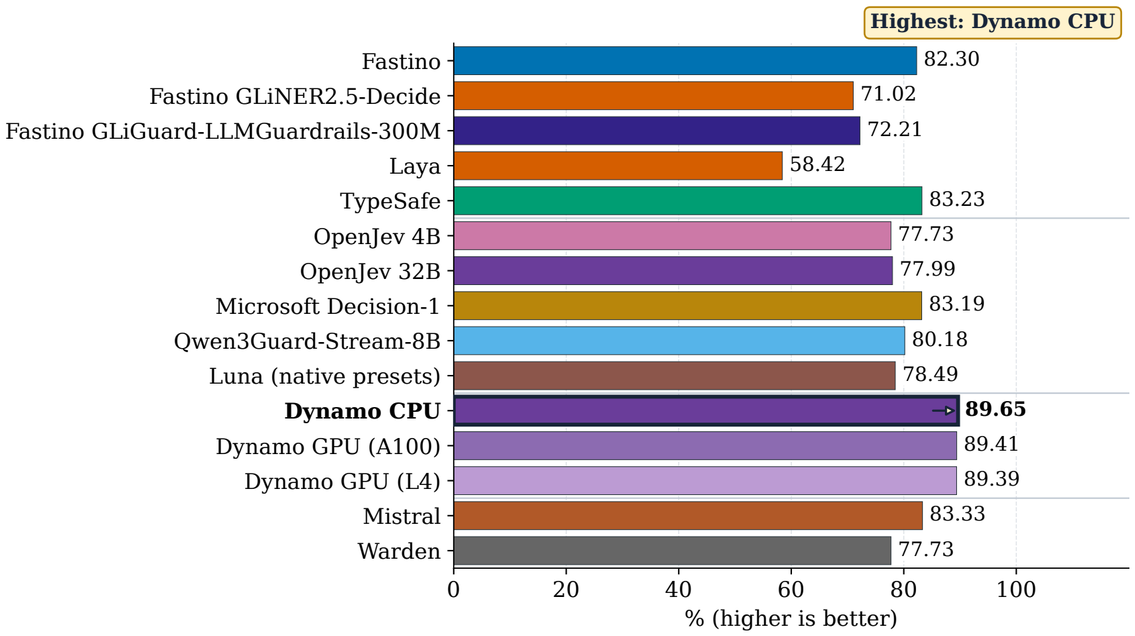

Precision

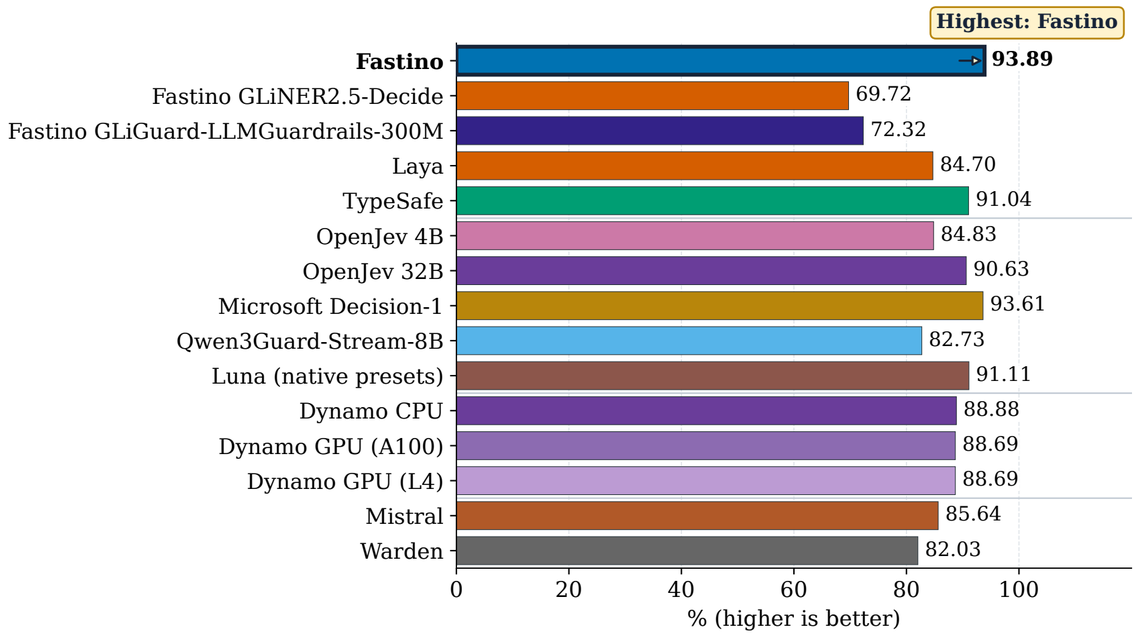

Recall

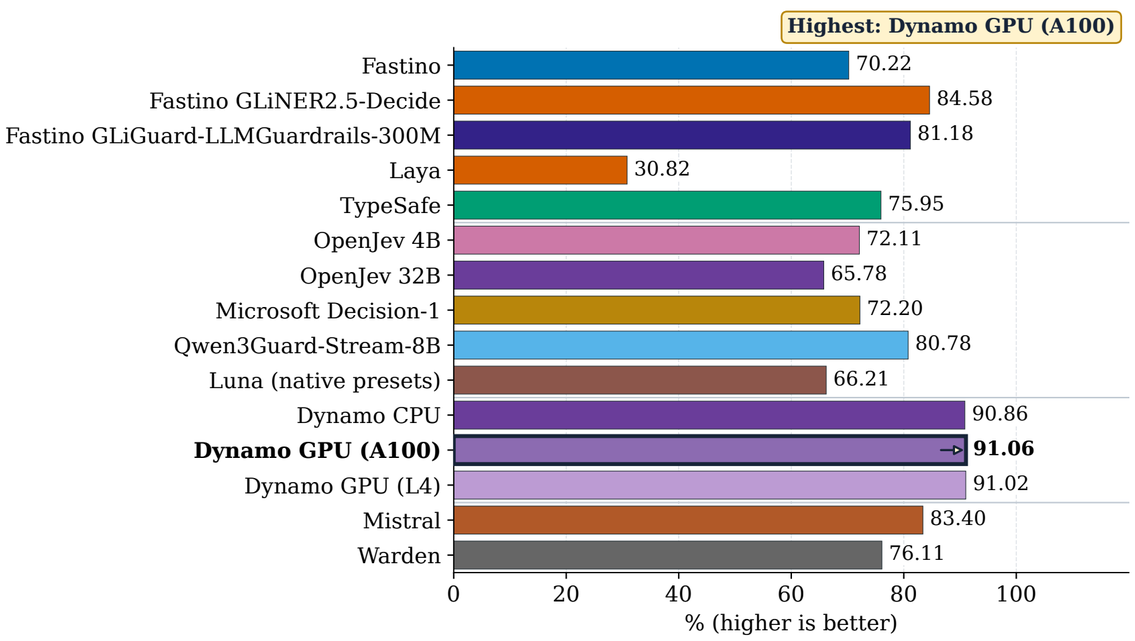

F1

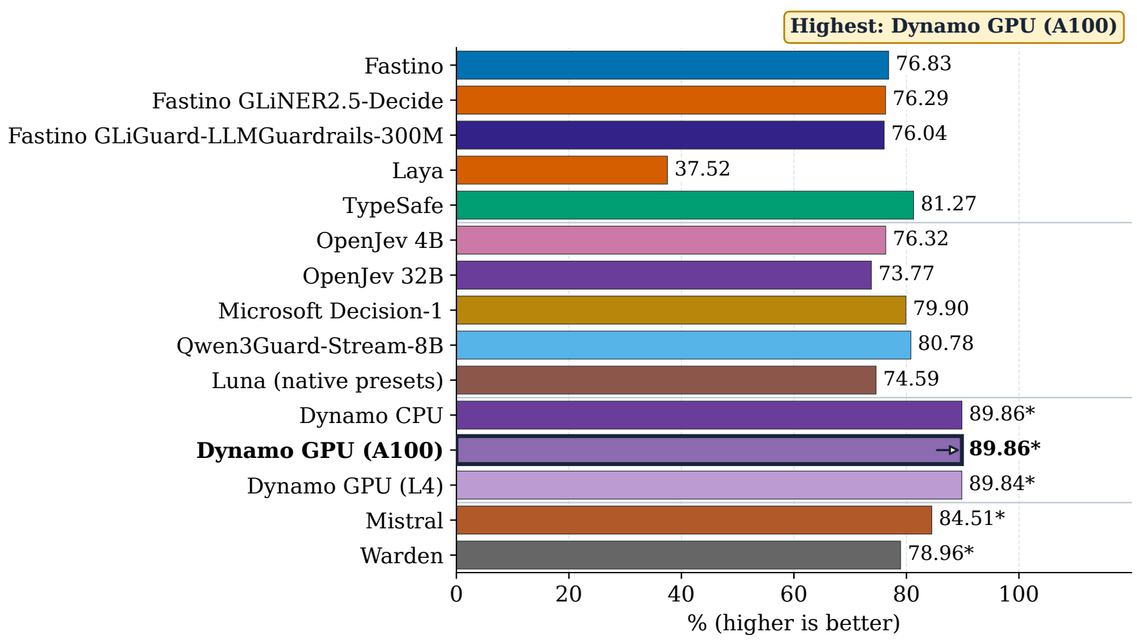

AUROC

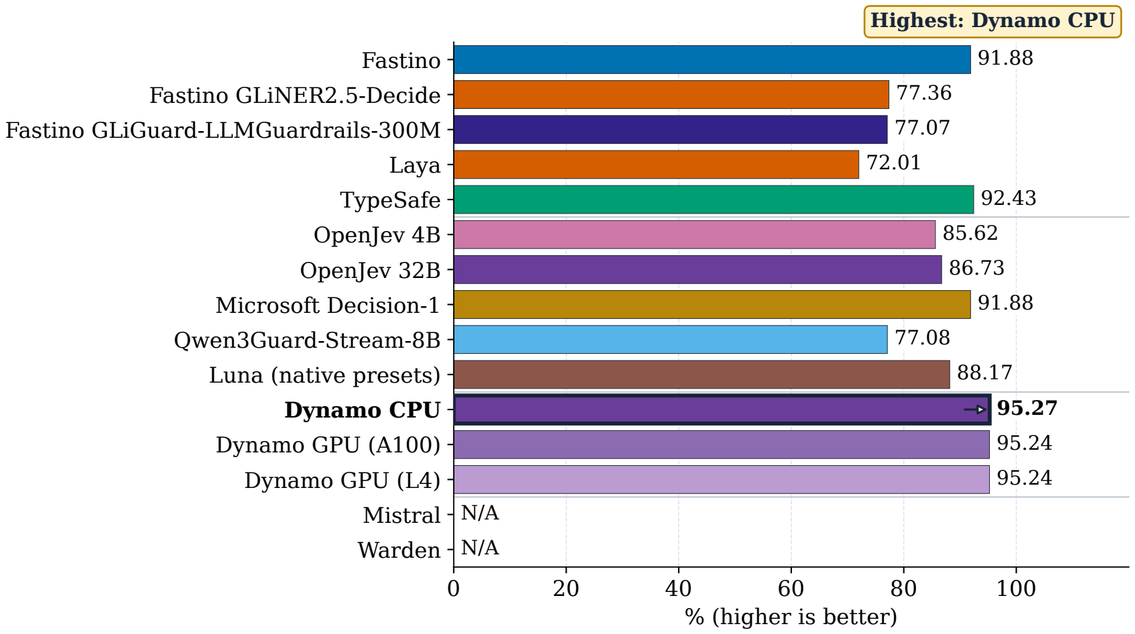

AVG latency (ms)

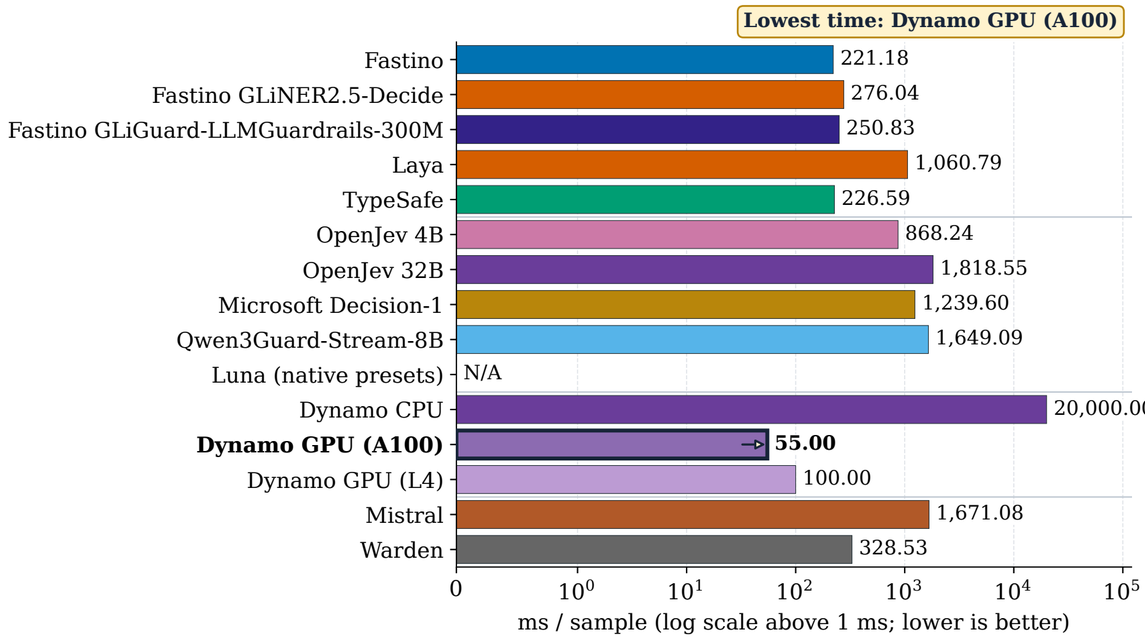

Computed models: no internal dataset was available; v1 / 0.5 means use only the three shared datasets.
Local latency uses v1 / 0.5 runs; OpenJev 4B / 32B and Decision-1 use their supplied summary latency.
Dynamo / Mistral / Warden: supplied averages; their evaluations may also include an internal dataset.
Its inclusion in their averages is unconfirmed; prompt, threshold and timing protocols were not supplied.
* F1 = 2PR / (P + R) from supplied averages; not a mean of per-dataset F1. N/A = unavailable.
Successful samples and timing protocols are not aligned; descriptive reference only, not a fair ranking.
Qwen3Guard-Stream-8B: binary safety labels; PII unsupported (historical proxies are audit-only).
Qwen mean latency: log elapsed/processed (failures included); CSV P95: successful requests. No ranking.
Luna: original task-specific presets, no custom prompts (prompt N/A); binary threshold 0.5.
Other models use dataset-specific v1/v2 prompts. Task-definition differences may affect results.
Luna PII micro-F1 is separate; binary PII and latency are N/A. Supplied metrics; scorer versions unverified.