Three-dataset model reference averages

ToxiGen / HateCheck / AEGIS2 | equal dataset weights | not sample-paired

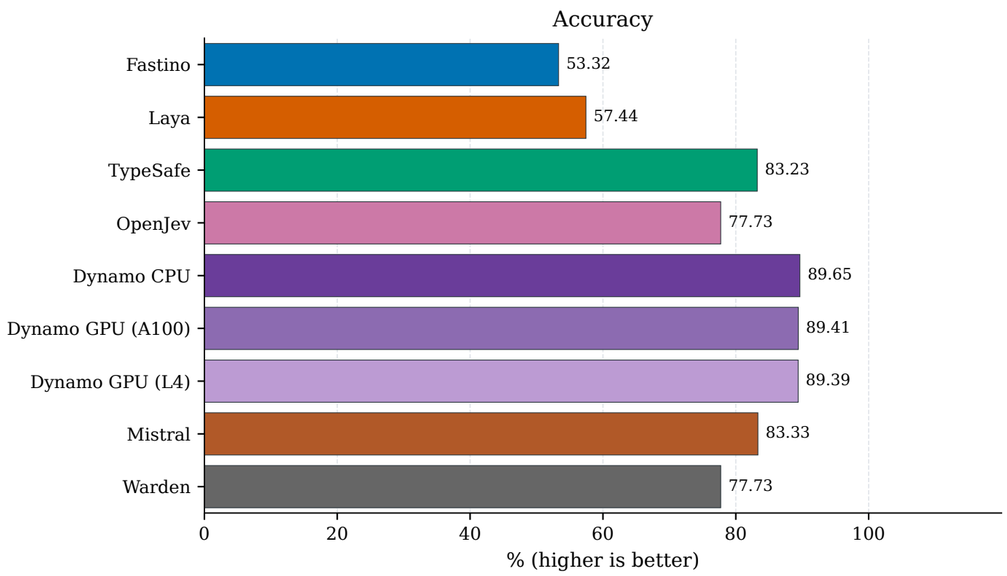

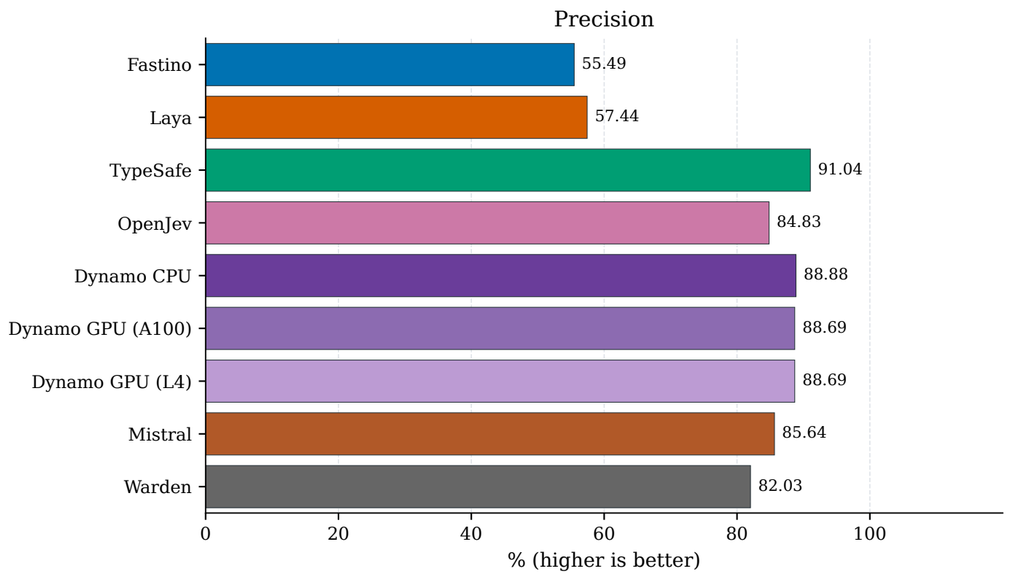

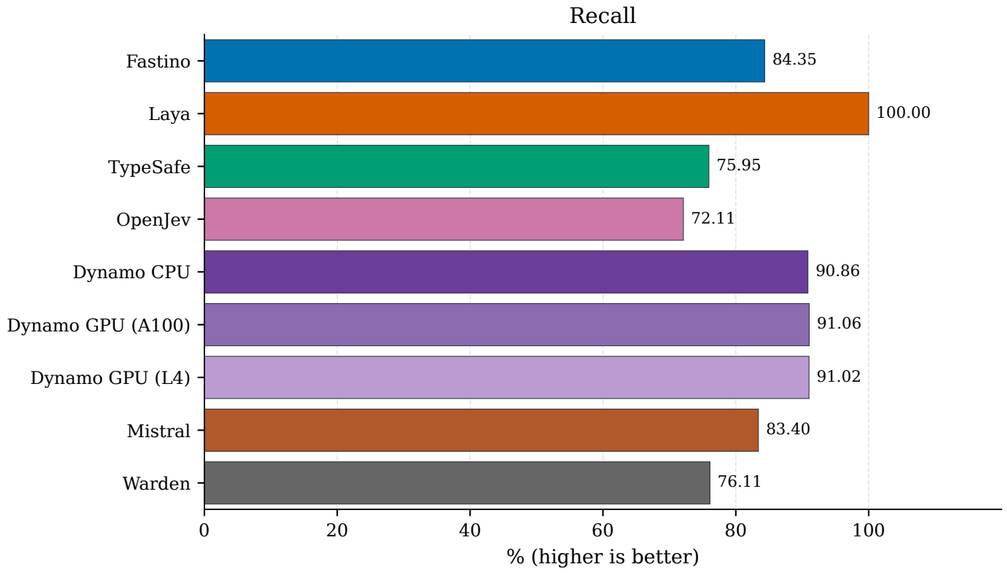

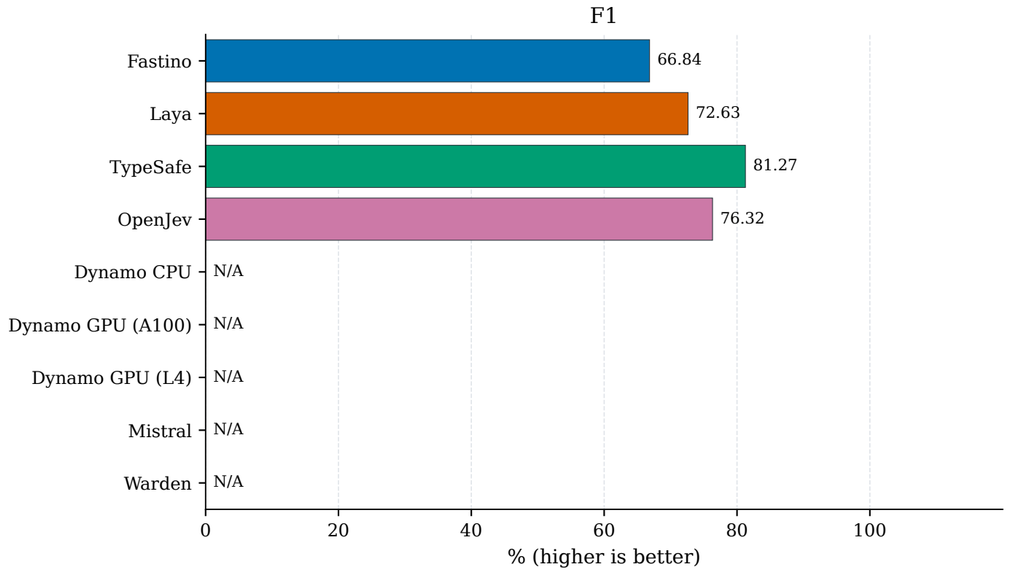

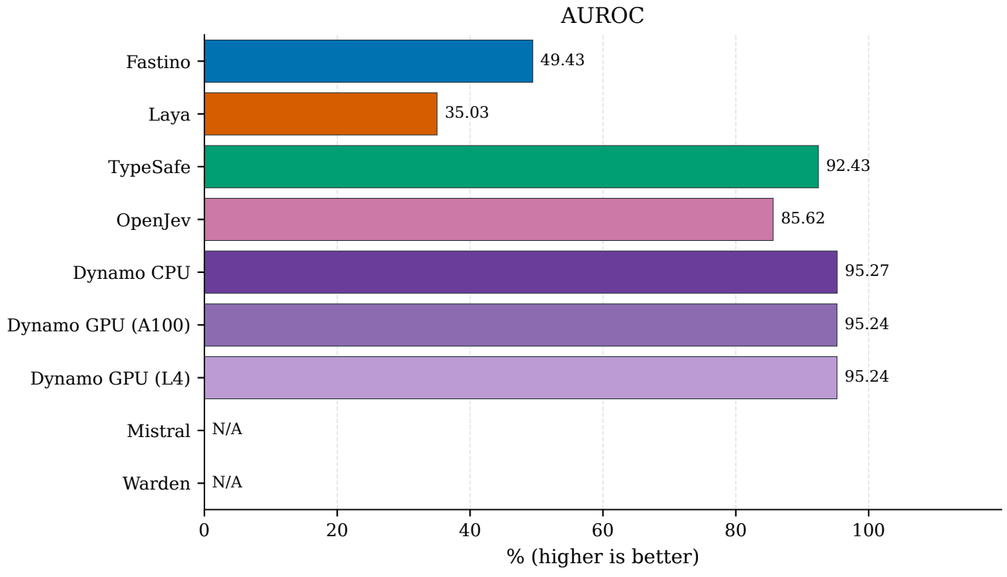

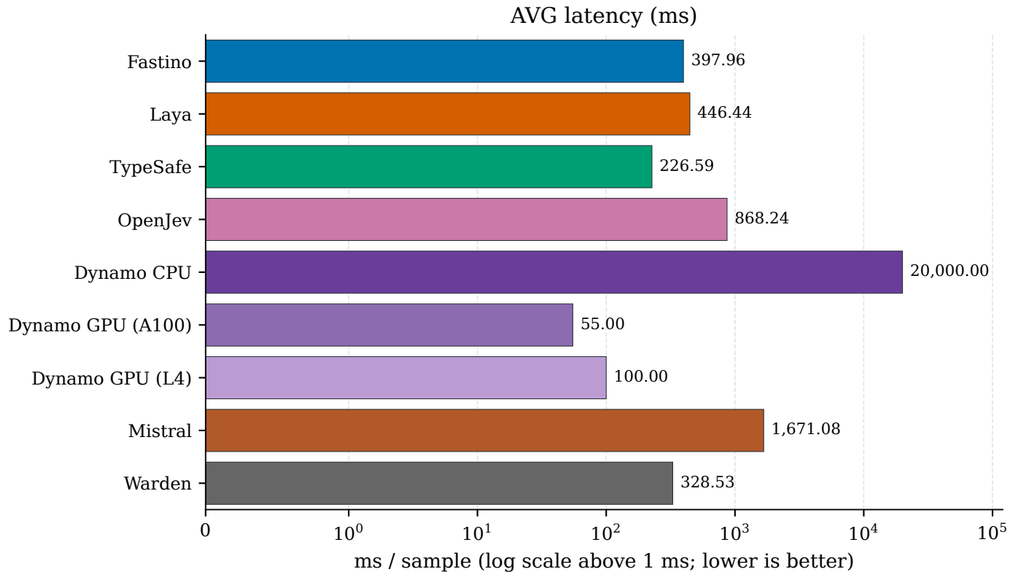

Original four models: prompt v1, threshold 0.5; quality and latency are equal-weight dataset means.
Local latency uses only the v1 / 0.5 run per dataset; OpenJev uses its supplied summary latency.
Dynamo / Mistral / Warden: user-supplied three-dataset means, scope confirmed by user (original '4' was a typo).
Their prompt, threshold and timing protocol were not supplied; means are shown once.
N/A = not supplied or required source data missing. F1 is not inferred from average precision / recall.
Successful samples and timing protocols are not aligned; descriptive reference only, not a fair ranking.