Dataset-Category Multi-Model Reference Averages (v1 / Threshold 0.5) | Luna: native presets (prompt N/A)

v1 / threshold = 0.5 | equal dataset weights within each category | NOT sample-paired

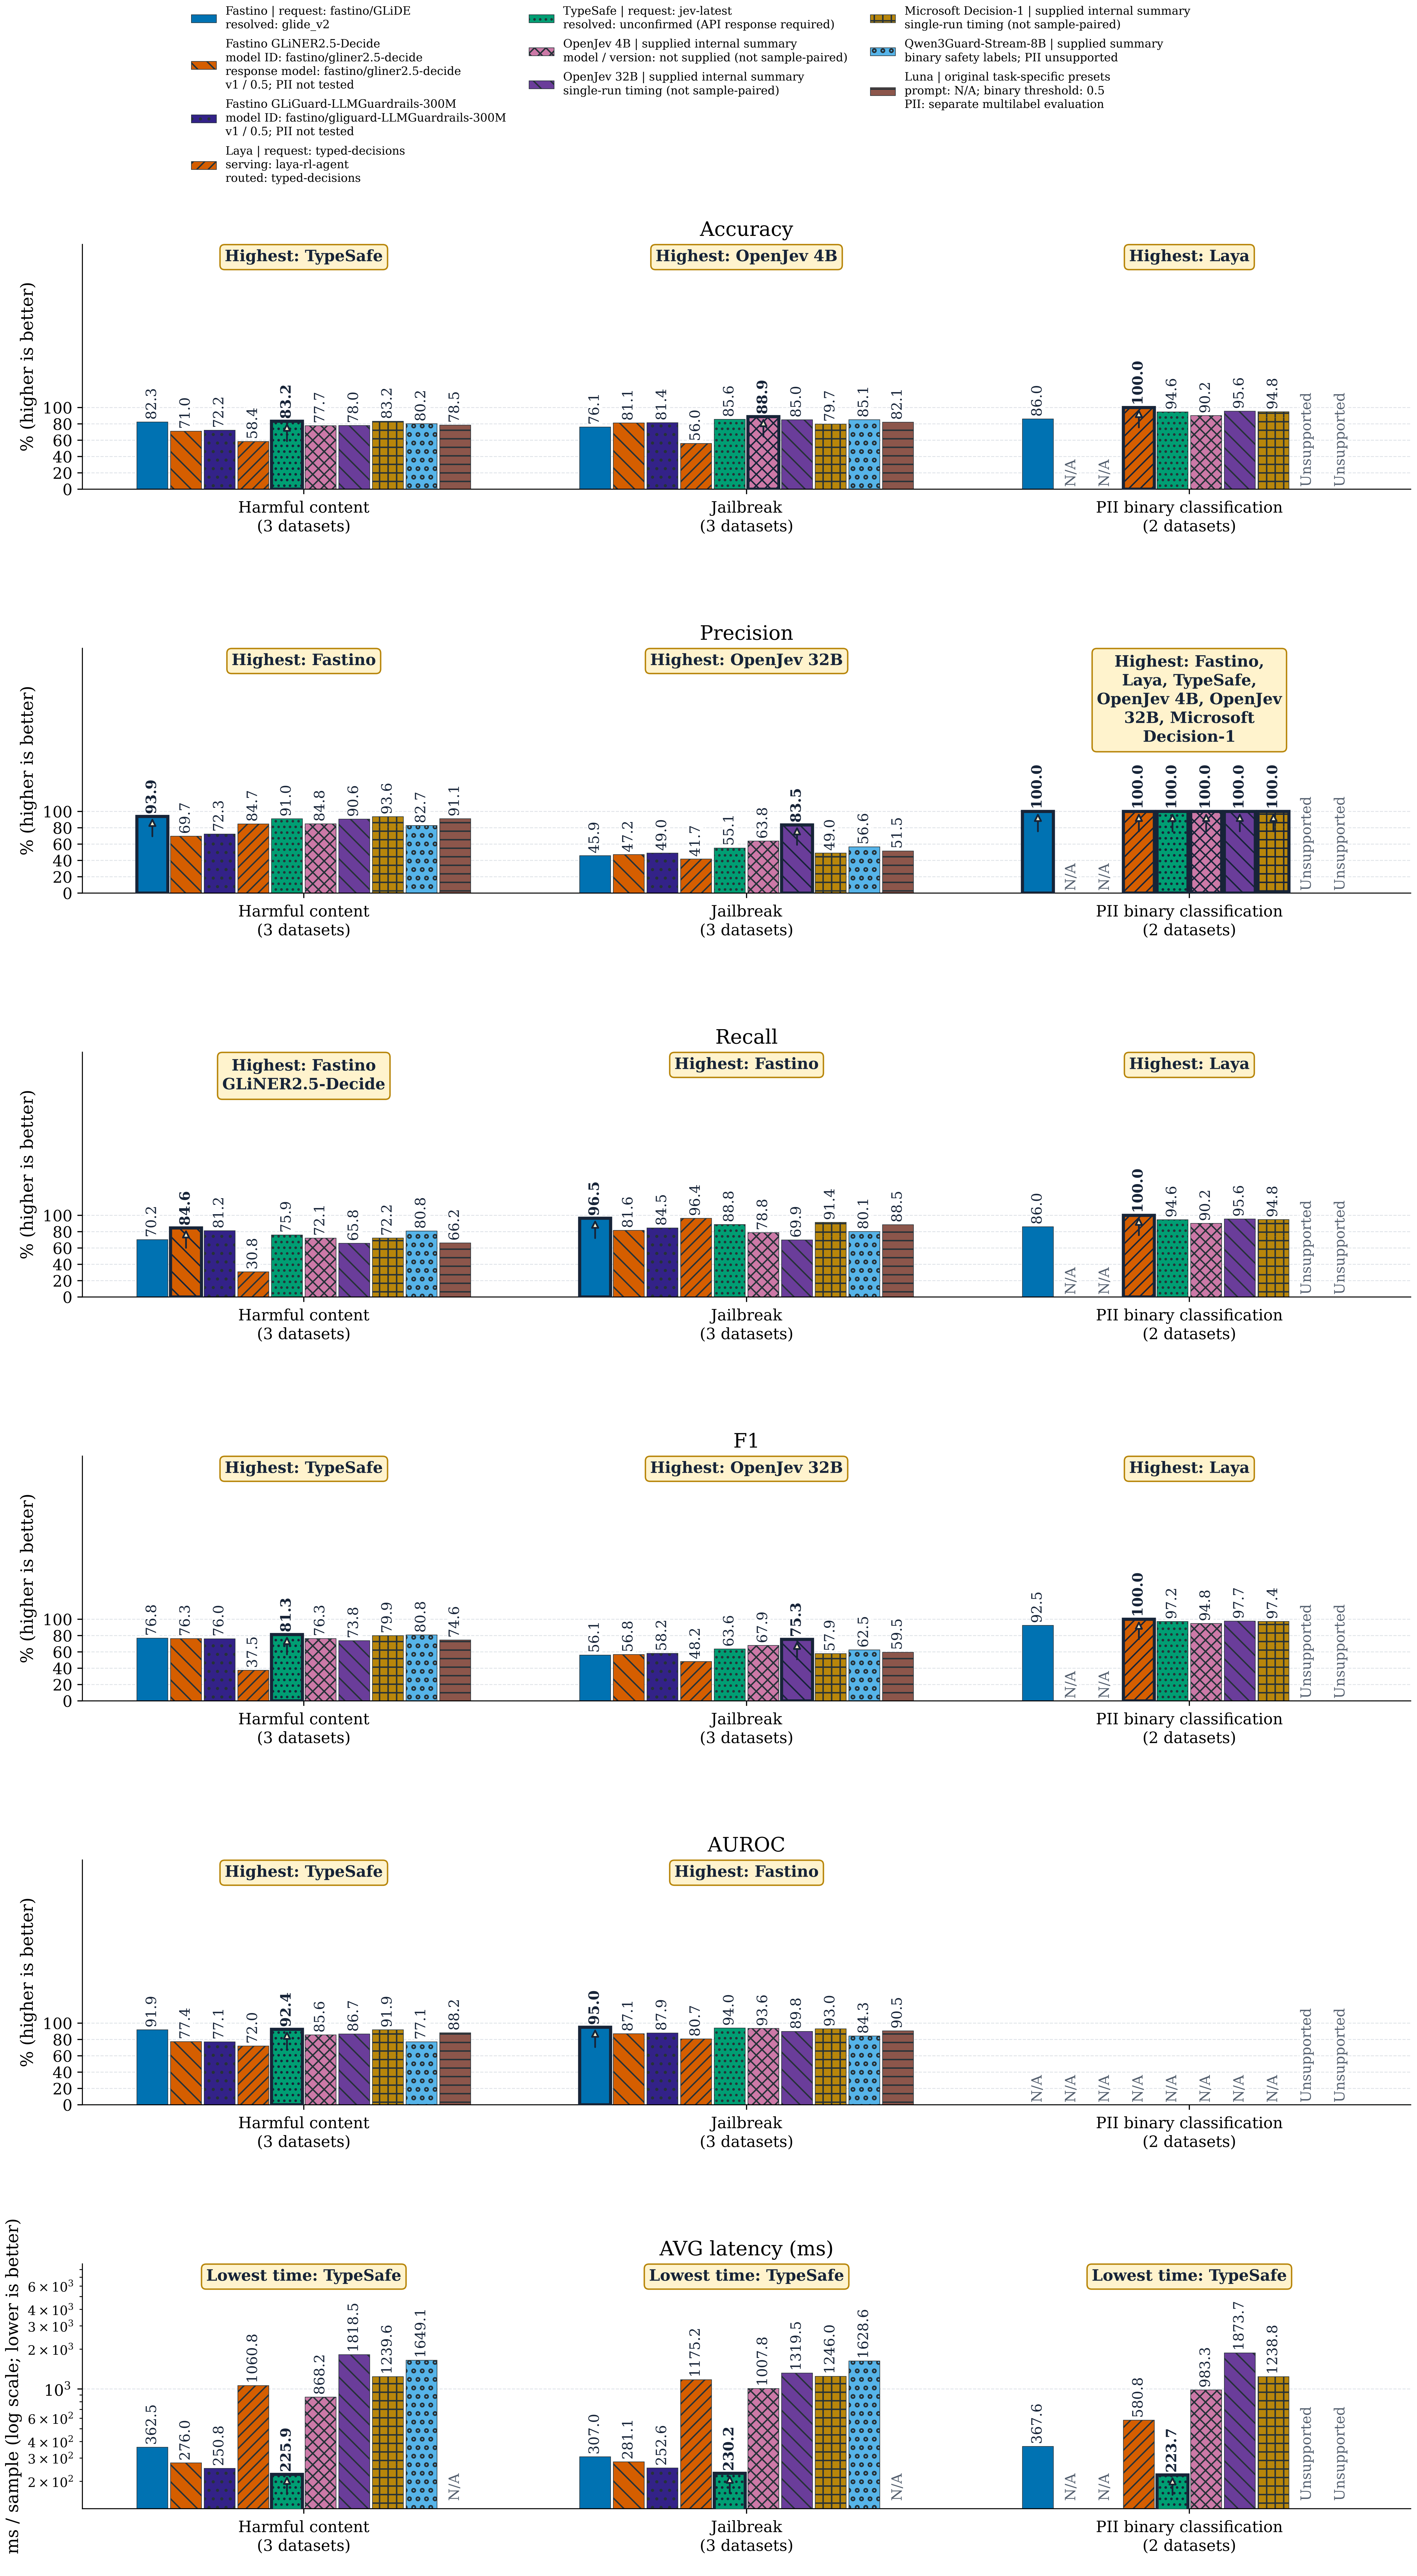

PII = ordinary binary classification (present vs absent), NOT masking quality.

Harmful content: ToxiGen, HateCheck, AEGIS2. Jailbreak: Detect-Jailbreak, Jailbreak-Classification, JailbreakHub.
PII: Nano-1K, Mini-10K; ordinary binary classification (present vs absent), not extraction or masking quality.
Current PII datasets contain positives only; AUROC is undefined.
Each metric is the arithmetic mean of per-dataset metrics, not pooled samples or F1 recomputed from means.
Missing required data or undefined metrics = N/A; category scopes are never reduced to available datasets.
Latency averages the same reference timings as the per-dataset chart; protocols differ and include failure waits.
Own successful samples / supplied aggregates; descriptive only, not a fair ranking.
Qwen3Guard-Stream-8B: binary safety labels; PII unsupported (historical proxies are audit-only).
Qwen mean latency: log elapsed/processed (failures included); CSV P95: successful requests. No ranking.
Luna: original task-specific presets, no custom prompts (prompt N/A); binary threshold 0.5.
Other models use dataset-specific v1/v2 prompts. Task-definition differences may affect results.
Luna PII micro-F1 is separate; binary PII and latency are N/A. Supplied metrics; scorer versions unverified.