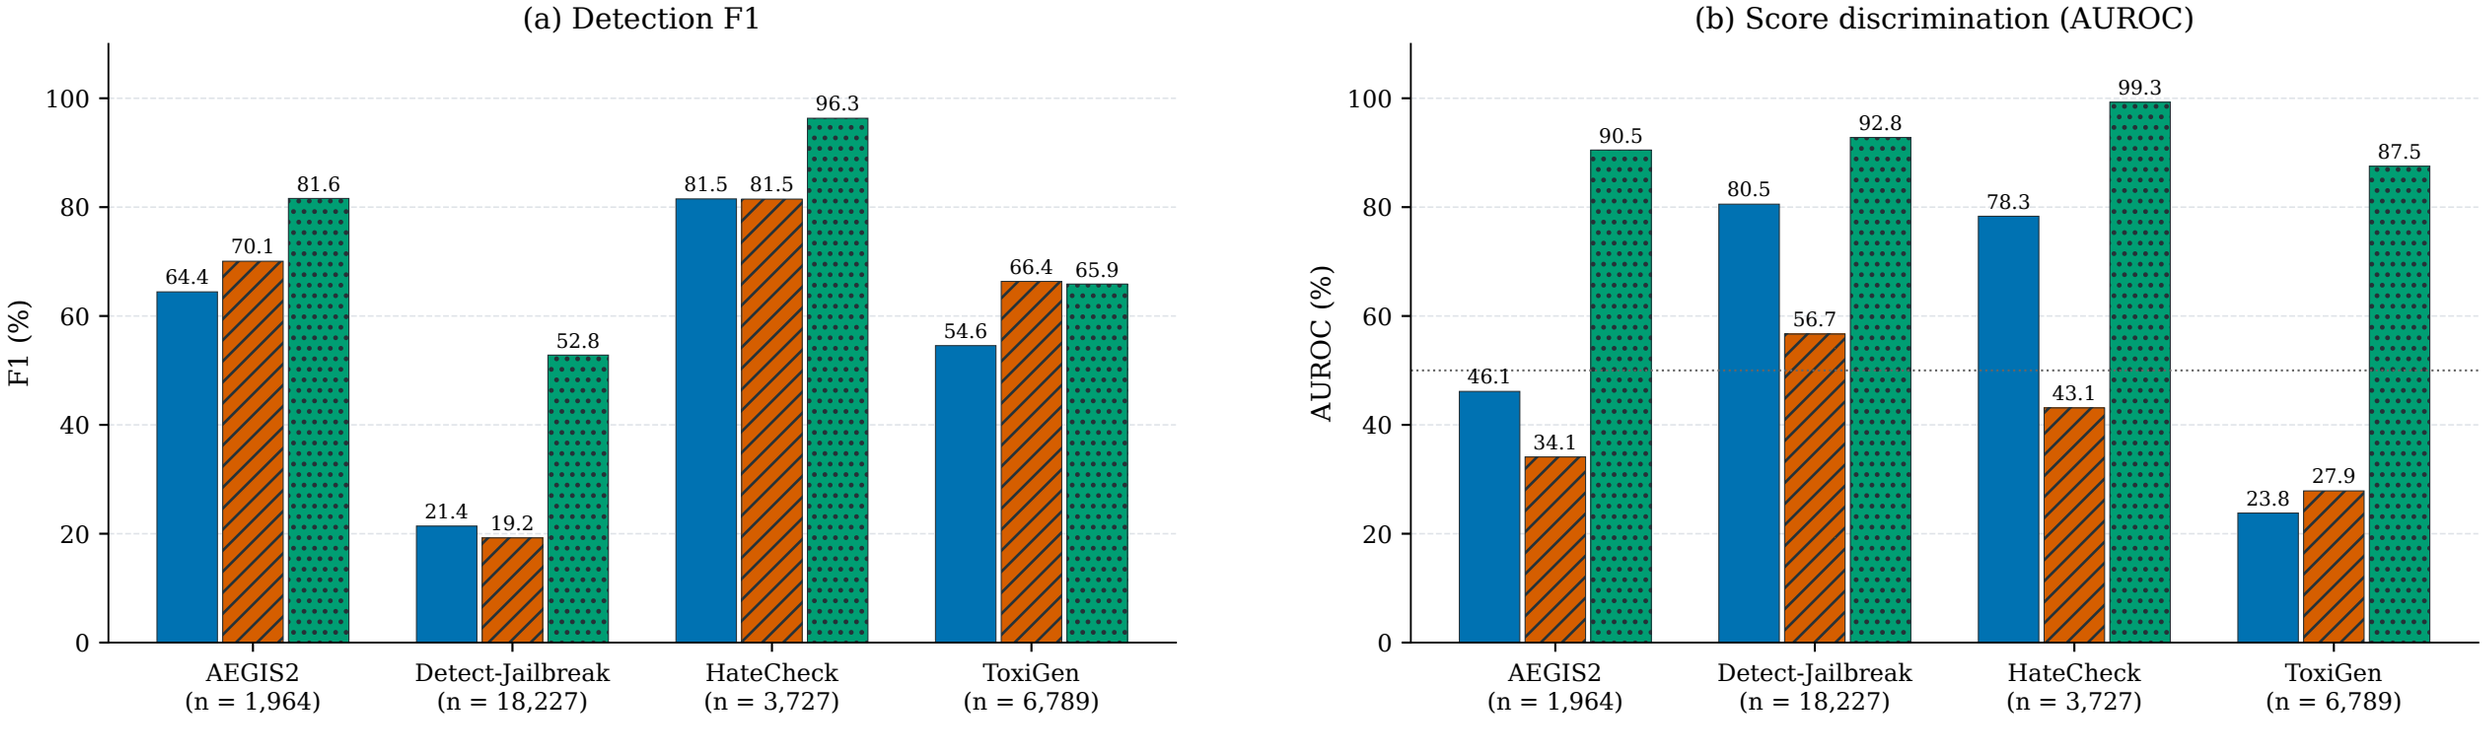
Safety classification: paired baseline

v1 | threshold = 0.5 | common completed non-PII datasets | higher is better



Each dataset uses the same successful sample IDs for all selected models.
Values are descriptive point estimates; repeated sweeps are not independent replicates.

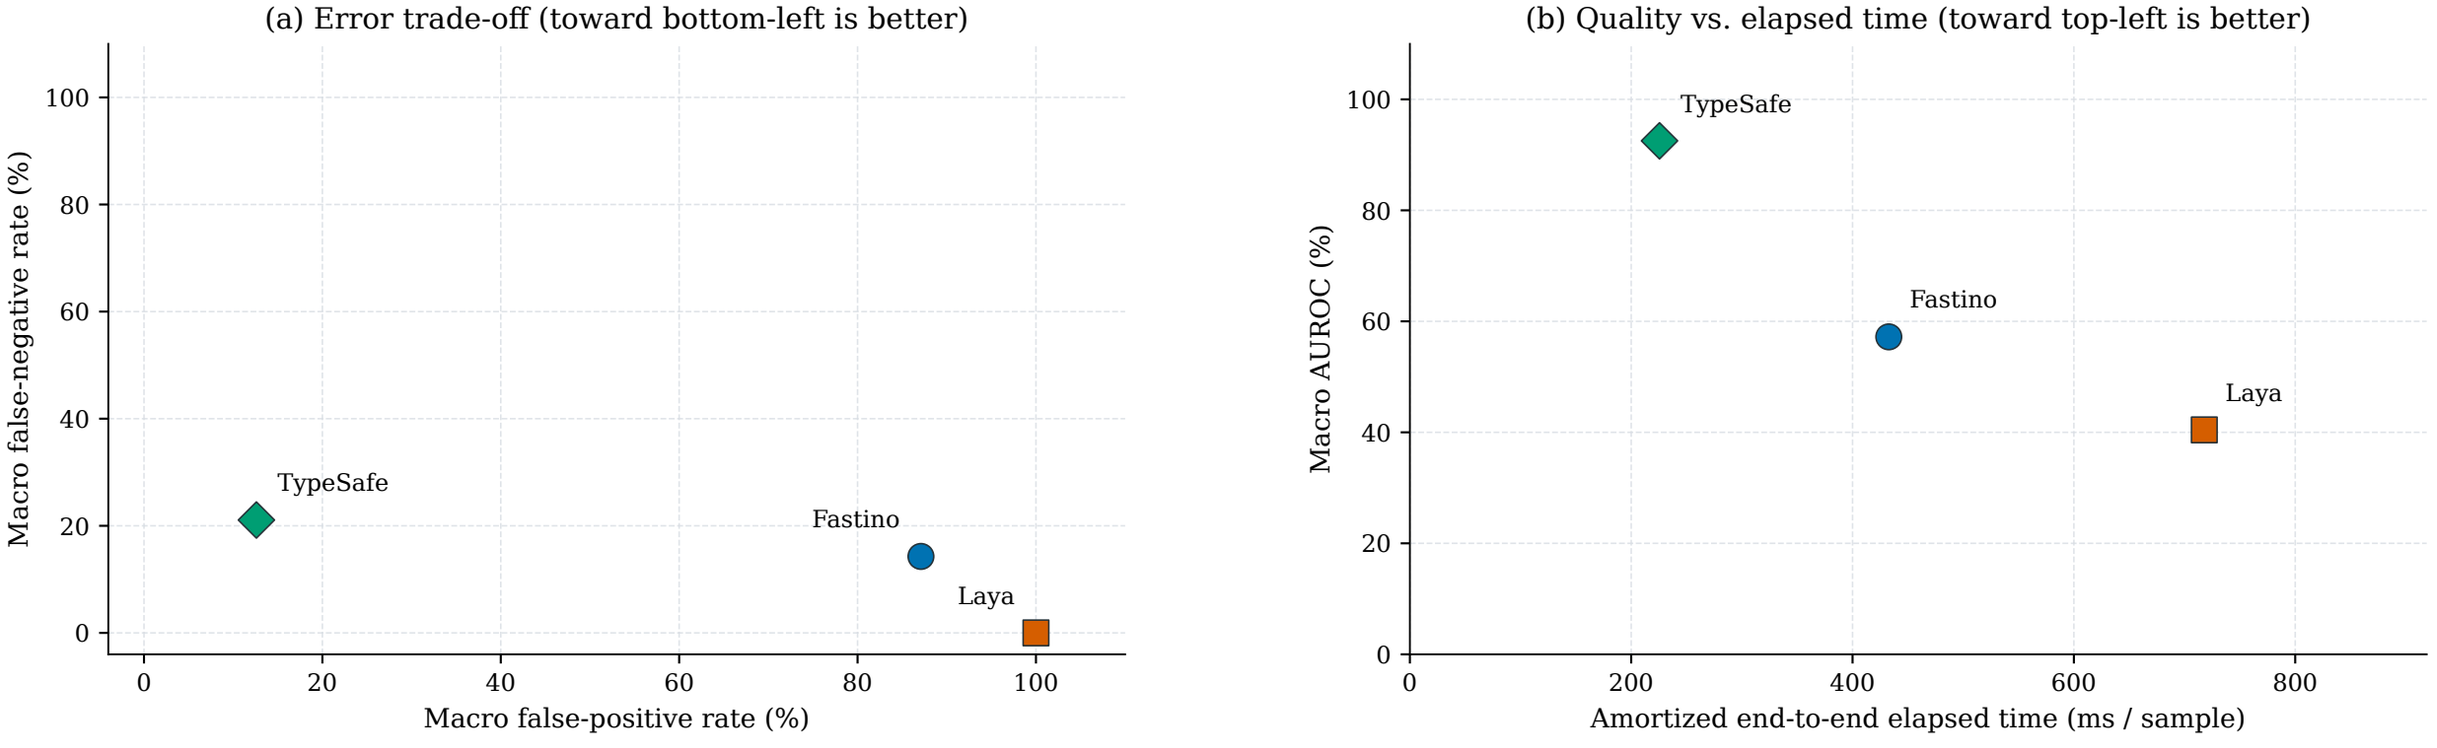
Fastino | request: fastino/GLiDE
resolved: glide_v2

Laya | request: typed-decisions
serving: laya-rl-agent
routed: typed-decisions

TypeSafe | request: jev-latest
resolved: unconfirmed (API response required)

Baseline quality, errors and elapsed time

v1 | threshold = 0.5 | 4 shared datasets | dataset-macro quality



Elapsed time = total logged elapsed / all processed baseline rows, including failed requests.
Runs occurred at different times; this is not pure inference latency or a request-level percentile.

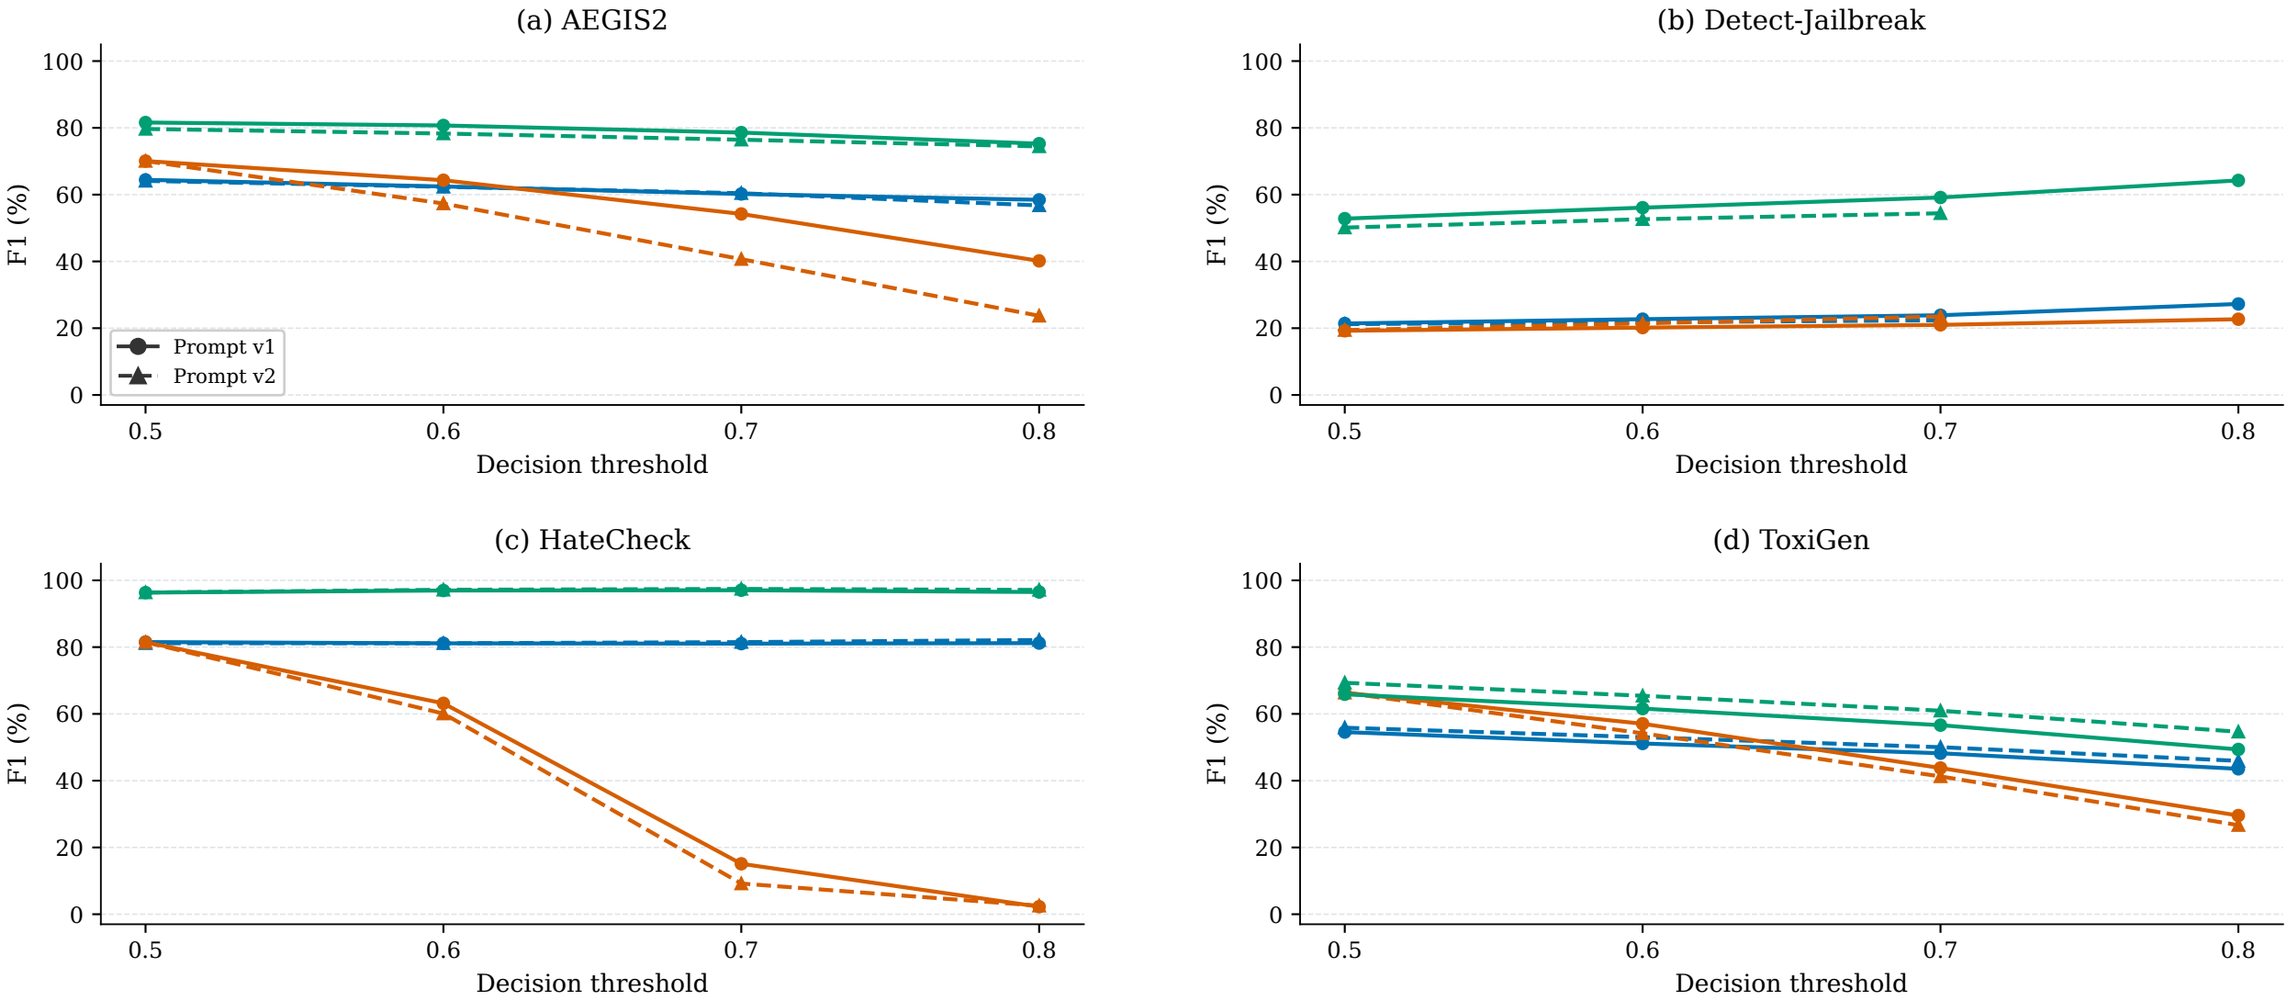
Fastino | request: fastino/GLiDE
resolved: glide_v2

Laya | request: typed-decisions
serving: laya-rl-agent
routed: typed-decisions

TypeSafe | request: jev-latest
resolved: unconfirmed (API response required)

Sensitivity to prompt version and decision threshold

Exploratory test-set sweeps, not held-out model selection



Points use within-configuration paired successful samples; error-driven sample sets can differ across points.
Missing configurations are gaps, not zero scores. Curves mix threshold effects and repeat-call variation.
Scores are not held fixed across thresholds; lines do not imply independent replicates or confidence intervals.

Fastino | request: fastino/GLiDE
resolved: glide_v2

Laya | request: typed-decisions
serving: laya-rl-agent
routed: typed-decisions

TypeSafe | request: jev-latest
resolved: unconfirmed (API response required)

Baseline evaluation coverage and request failures

v1 | threshold = 0.5 | counts are successful / reference rows



Gray = no saved predictions; hatched = complete CSV with no successful requests, not a missing run.
PII rows measure only evaluation coverage, not masking quality. Laya's suite is incomplete.

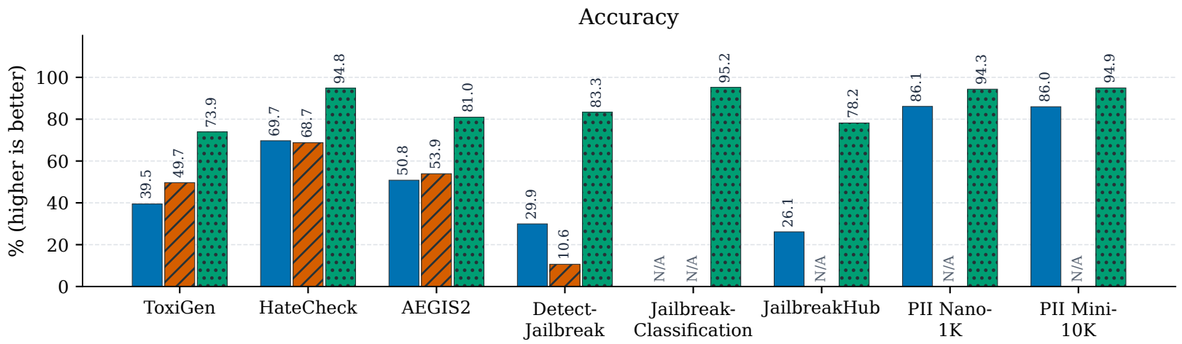
Fastino | request: fastino/GLiDE
resolved: glide_v2

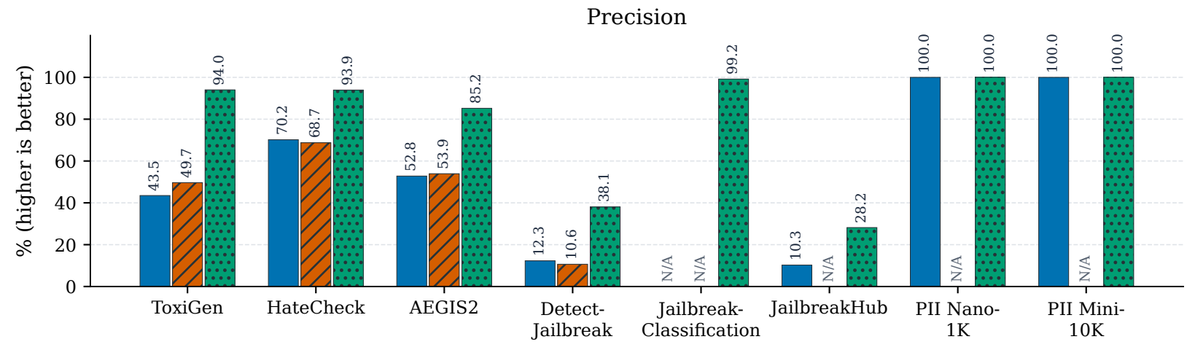
Laya | request: typed-decisions
serving: laya-rl-agent
routed: typed-decisions

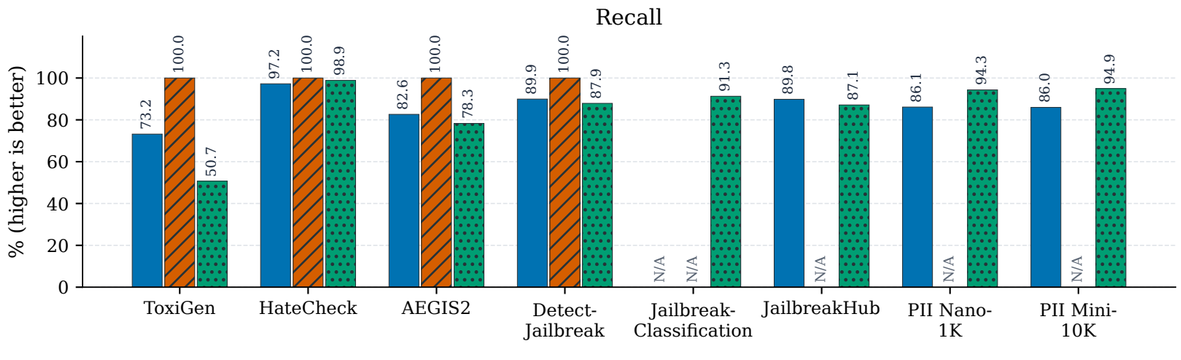
TypeSafe | request: jev-latest
resolved: unconfirmed (API response required)

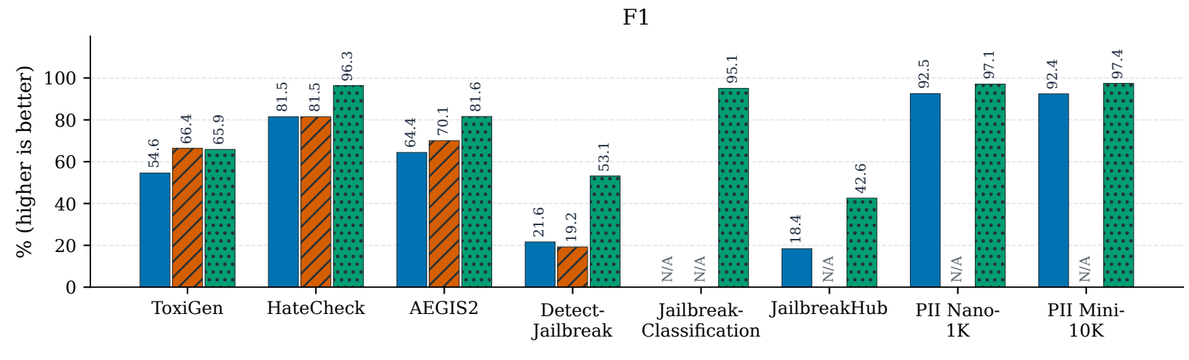
Per-dataset classification metrics and average latency

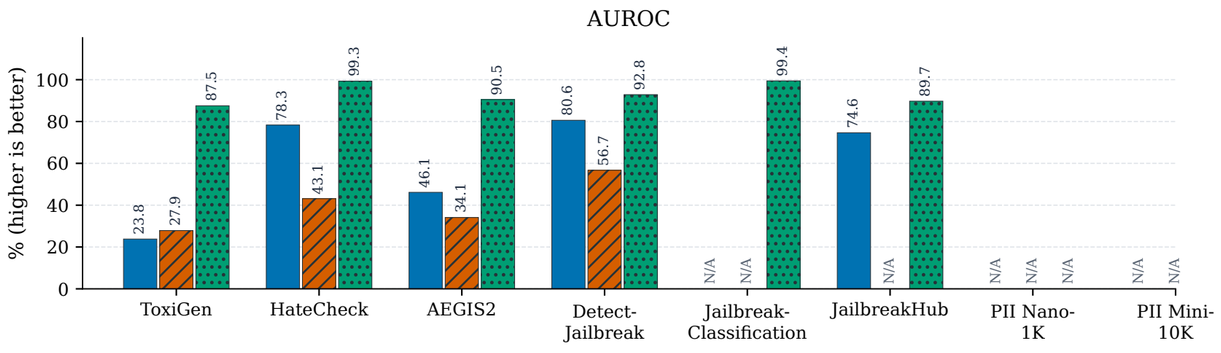
v1 | threshold = 0.5 | each model's own successful samples (not paired)

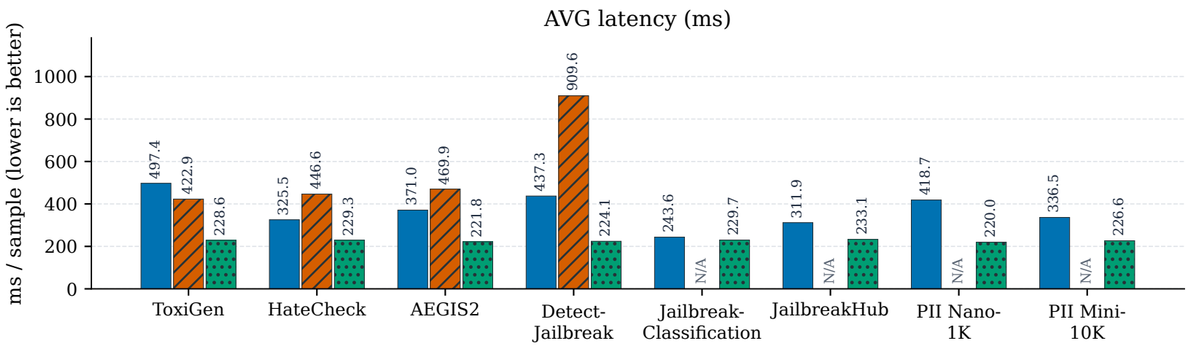












N/A = missing result, no successful predictions, undefined metric, or no aligned complete timing log.
PII datasets contain positives only: detection metrics do not measure masking quality; AUROC is undefined.
Latency = logged end-to-end elapsed / all processed rows, including failed requests; not pure inference time.
Successful sample sets can differ between models; use the paired baseline for fair quality comparisons.

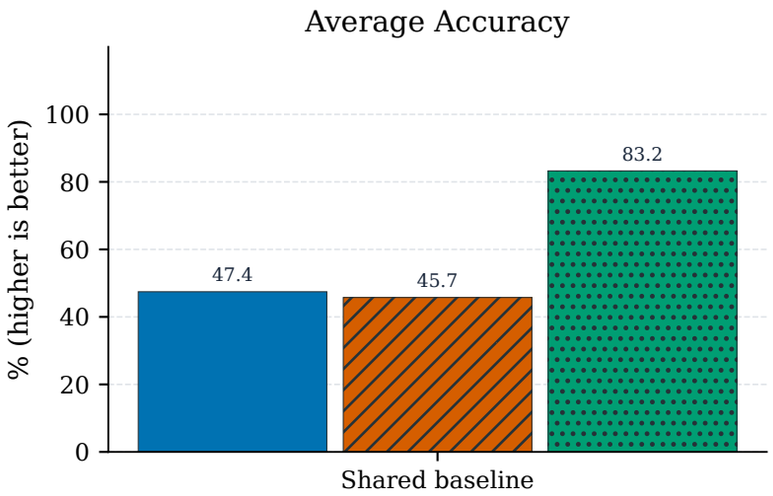
Fastino | request: fastino/GLiDE
resolved: glide_v2

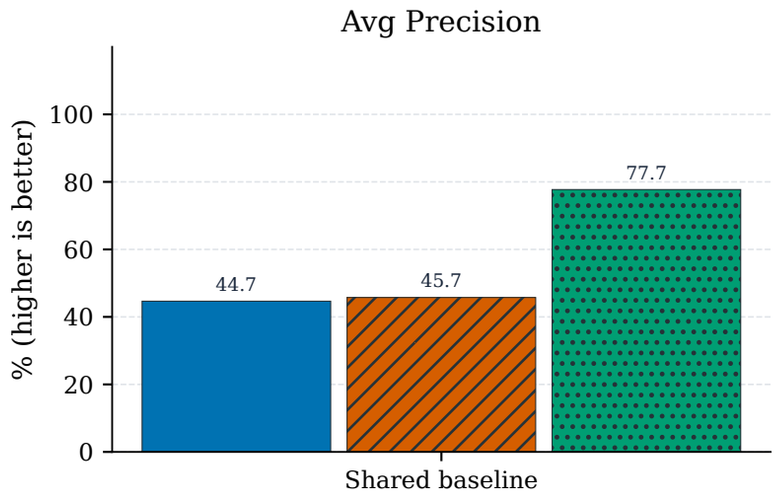
Laya | request: typed-decisions
serving: laya-rl-agent
routed: typed-decisions

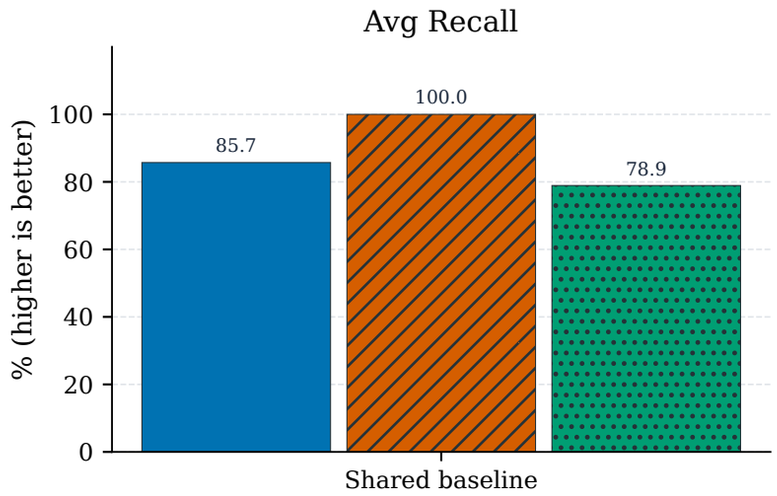
TypeSafe | request: jev-latest
resolved: unconfirmed (API response required)

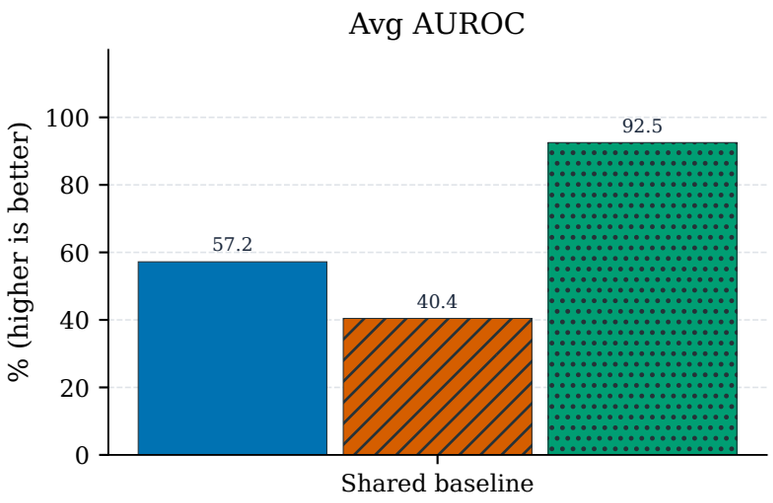
Model averages on the shared baseline

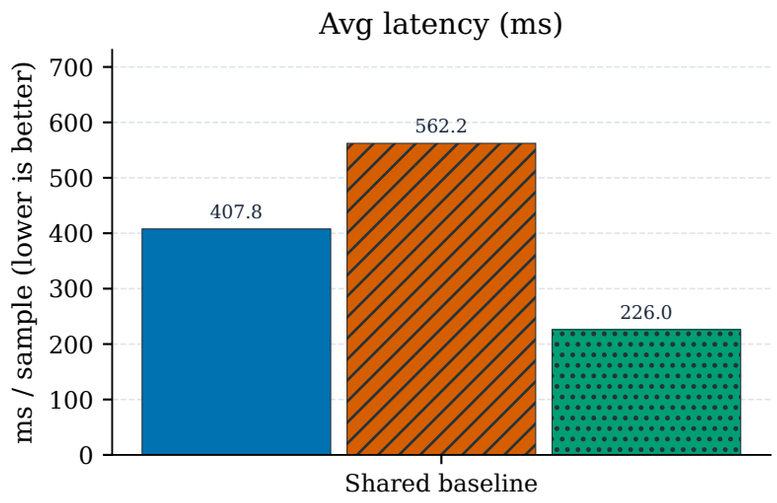
v1 | threshold = 0.5 | 4 shared non-PII datasets | equal dataset weights











Quality: common successful sample IDs within each dataset, then an unweighted mean across datasets.
Latency: unweighted mean of per-dataset log elapsed / processed rows, including failures (not inference time).
N/A timing means at least one baseline dataset has no aligned complete timing log; missing is never zero.

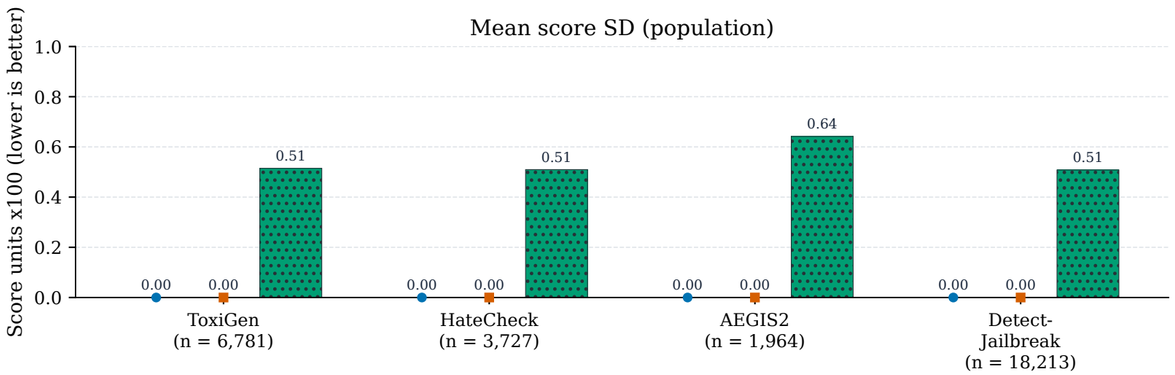
Fastino | request: fastino/GLiDE
resolved: glide_v2

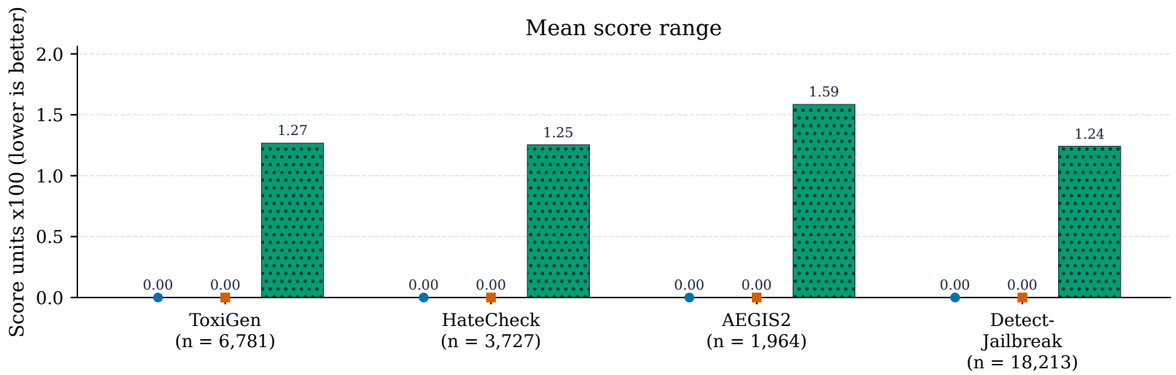
Laya | request: typed-decisions
serving: laya-rl-agent
routed: typed-decisions

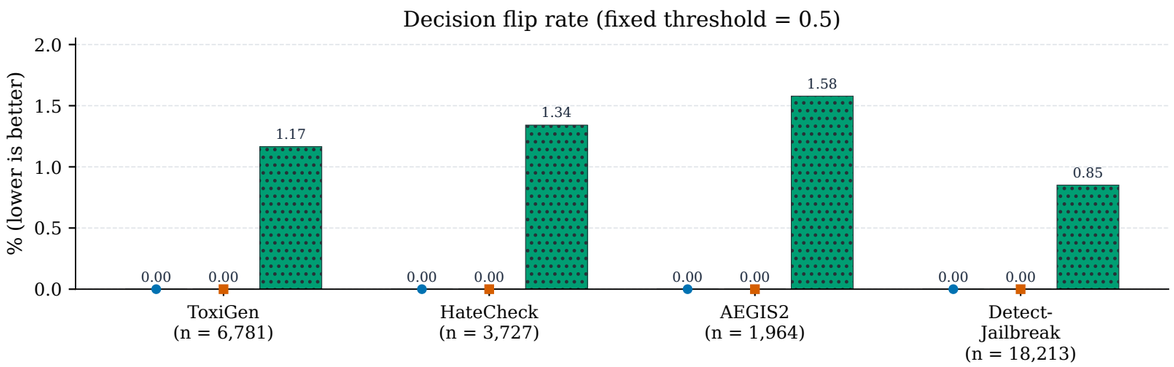
TypeSafe | request: jev-latest
resolved: unconfirmed (API response required)

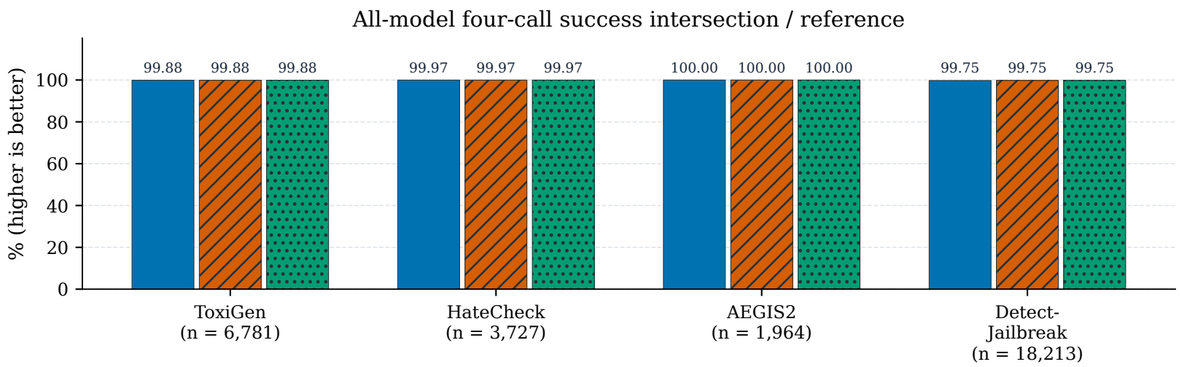
Repeat-call consistency: prompt v1

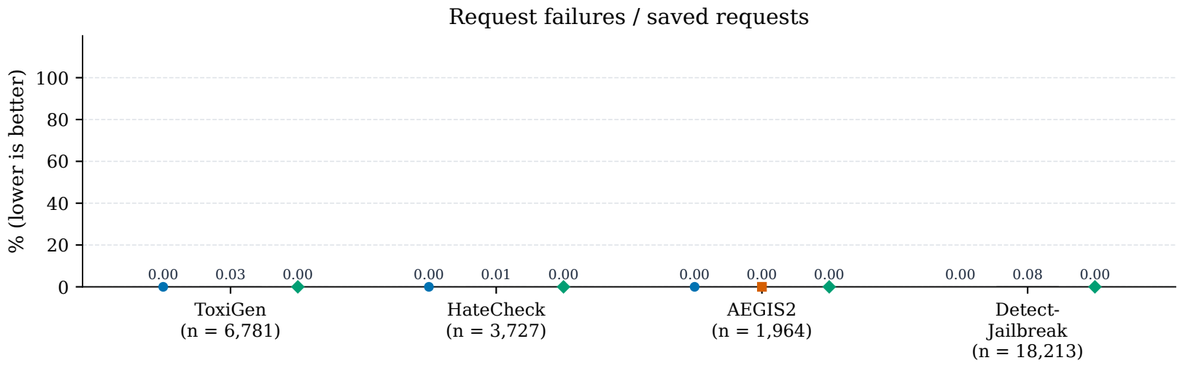
4 shared datasets | four calls per model | identical successful sample IDs across all models



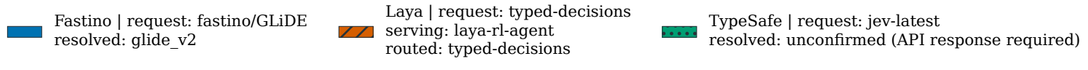






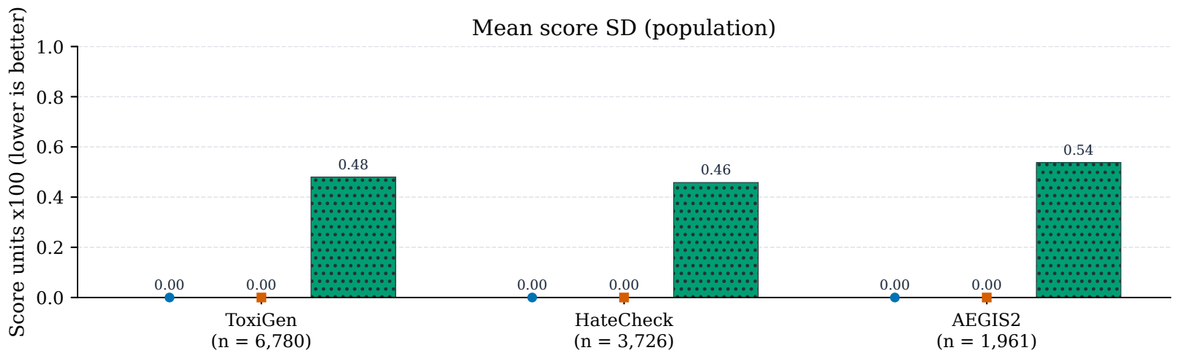


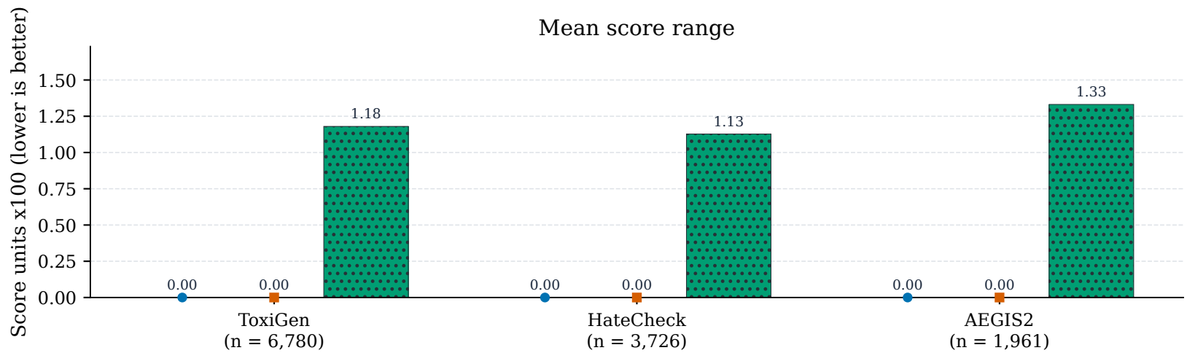
Score SD and range are displayed x100; flip rate uses score >= 0.5 for ALL four calls (not saved predictions).
Only datasets with complete four-call CSVs for every model and nonempty common successful IDs are shown.
Baseline markers show measured zeros, not missing data. Coverage uses all-model common IDs; failures use all saved requests.
Historical backend / routing / time changes may contribute. Consistency is not accuracy; this is exploratory.

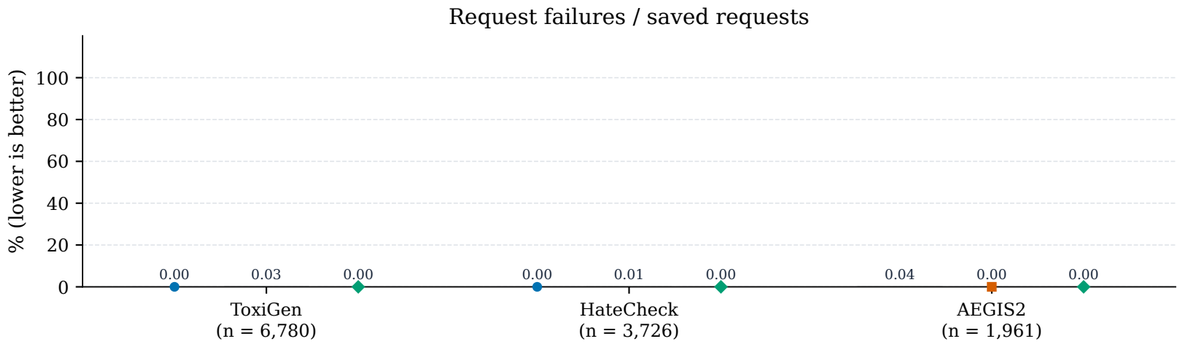


Repeat-call consistency: prompt v2

3 shared datasets | four calls per model | identical successful sample IDs across all models







Score SD and range are displayed x100; flip rate uses score >= 0.5 for ALL four calls (not saved predictions).
Only datasets with complete four-call CSVs for every model and nonempty common successful IDs are shown.
Baseline markers show measured zeros, not missing data. Coverage uses all-model common IDs; failures use all saved requests.
Historical backend / routing / time changes may contribute. Consistency is not accuracy; this is exploratory.

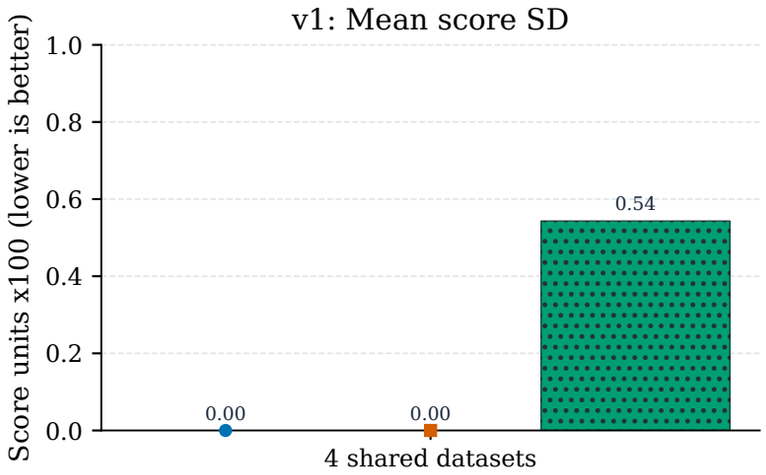
Fastino | request: fastino/GLiDE
resolved: glide_v2

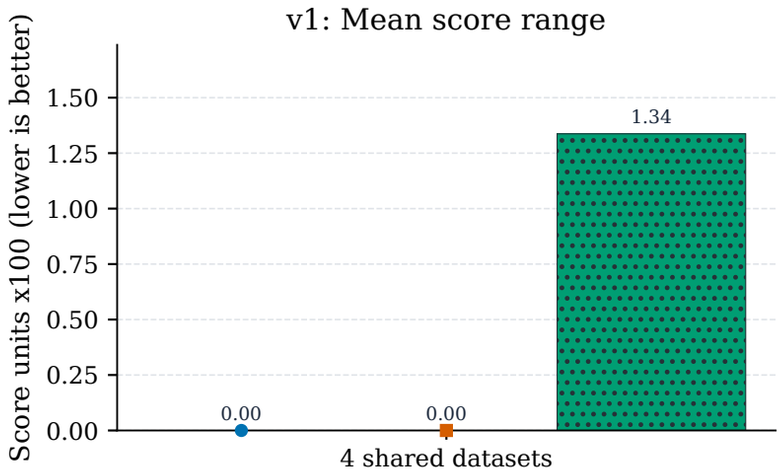
Laya | request: typed-decisions
serving: laya-rl-agent
routed: typed-decisions

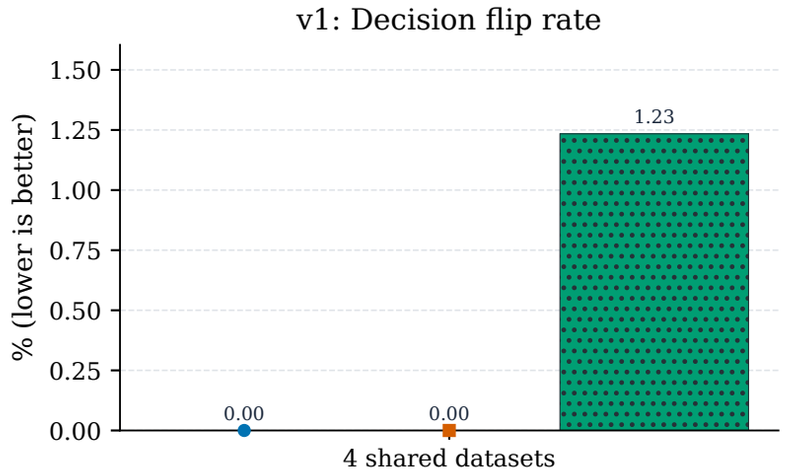
TypeSafe | request: jev-latest
resolved: unconfirmed (API response required)

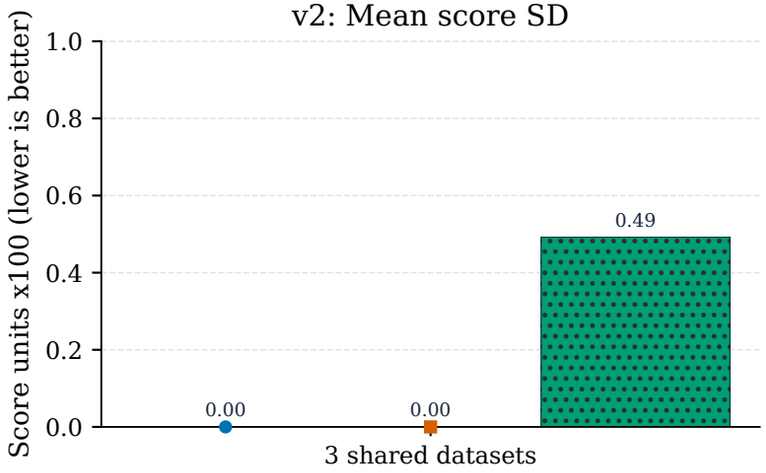
Model repeat-call consistency: paired dataset-macro averages

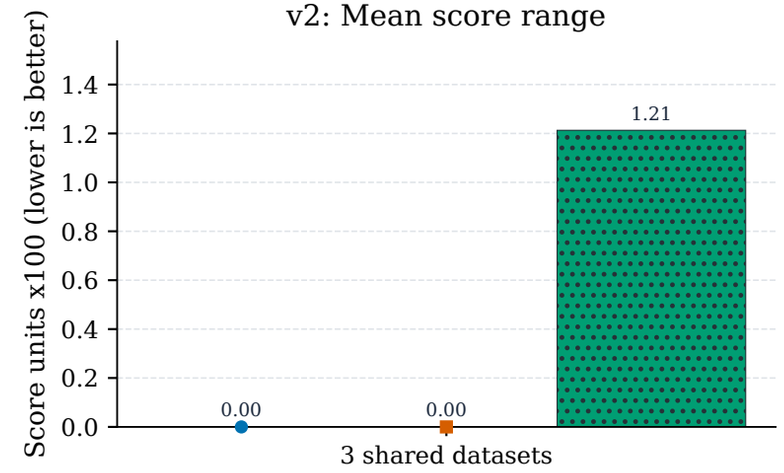
Prompt v1 and v2 are separate | common four-call successful IDs | shared non-PII datasets only

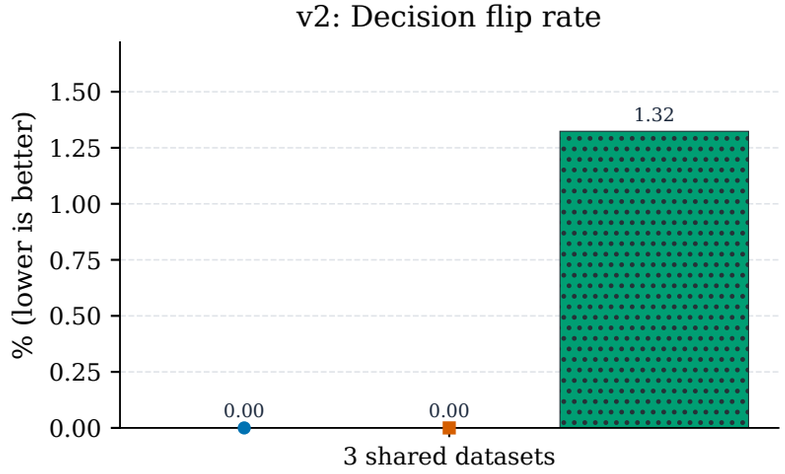












Equal dataset weights. Score SD / range displayed x100; decision flips use the fixed threshold 0.5.
Baseline markers show measured zeros. No eligible shared datasets = N/A; datasets and IDs may differ between prompts.
See the report for shared sample counts, coverage and failure rates; low variation does not imply correctness.

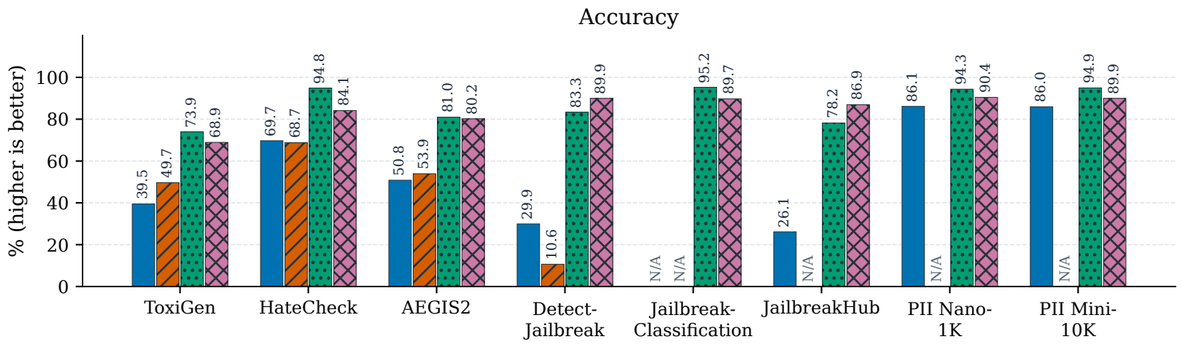
Fastino | request: fastino/GLiDE
resolved: glide_v2

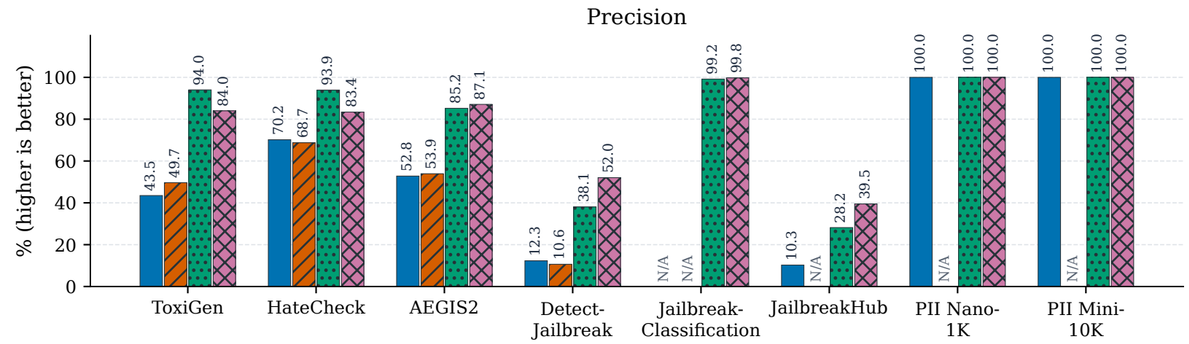
Laya | request: typed-decisions
serving: laya-rl-agent
routed: typed-decisions

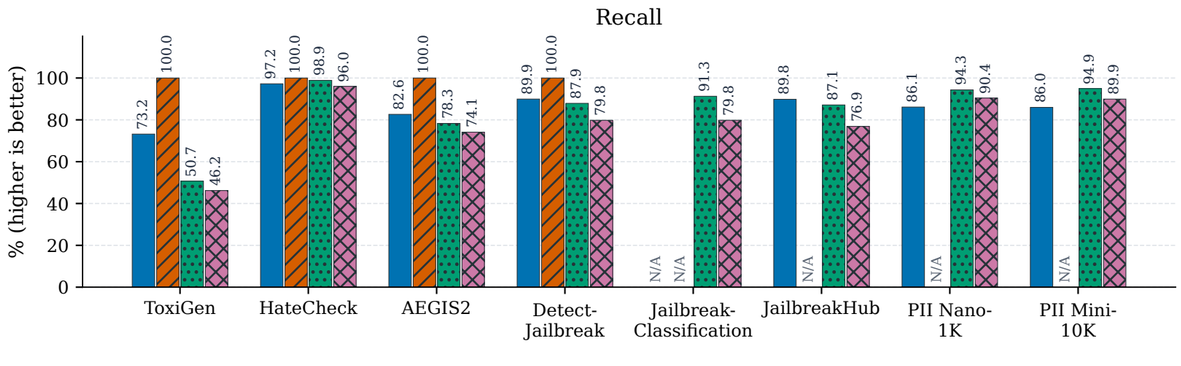
TypeSafe | request: jev-latest
resolved: unconfirmed (API response required)

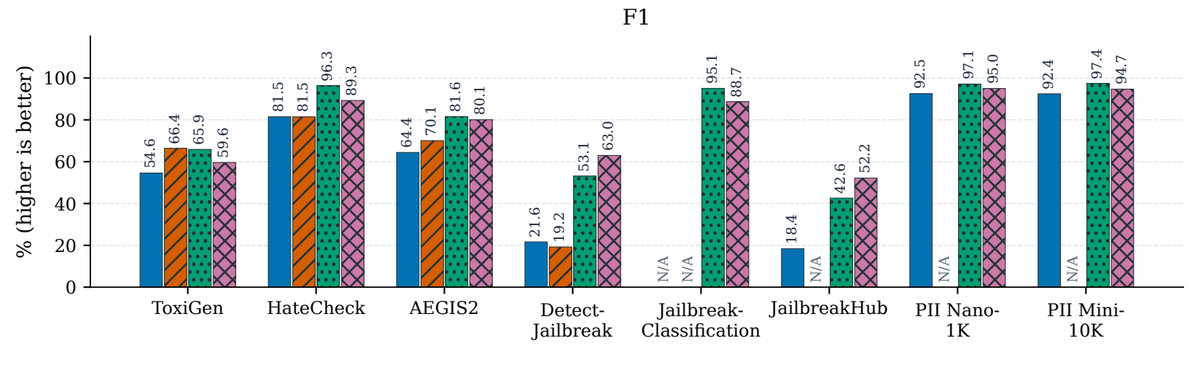
Four-model reference comparison: prompt v1

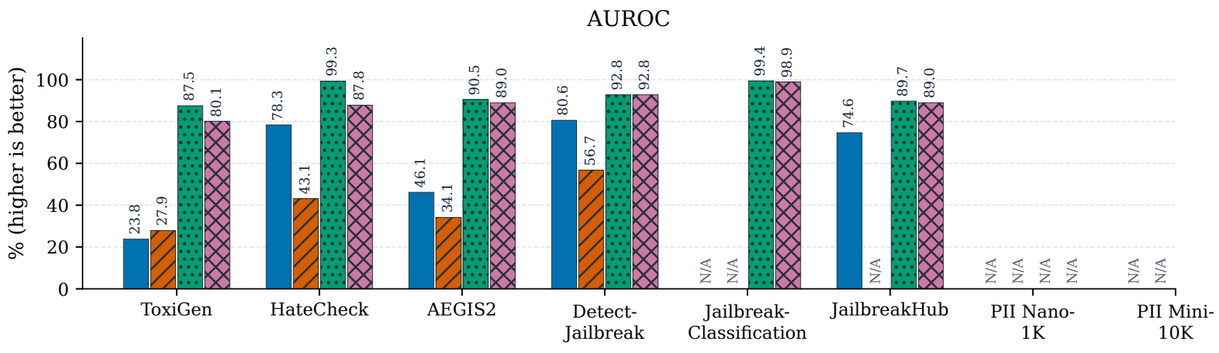
threshold = 0.5 | own successful samples | OpenJev: supplied aggregate, NOT sample-paired

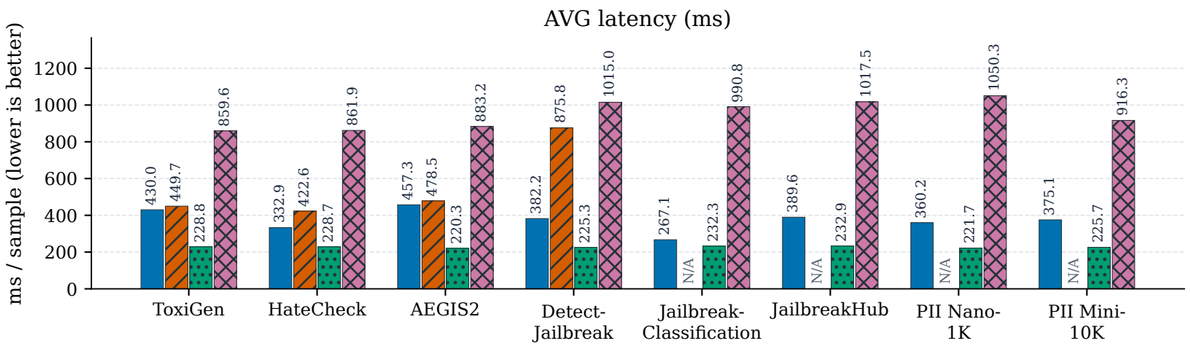












OpenJev values are transcribed from the supplied summary; valid/error counts and model version are unconfirmed.
Latency for each model = mean of elapsed / processed over four complete runs (including failed requests).
N/A = missing quality result, single-class AUROC, or incomplete four-run timing; not zero.
Descriptive only, not a fair ranking. PII values measure presence detection, not masking quality.

Fastino | request: fastino/GLiDE
resolved: glide_v2

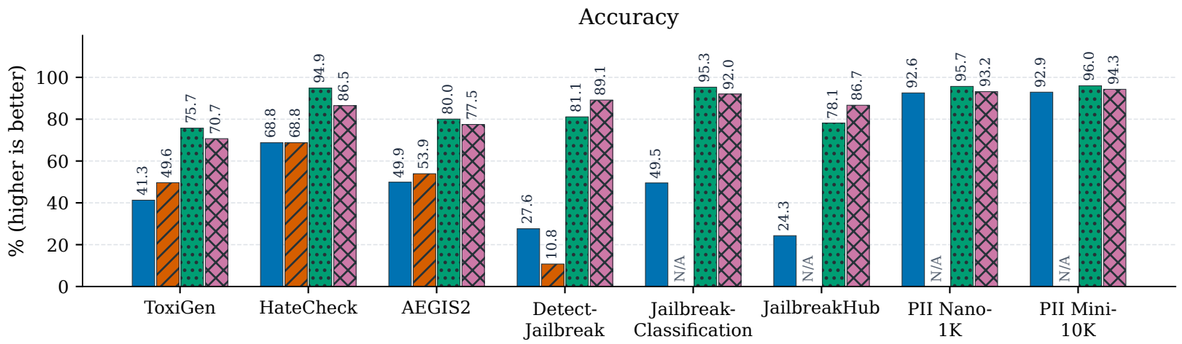
Laya | request: typed-decisions
serving: laya-rl-agent
routed: typed-decisions

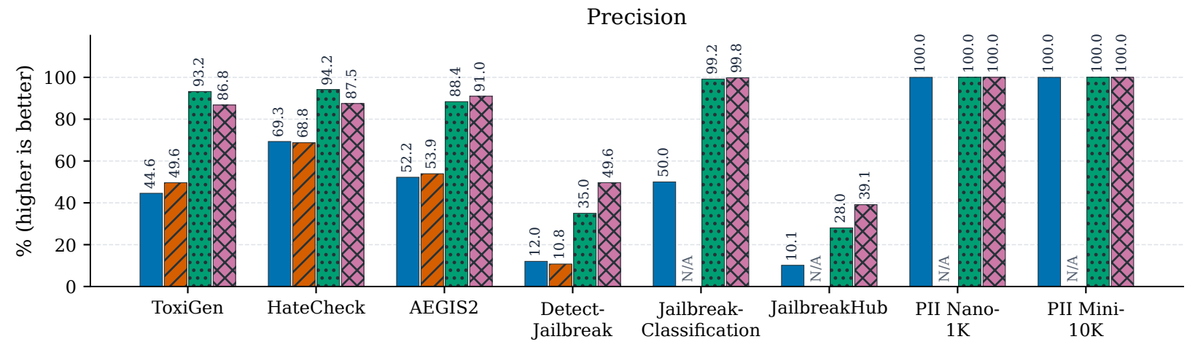
TypeSafe | request: jev-latest
resolved: unconfirmed (API response required)

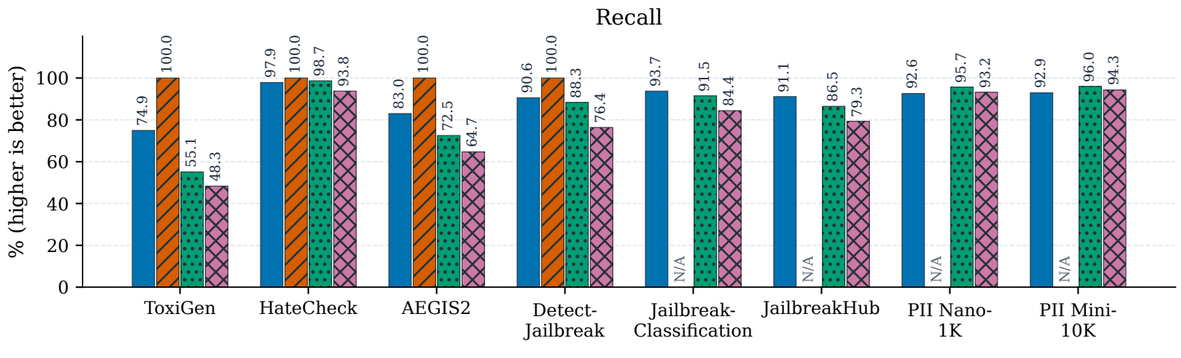
OpenJev | supplied internal summary
model / version: not supplied (not sample-paired)

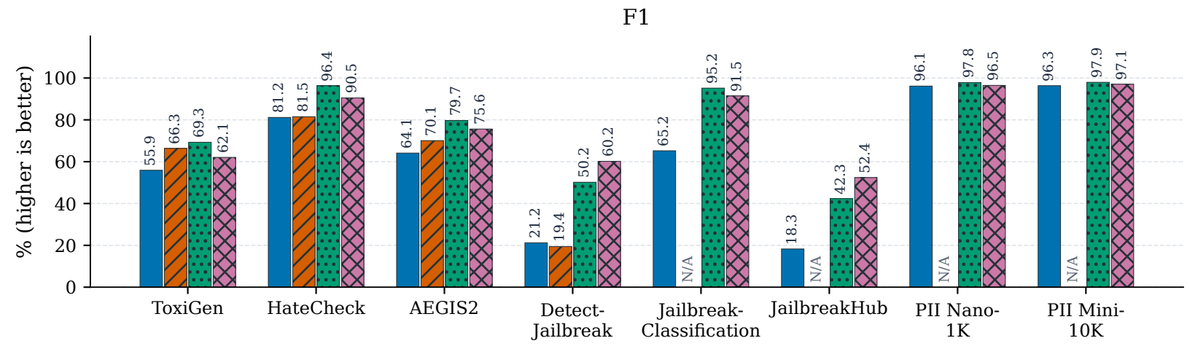
Four-model reference comparison: prompt v2

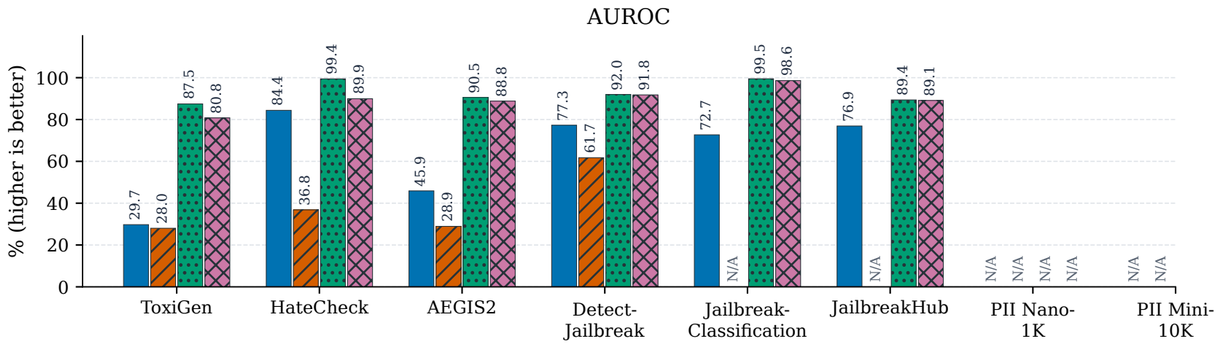
threshold = 0.5 | own successful samples | OpenJev: supplied aggregate, NOT sample-paired

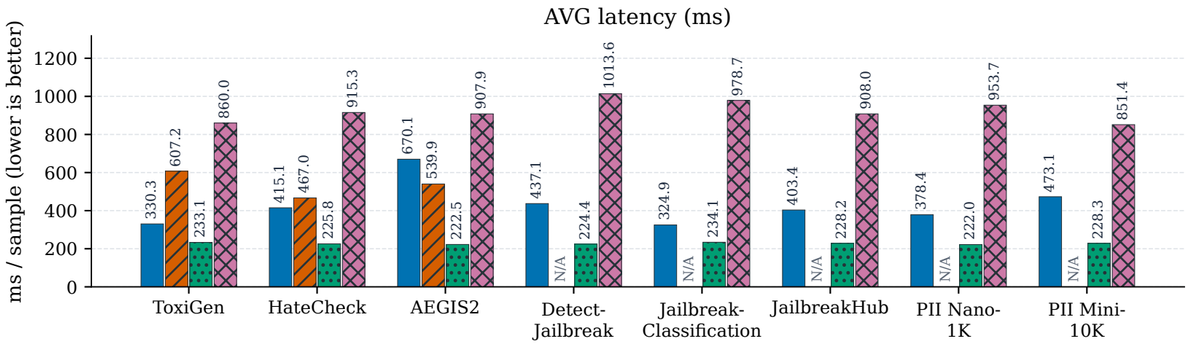












OpenJev values are transcribed from the supplied summary; valid/error counts and model version are unconfirmed.
Latency for each model = mean of elapsed / processed over four complete runs (including failed requests).
N/A = missing quality result, single-class AUROC, or incomplete four-run timing; not zero.
Descriptive only, not a fair ranking. PII values measure presence detection, not masking quality.

Fastino | request: fastino/GLiDE
resolved: glide_v2

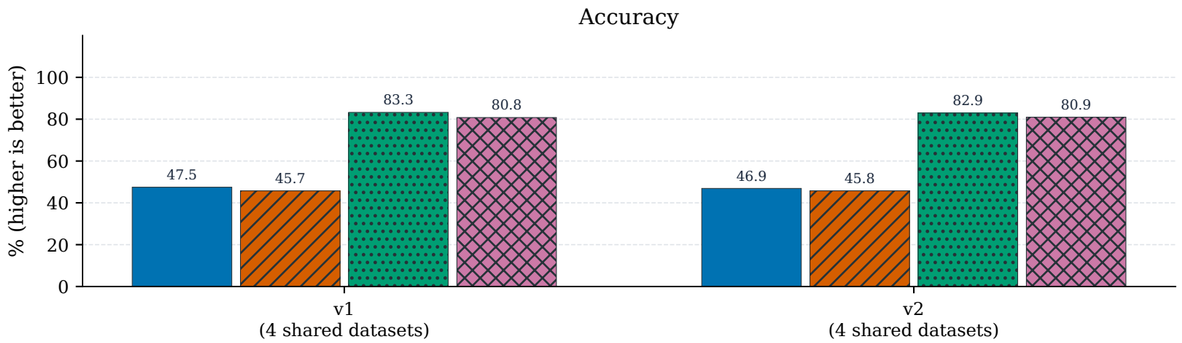
Laya | request: typed-decisions
serving: laya-rl-agent
routed: typed-decisions

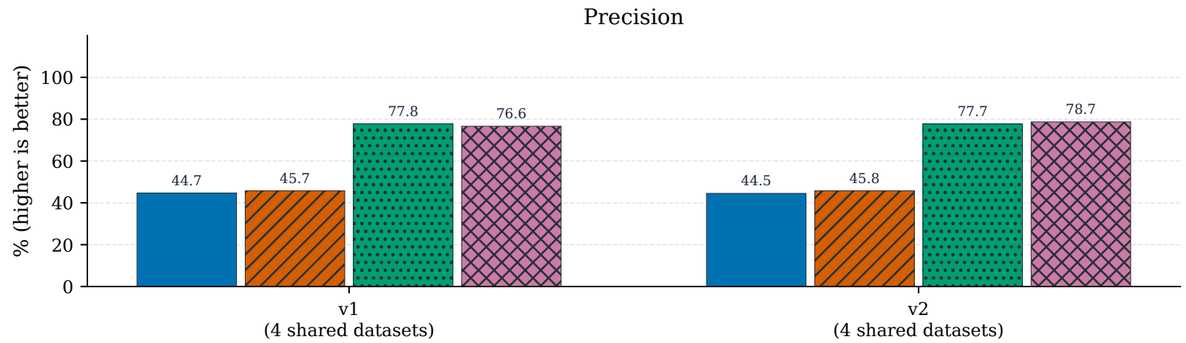
TypeSafe | request: jev-latest
resolved: unconfirmed (API response required)

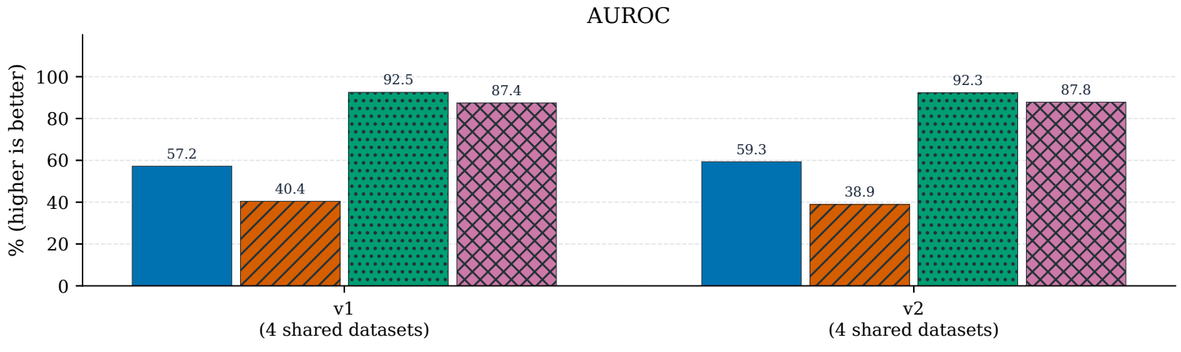
OpenJev | supplied internal summary
model / version: not supplied (not sample-paired)

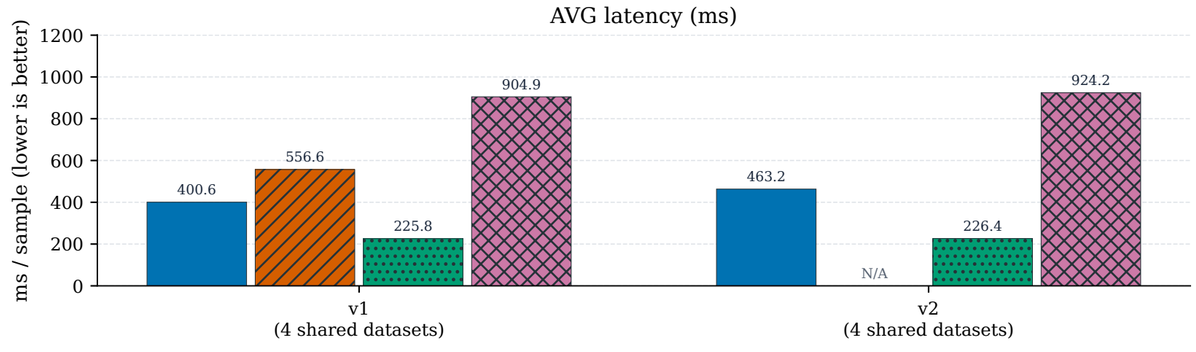
Four-model reference averages (not sample-paired)

threshold = 0.5 | shared non-PII dataset names | equal dataset weights | prompts kept separate









Each model uses its own successful samples; no common OpenJev sample IDs are available. Not a fair ranking.
Latency = mean per-run elapsed / processed across four complete runs, then an equal-weight dataset mean.
Missing any selected dataset's four-run timing makes average latency N/A; failure waiting is included.
OpenJev summary precision is retained as supplied; raw predictions and repeat-call stability are unavailable.

Fastino | request: fastino/GLiDE
resolved: glide_v2

Laya | request: typed-decisions
serving: laya-rl-agent
routed: typed-decisions

TypeSafe | request: jev-latest
resolved: unconfirmed (API response required)

OpenJev | supplied internal summary
model / version: not supplied (not sample-paired)