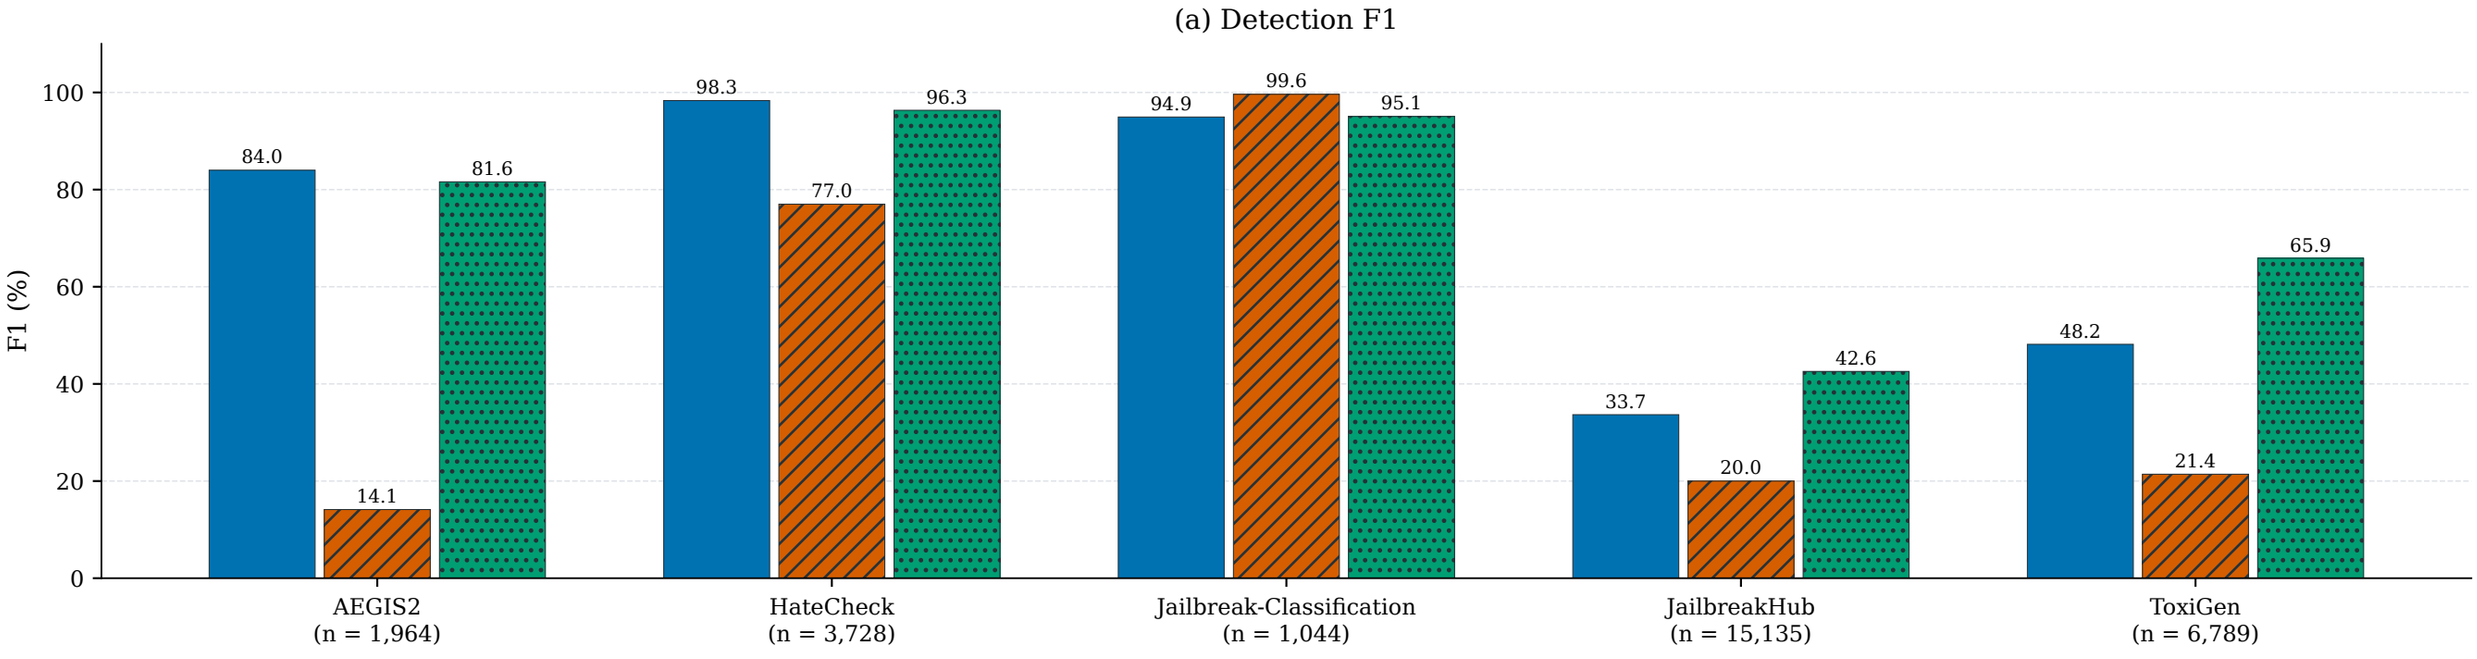
Sample-Aligned Classification Comparison

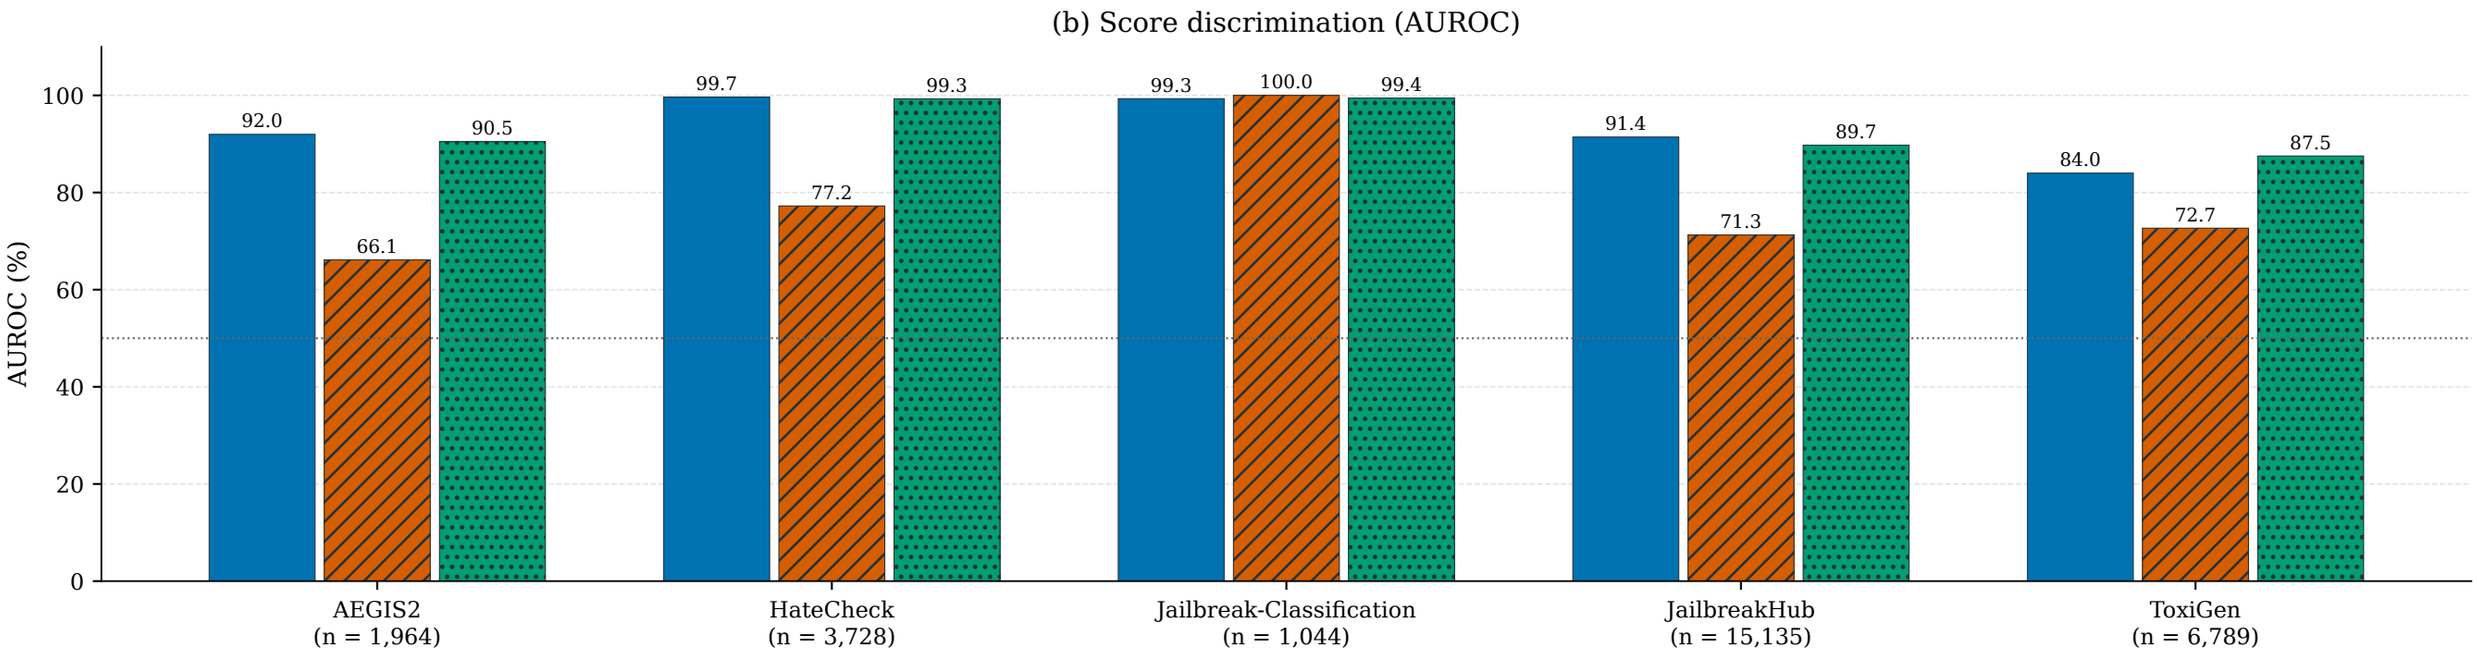
v1 | threshold = 0.5 | common completed non-PII datasets | higher is better

Fastino | request: fastino/GLiDE
resolved: glide_v2

Laya | request: typed-decisions
serving: laya-rl-agent
routed: typed-decisions

TypeSafe | request: jev-latest
resolved: unconfirmed (API response required)

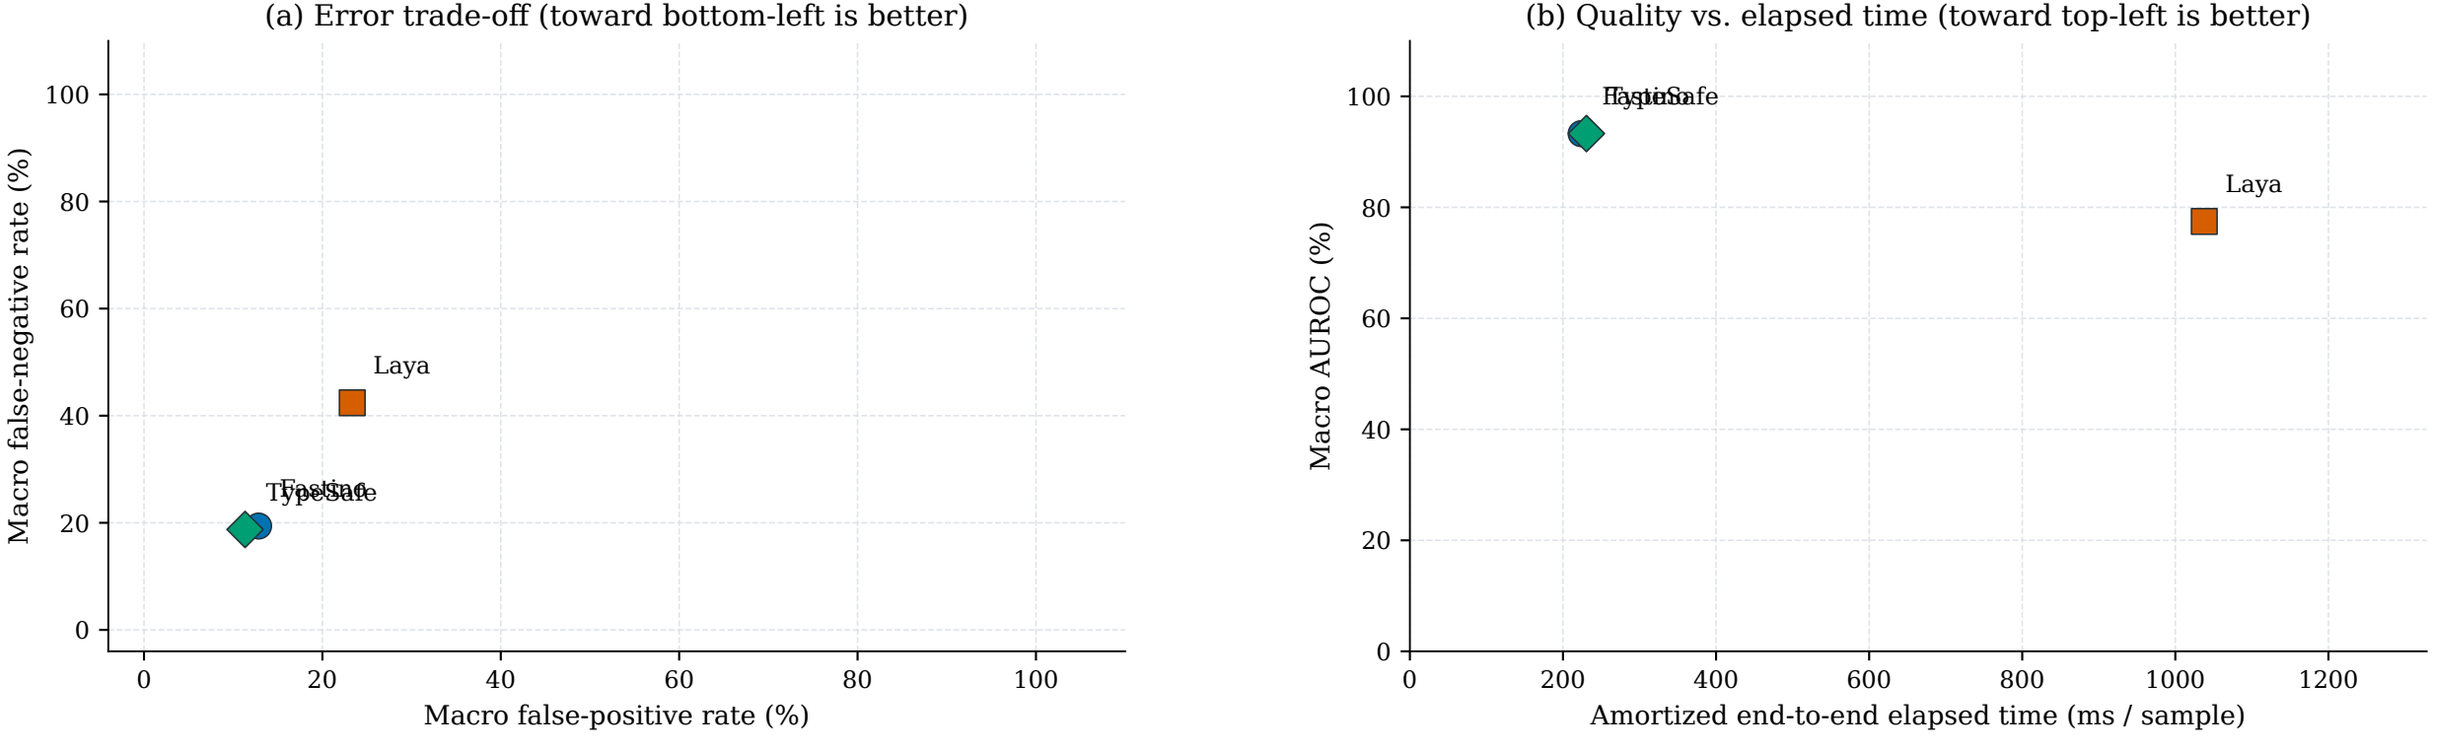




Each dataset uses the same successful sample IDs for all selected models.
Values are descriptive point estimates; repeated sweeps are not independent replicates.

Error Rates and Latency Trade-offs

v1 | threshold = 0.5 | (a) 5 shared datasets | (b) 5 shared fully timed datasets



Panel (b) recomputes BOTH macro AUROC and sample-weighted elapsed on the same shared timed subset:
AEGIS2, HateCheck, Jailbreak-Classification, JailbreakHub, ToxiGen. No four-run timing averages substitute for missing v1 / 0.5 timing.
Elapsed time = total logged elapsed / all processed subset rows, including failed requests.
Runs occurred at different times; this is not pure inference latency or a request-level percentile.

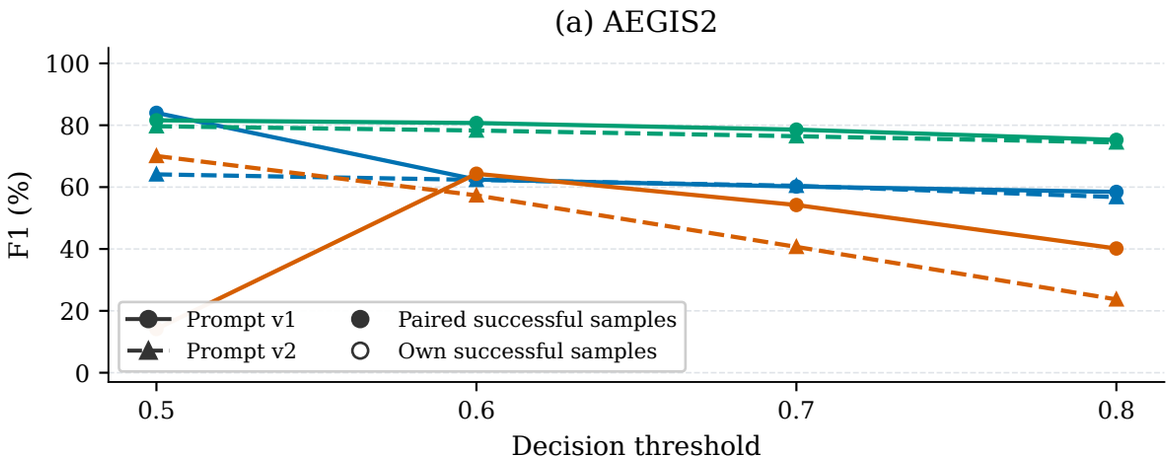
Fastino | request: fastino/GLiDE
resolved: glide_v2

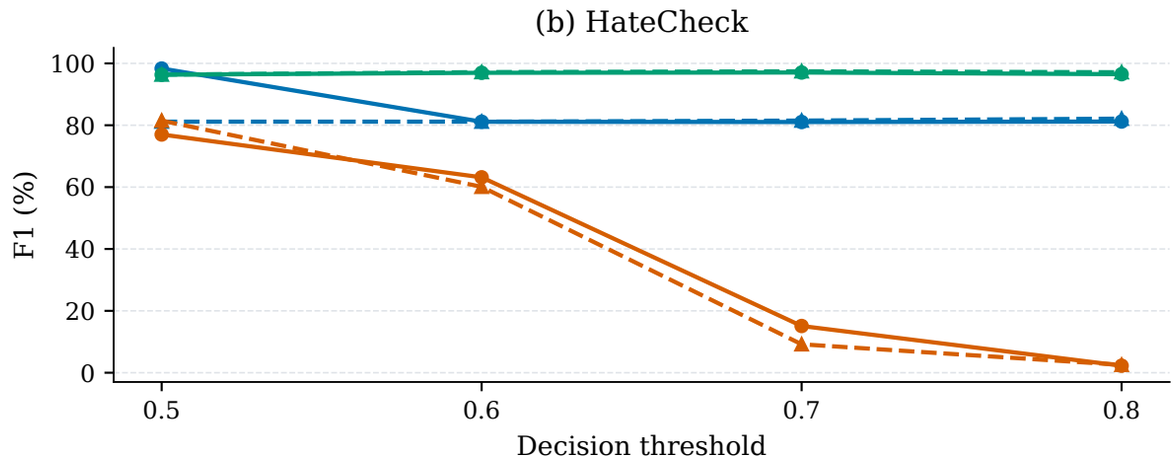
Laya | request: typed-decisions
serving: laya-rl-agent
routed: typed-decisions

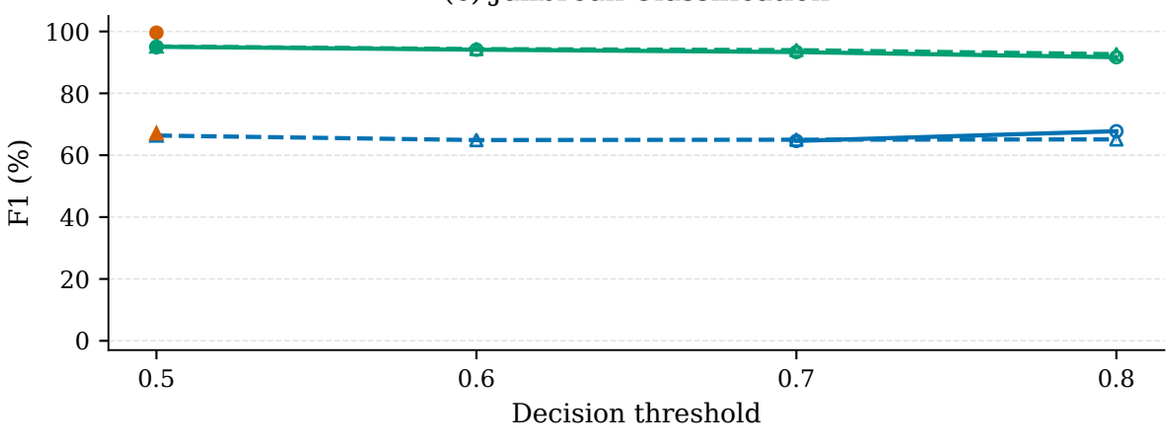
TypeSafe | request: jev-latest
resolved: unconfirmed (API response required)

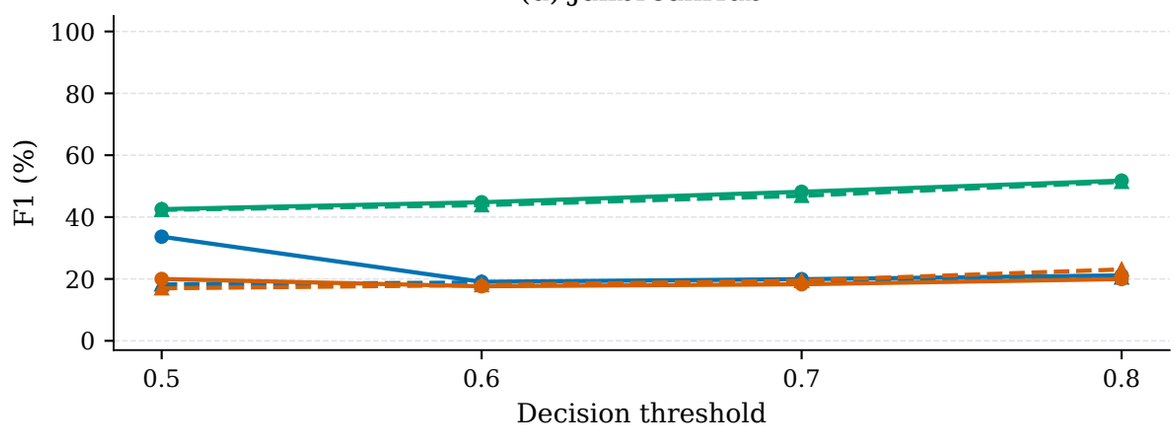
Prompt and Threshold Sensitivity (Fastino, Laya, TypeSafe)

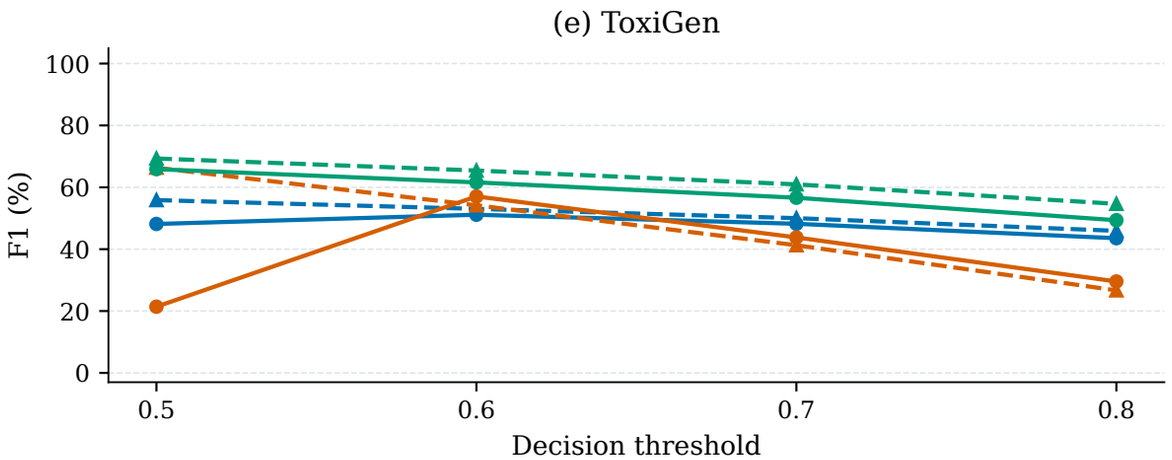
Exploratory test-set sweeps, not held-out model selection











Filled markers: paired successful samples; hollow: own successes when paired metrics are unavailable.
Sample sets can differ across points; this is not a fair ranking. Missing or all-failed runs remain gaps.
Curves mix threshold effects and repeat-call variation.
Scores are not held fixed across thresholds; lines do not imply independent replicates or confidence intervals.

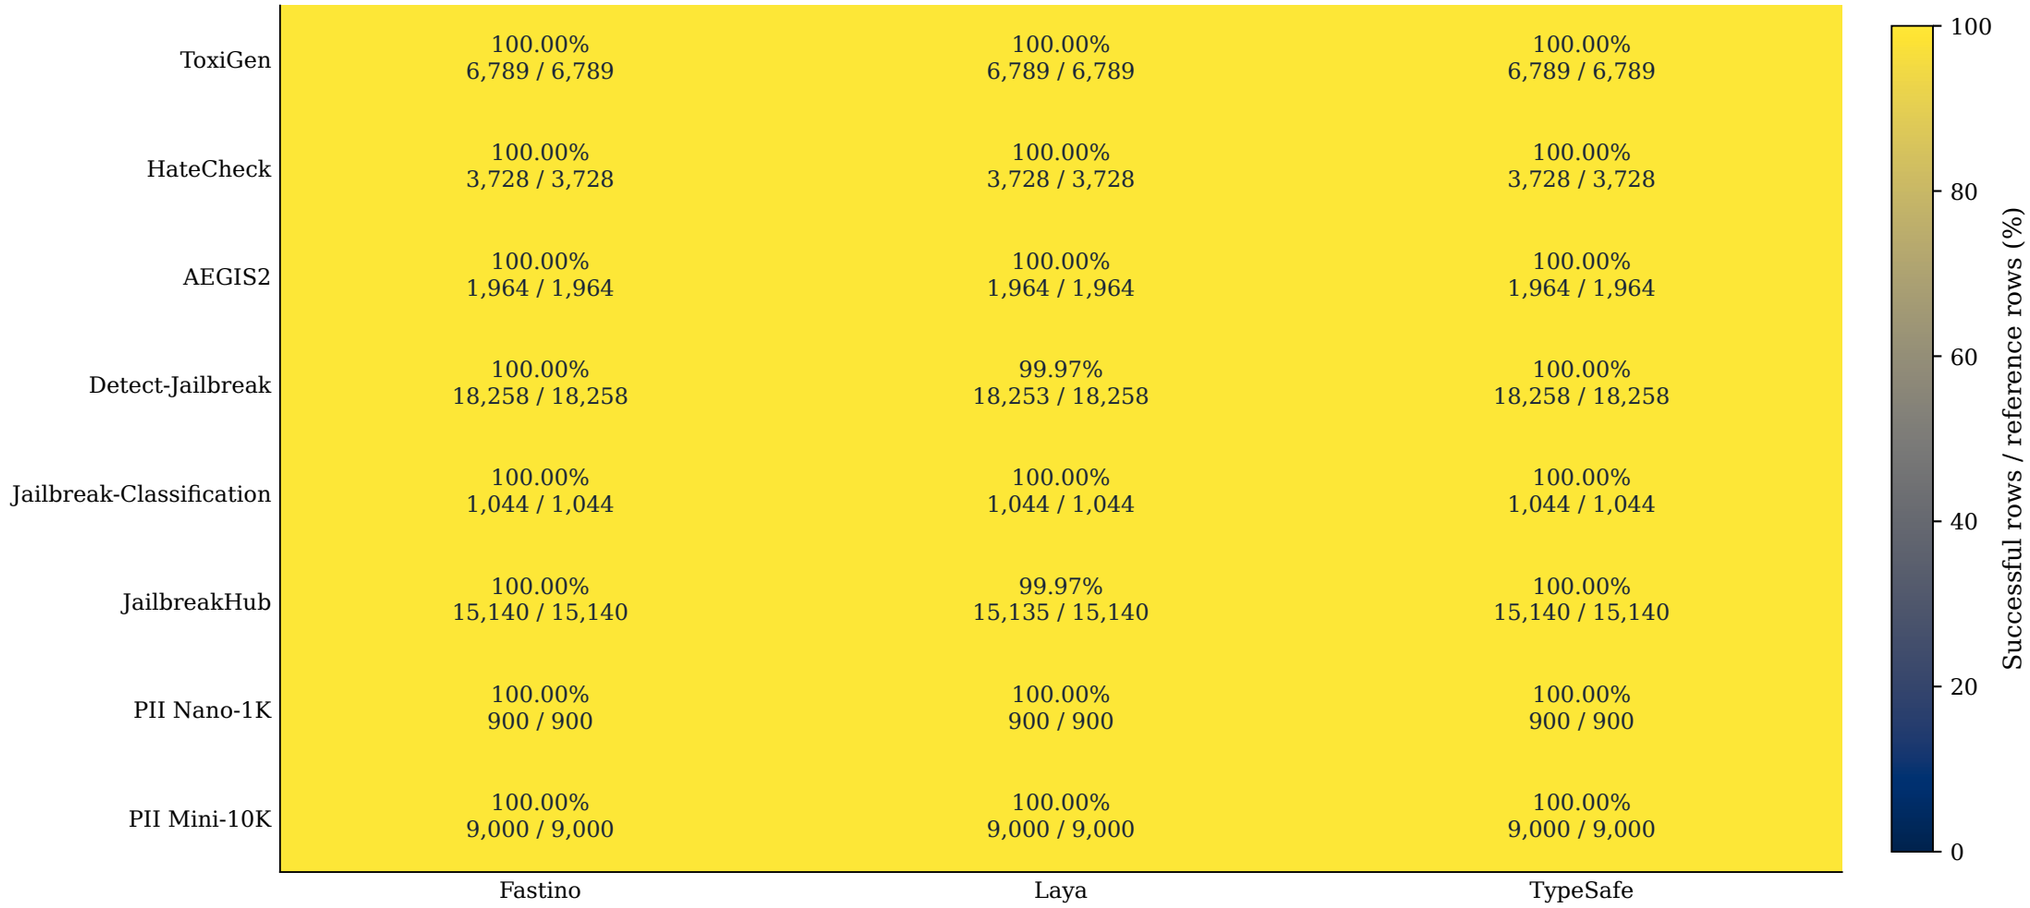
Fastino | request: fastino/GLiDE
resolved: glide_v2

Laya | request: typed-decisions
serving: laya-rl-agent
routed: typed-decisions

TypeSafe | request: jev-latest
resolved: unconfirmed (API response required)

Evaluation Coverage and Request Failures

v1 | threshold = 0.5 | counts are successful / reference rows



Fastino | request: fastino/GLiDE
resolved: glide_v2

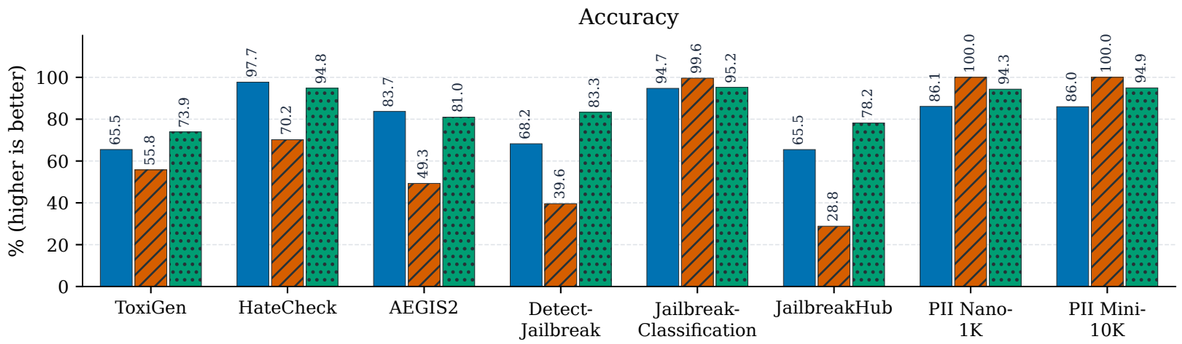
Laya | request: typed-decisions
serving: laya-rl-agent
routed: typed-decisions

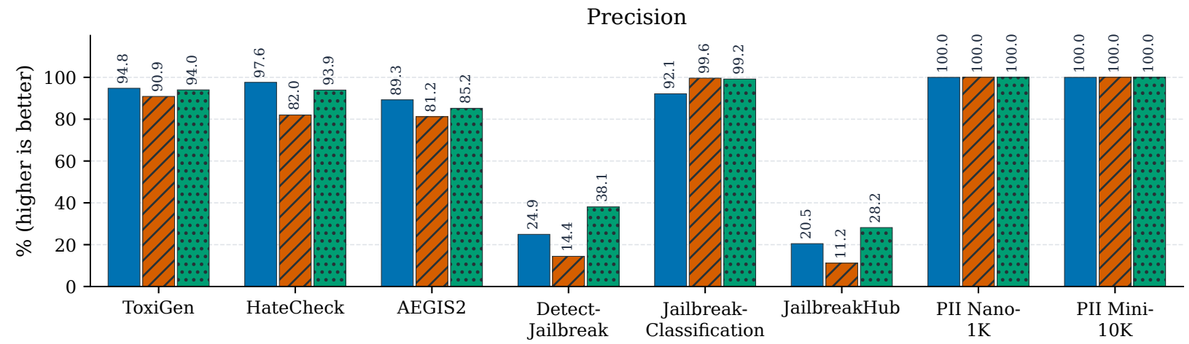
TypeSafe | request: jev-latest
resolved: unconfirmed (API response required)

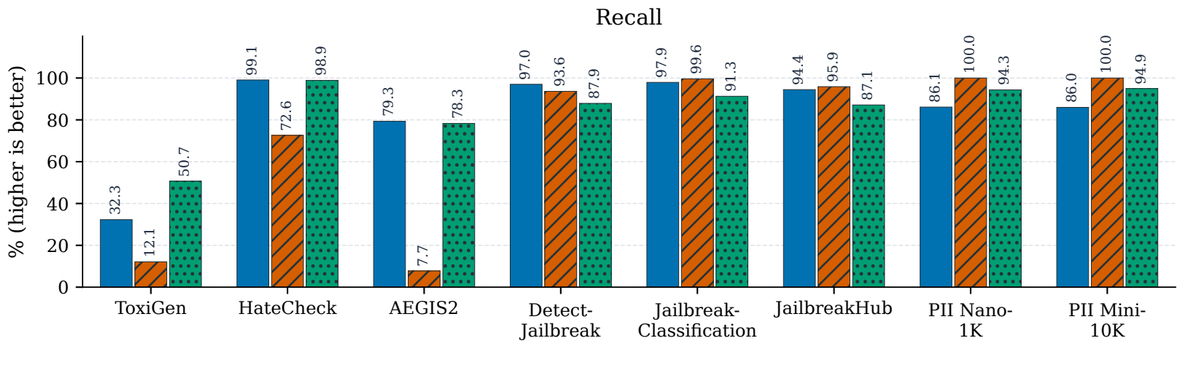
Gray = no saved predictions; hatched = complete CSV with no successful requests, not a missing run.
PII is ordinary binary classification (PII present vs absent); rows show coverage, not masking quality.

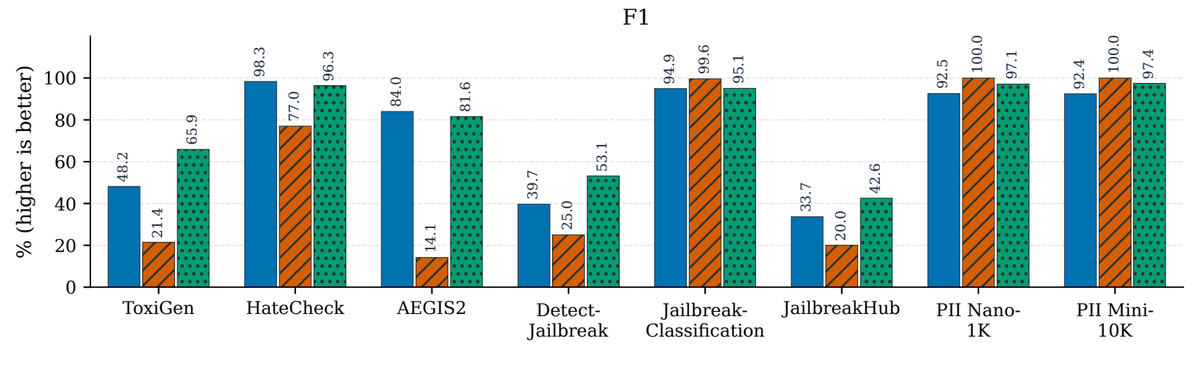
Per-Dataset Performance Metrics (Fastino, Laya, TypeSafe)

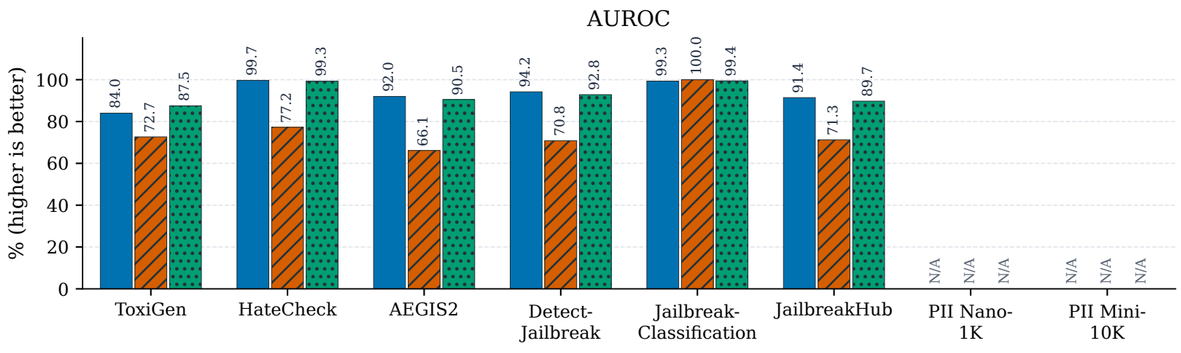
v1 | threshold = 0.5 | each model's own successful samples (not paired)

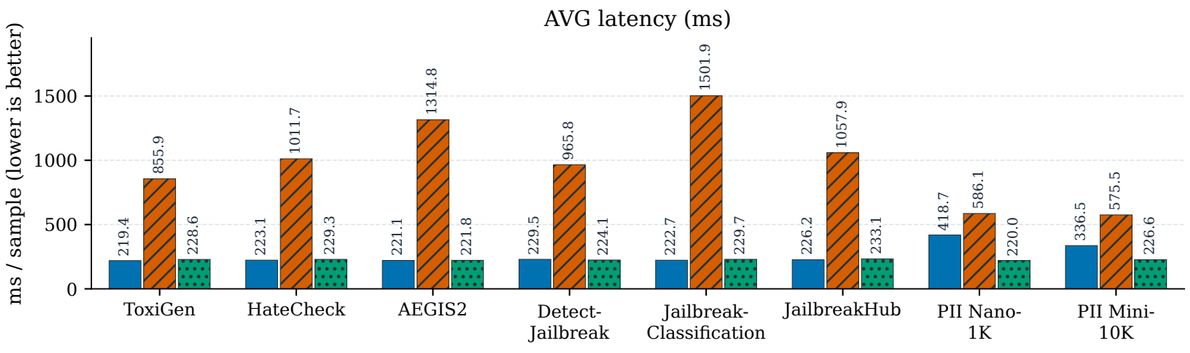












N/A = missing result, no successful predictions, undefined metric, or no aligned complete timing log.
PII is ordinary binary classification (PII present vs absent), not extraction or masking quality.
PII datasets contain positives only; AUROC is undefined.
Latency = logged end-to-end elapsed / all processed rows, including failed requests; not pure inference time.
Successful sample sets can differ between models; use the paired baseline for fair quality comparisons.

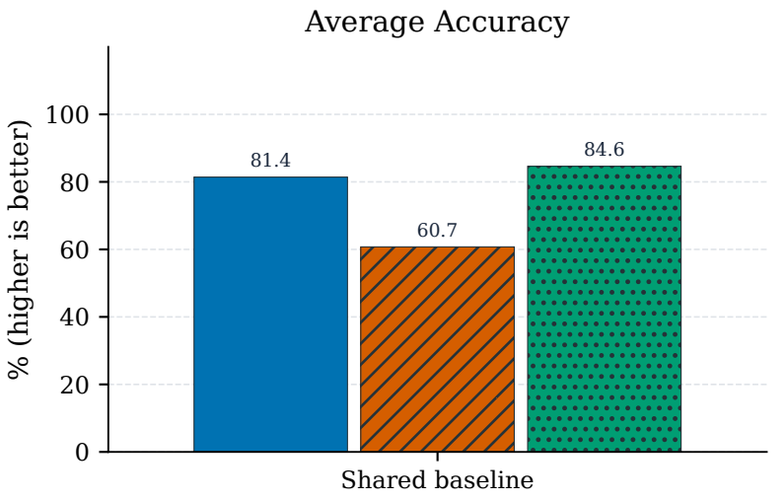
Fastino | request: fastino/GLiDE
resolved: glide_v2

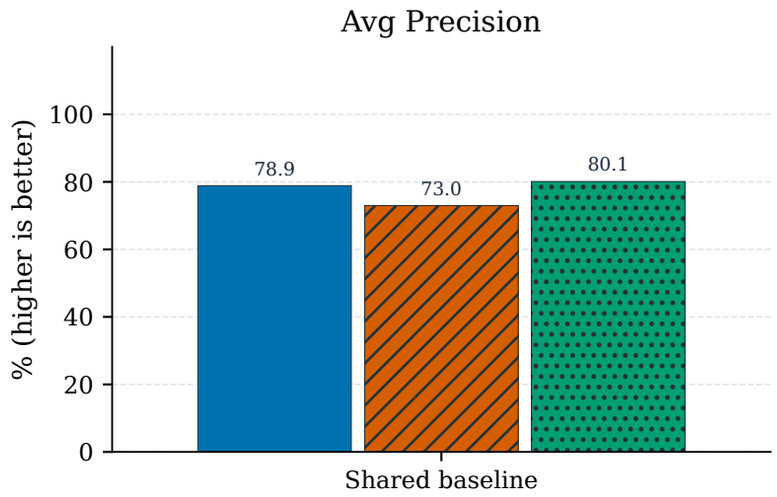
Laya | request: typed-decisions
serving: laya-rl-agent
routed: typed-decisions

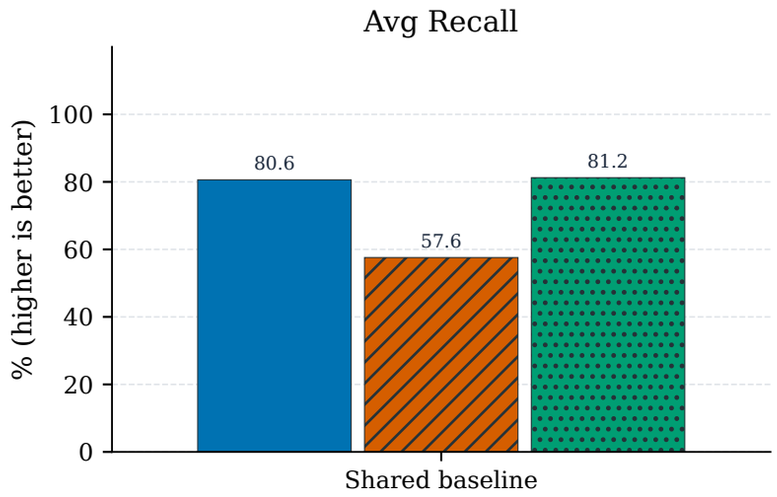
TypeSafe | request: jev-latest
resolved: unconfirmed (API response required)

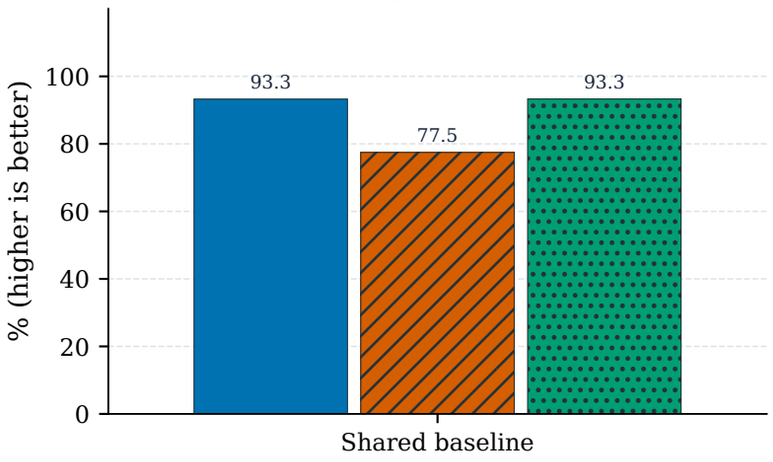
Sample-Aligned Model Averages (Fastino, Laya, TypeSafe)

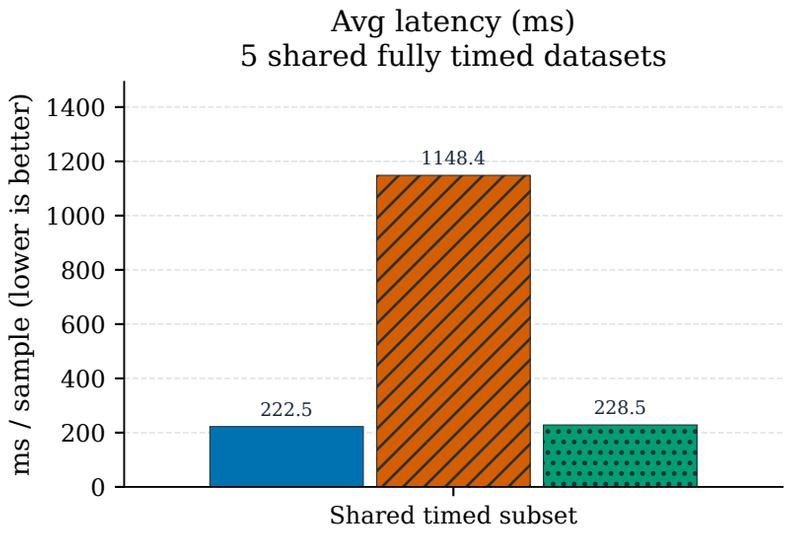
v1 | threshold = 0.5 | quality: 5 shared non-PII datasets | latency: 5 shared fully timed datasets | equal dataset weights











Quality: common successful sample IDs within each dataset, then an unweighted mean across datasets.
Latency: unweighted mean of per-dataset log elapsed / processed rows on the SAME timed subset for all models: AEGIS2, HateCheck, Jailbreak-Classification, JailbreakHub, ToxiGen. Quality and latency dataset scopes can differ.
Elapsed includes failures and network overhead, not pure inference time. No four-run timing substitutes.
N/A timing means no shared fully timed baseline datasets; missing is never zero.

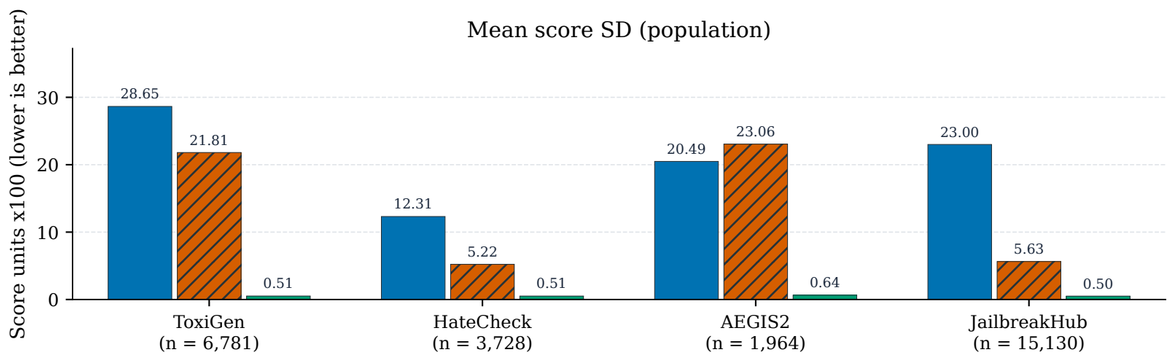
Fastino | request: fastino/GLiDE
resolved: glide_v2

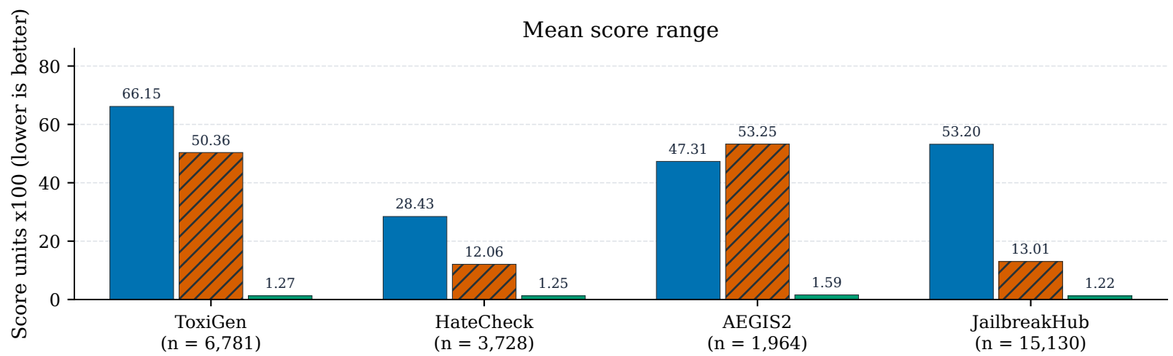
Laya | request: typed-decisions
serving: laya-rl-agent
routed: typed-decisions

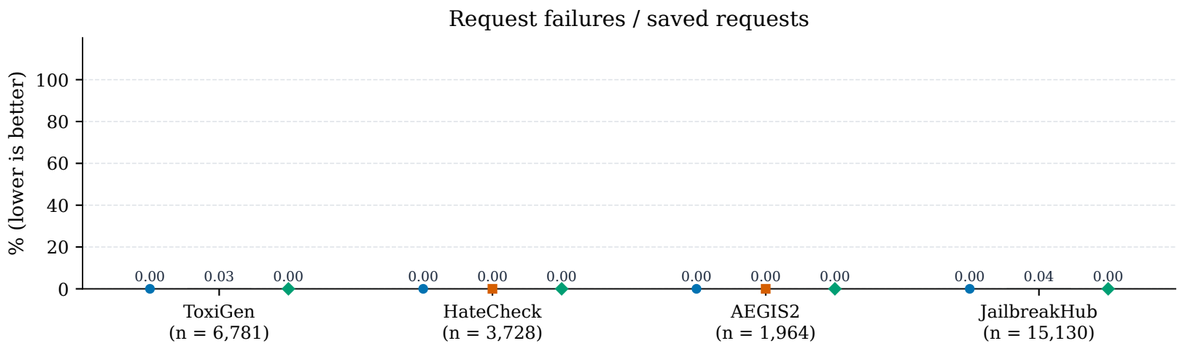
TypeSafe | request: jev-latest
resolved: unconfirmed (API response required)

Repeat-Call Consistency: v1 (Fastino, Laya, TypeSafe)

4 shared datasets | four calls per model | identical successful sample IDs across all models







Score SD and range are displayed x100; flip rate uses score >= 0.5 for ALL four calls (not saved predictions).
Only datasets with complete four-call CSVs for every model and nonempty common successful IDs are shown.
Baseline markers show measured zeros, not missing data. Coverage uses all-model common IDs; failures use all saved requests.
Historical backend / routing / time changes may contribute. Consistency is not accuracy; this is exploratory.

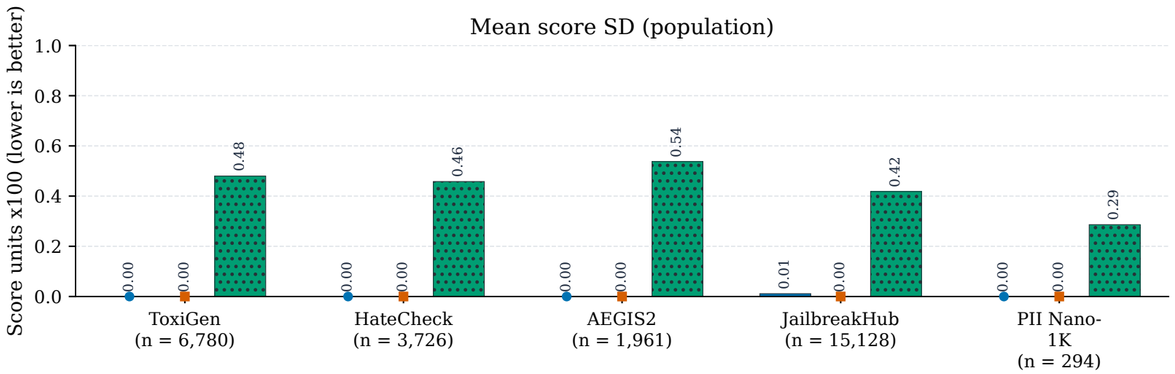
Fastino | request: fastino/GLiDE
resolved: glide_v2

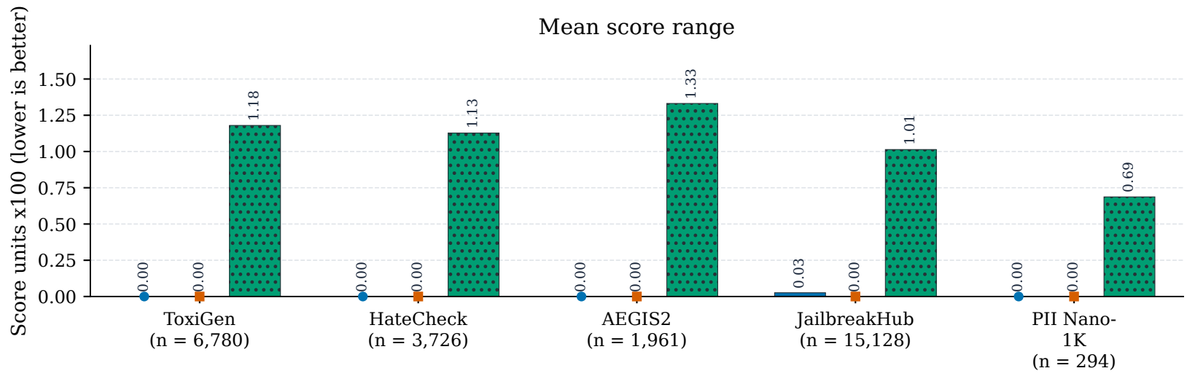
Laya | request: typed-decisions
serving: laya-rl-agent
routed: typed-decisions

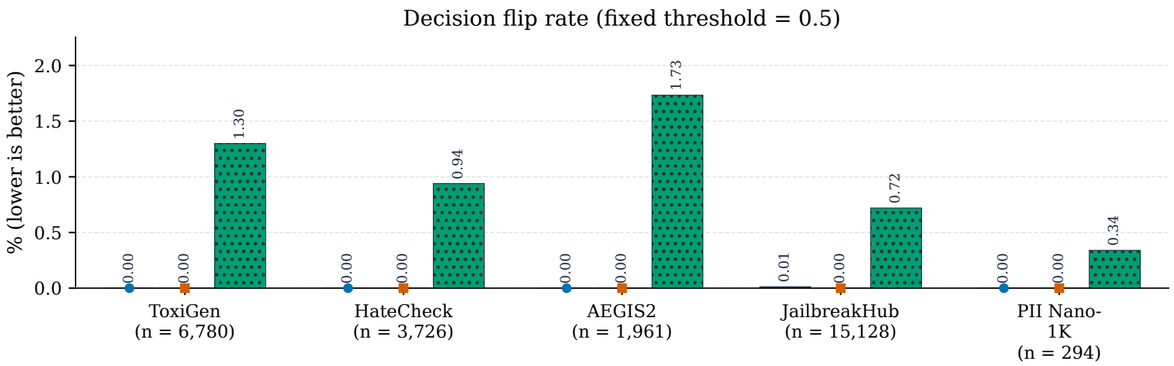
TypeSafe | request: jev-latest
resolved: unconfirmed (API response required)

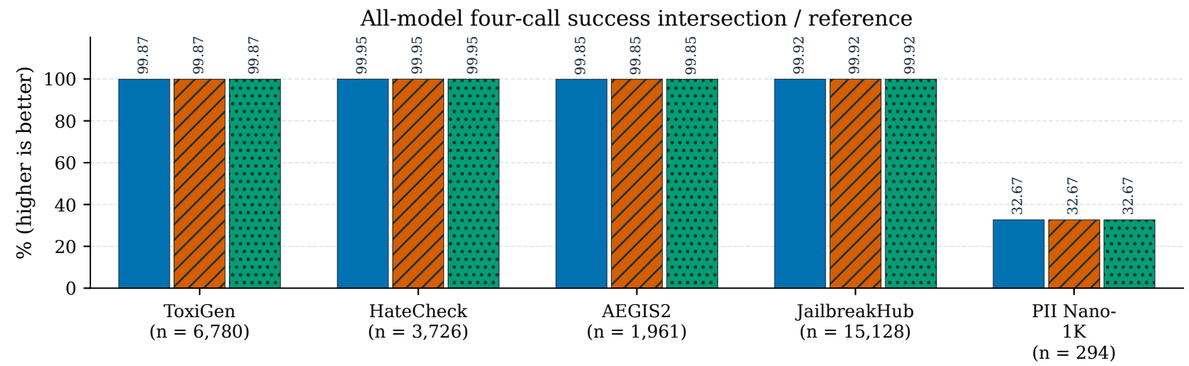
Repeat-Call Consistency: v2 (Fastino, Laya, TypeSafe)

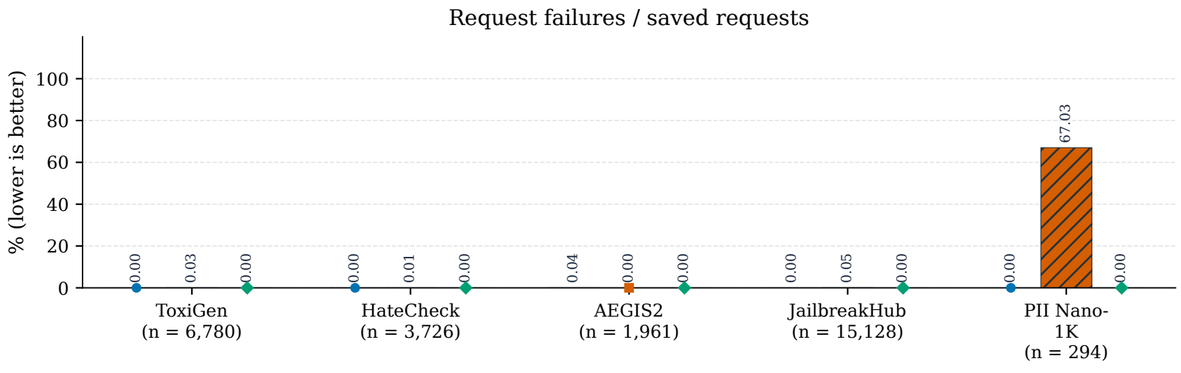
5 shared datasets | four calls per model | identical successful sample IDs across all models











Score SD and range are displayed x100; flip rate uses score >= 0.5 for ALL four calls (not saved predictions).
Only datasets with complete four-call CSVs for every model and nonempty common successful IDs are shown.
Baseline markers show measured zeros, not missing data. Coverage uses all-model common IDs; failures use all saved requests.
Historical backend / routing / time changes may contribute. Consistency is not accuracy; this is exploratory.

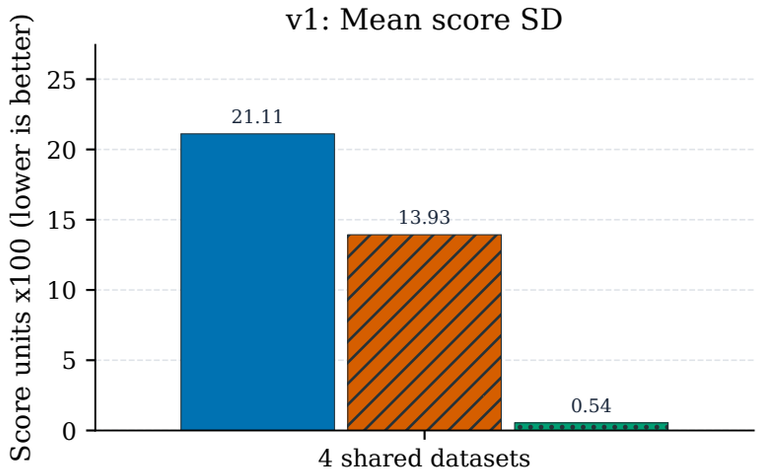
Fastino | request: fastino/GLiDE
resolved: glide_v2

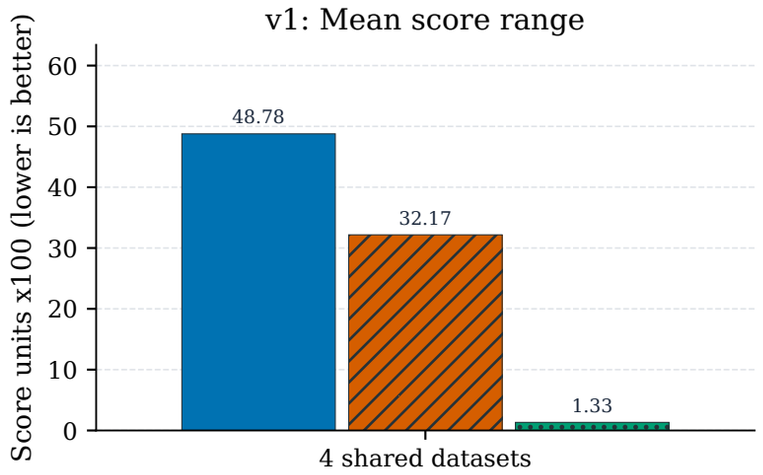
Laya | request: typed-decisions
serving: laya-rl-agent
routed: typed-decisions

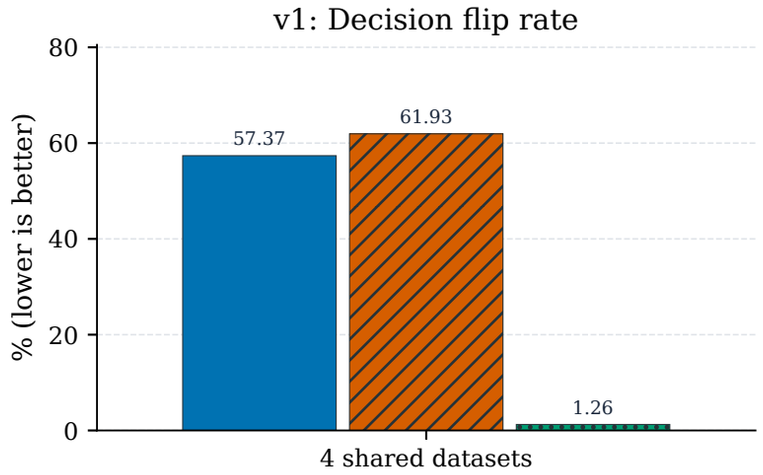
TypeSafe | request: jev-latest
resolved: unconfirmed (API response required)

Repeat-Call Consistency Averages (Fastino, Laya, TypeSafe)

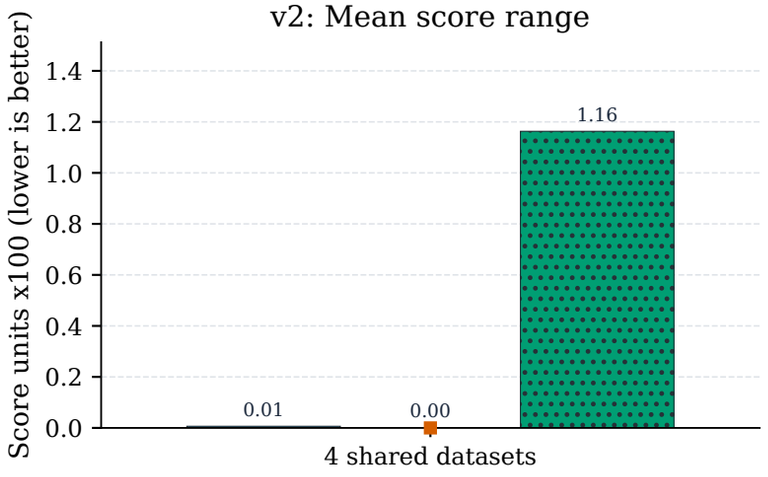
Prompt v1 and v2 are separate | common four-call successful IDs | shared non-PII datasets only

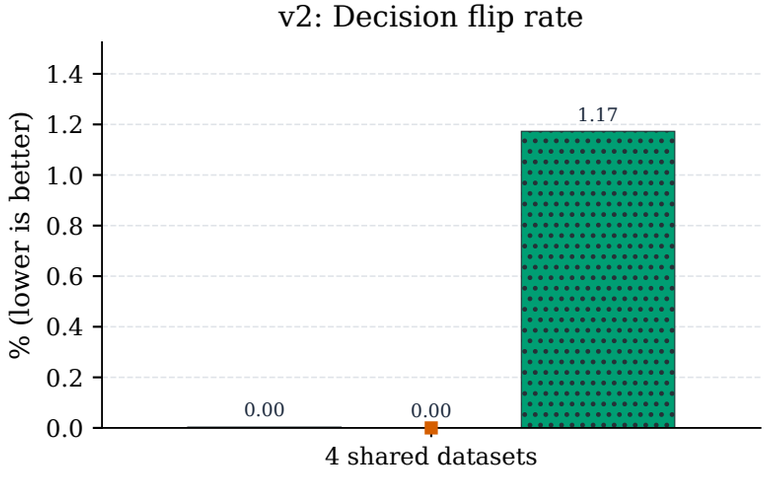












Equal dataset weights. Score SD / range displayed x100; decision flips use the fixed threshold 0.5.
Baseline markers show measured zeros. No eligible shared datasets = N/A; datasets and IDs may differ between prompts.
See the report for shared sample counts, coverage and failure rates; low variation does not imply correctness.

Fastino | request: fastino/GLiDE
resolved: glide_v2

Laya | request: typed-decisions
serving: laya-rl-agent
routed: typed-decisions

TypeSafe | request: jev-latest
resolved: unconfirmed (API response required)

Dataset-Category Multi-Model Reference Averages (v1 / Threshold 0.5) | Luna: native presets (prompt N/A) | Straiker: threshold N/A

v1 / threshold = 0.5 | equal dataset weights within each category | NOT sample-paired

Fastino | request: fastino/GLiDE
resolved: glide_v2

Fastino GLiNER2.5-Decide
model ID: fastino/gliner2.5-decide
response model: fastino/gliner2.5-decide
v1 / 0.5; PII not tested

Fastino GLiGuard-LLMGuardrails-300M
model ID: fastino/gliguard-LLMGuardrails-300M
v1 / 0.5; PII not tested

Laya | request: typed-decisions
serving: laya-rl-agent
routed: typed-decisions

TypeSafe | request: jev-latest
resolved: unconfirmed (API response required)

OpenJev 4B | supplied internal summary
model / version: not supplied (not sample-paired)

OpenJev 32B | supplied internal summary
single-run timing (not sample-paired)

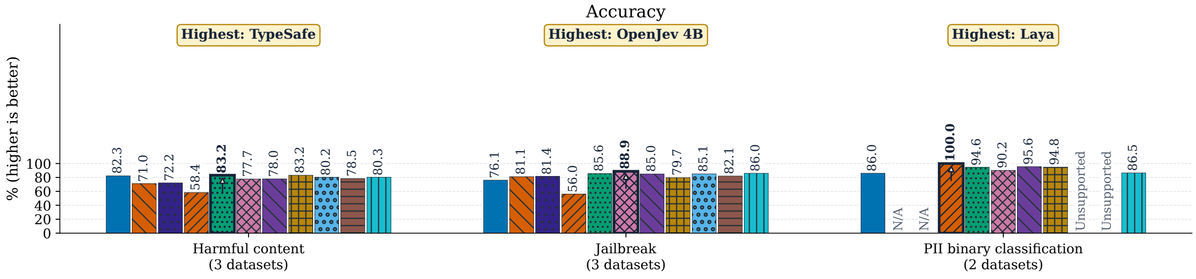
Microsoft Decision-1 | supplied internal summary
single-run timing (not sample-paired)

Qwen3Guard-Stream-8B | supplied summary
binary safety labels; PII unsupported

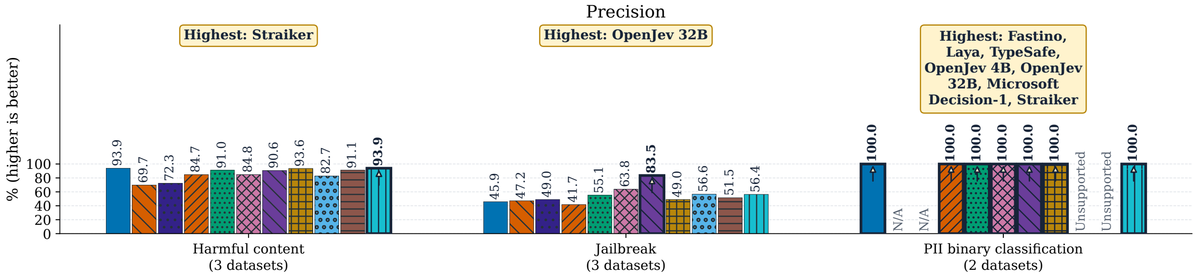
Luna | original task-specific presets
prompt: N/A; binary threshold: 0.5
PII: separate multilabel evaluation

Straiker | supplied aggregate
prompt: v1; threshold: N/A
single-run log elapsed/processed

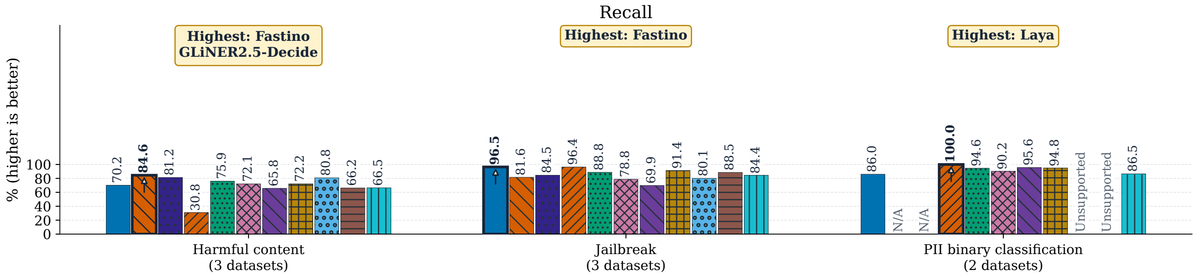
Accuracy



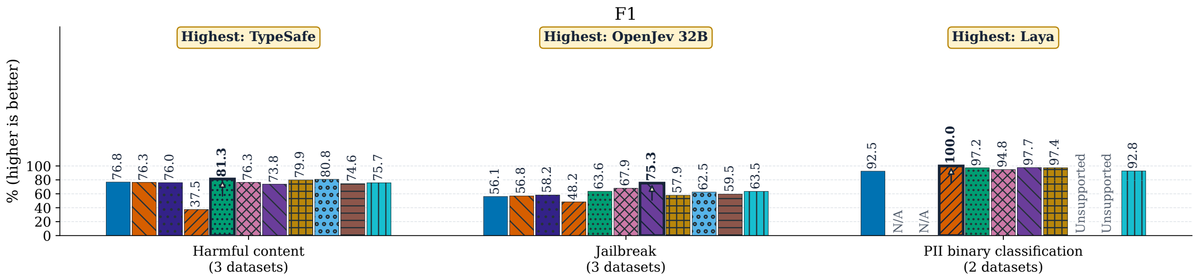
Precision



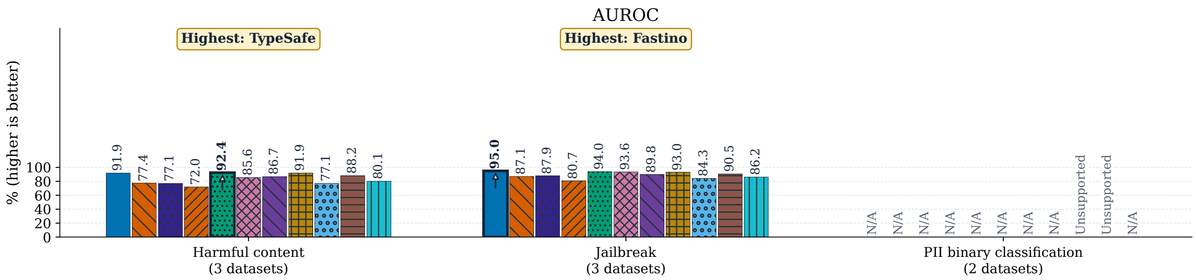
Recall



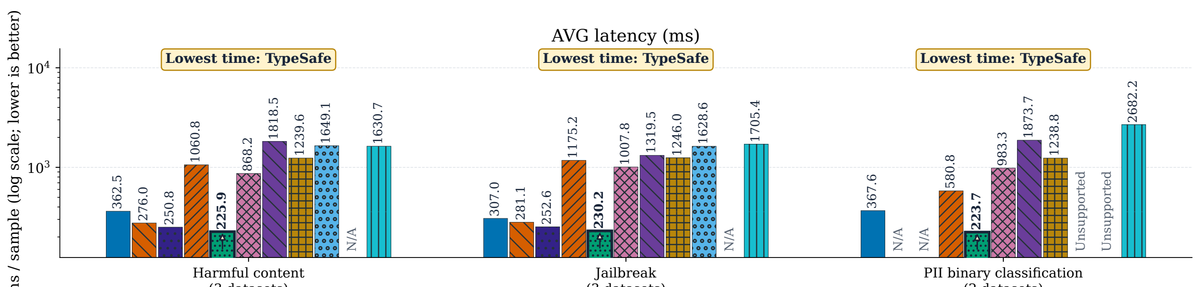
F1



AUROC



AVG latency (ms)



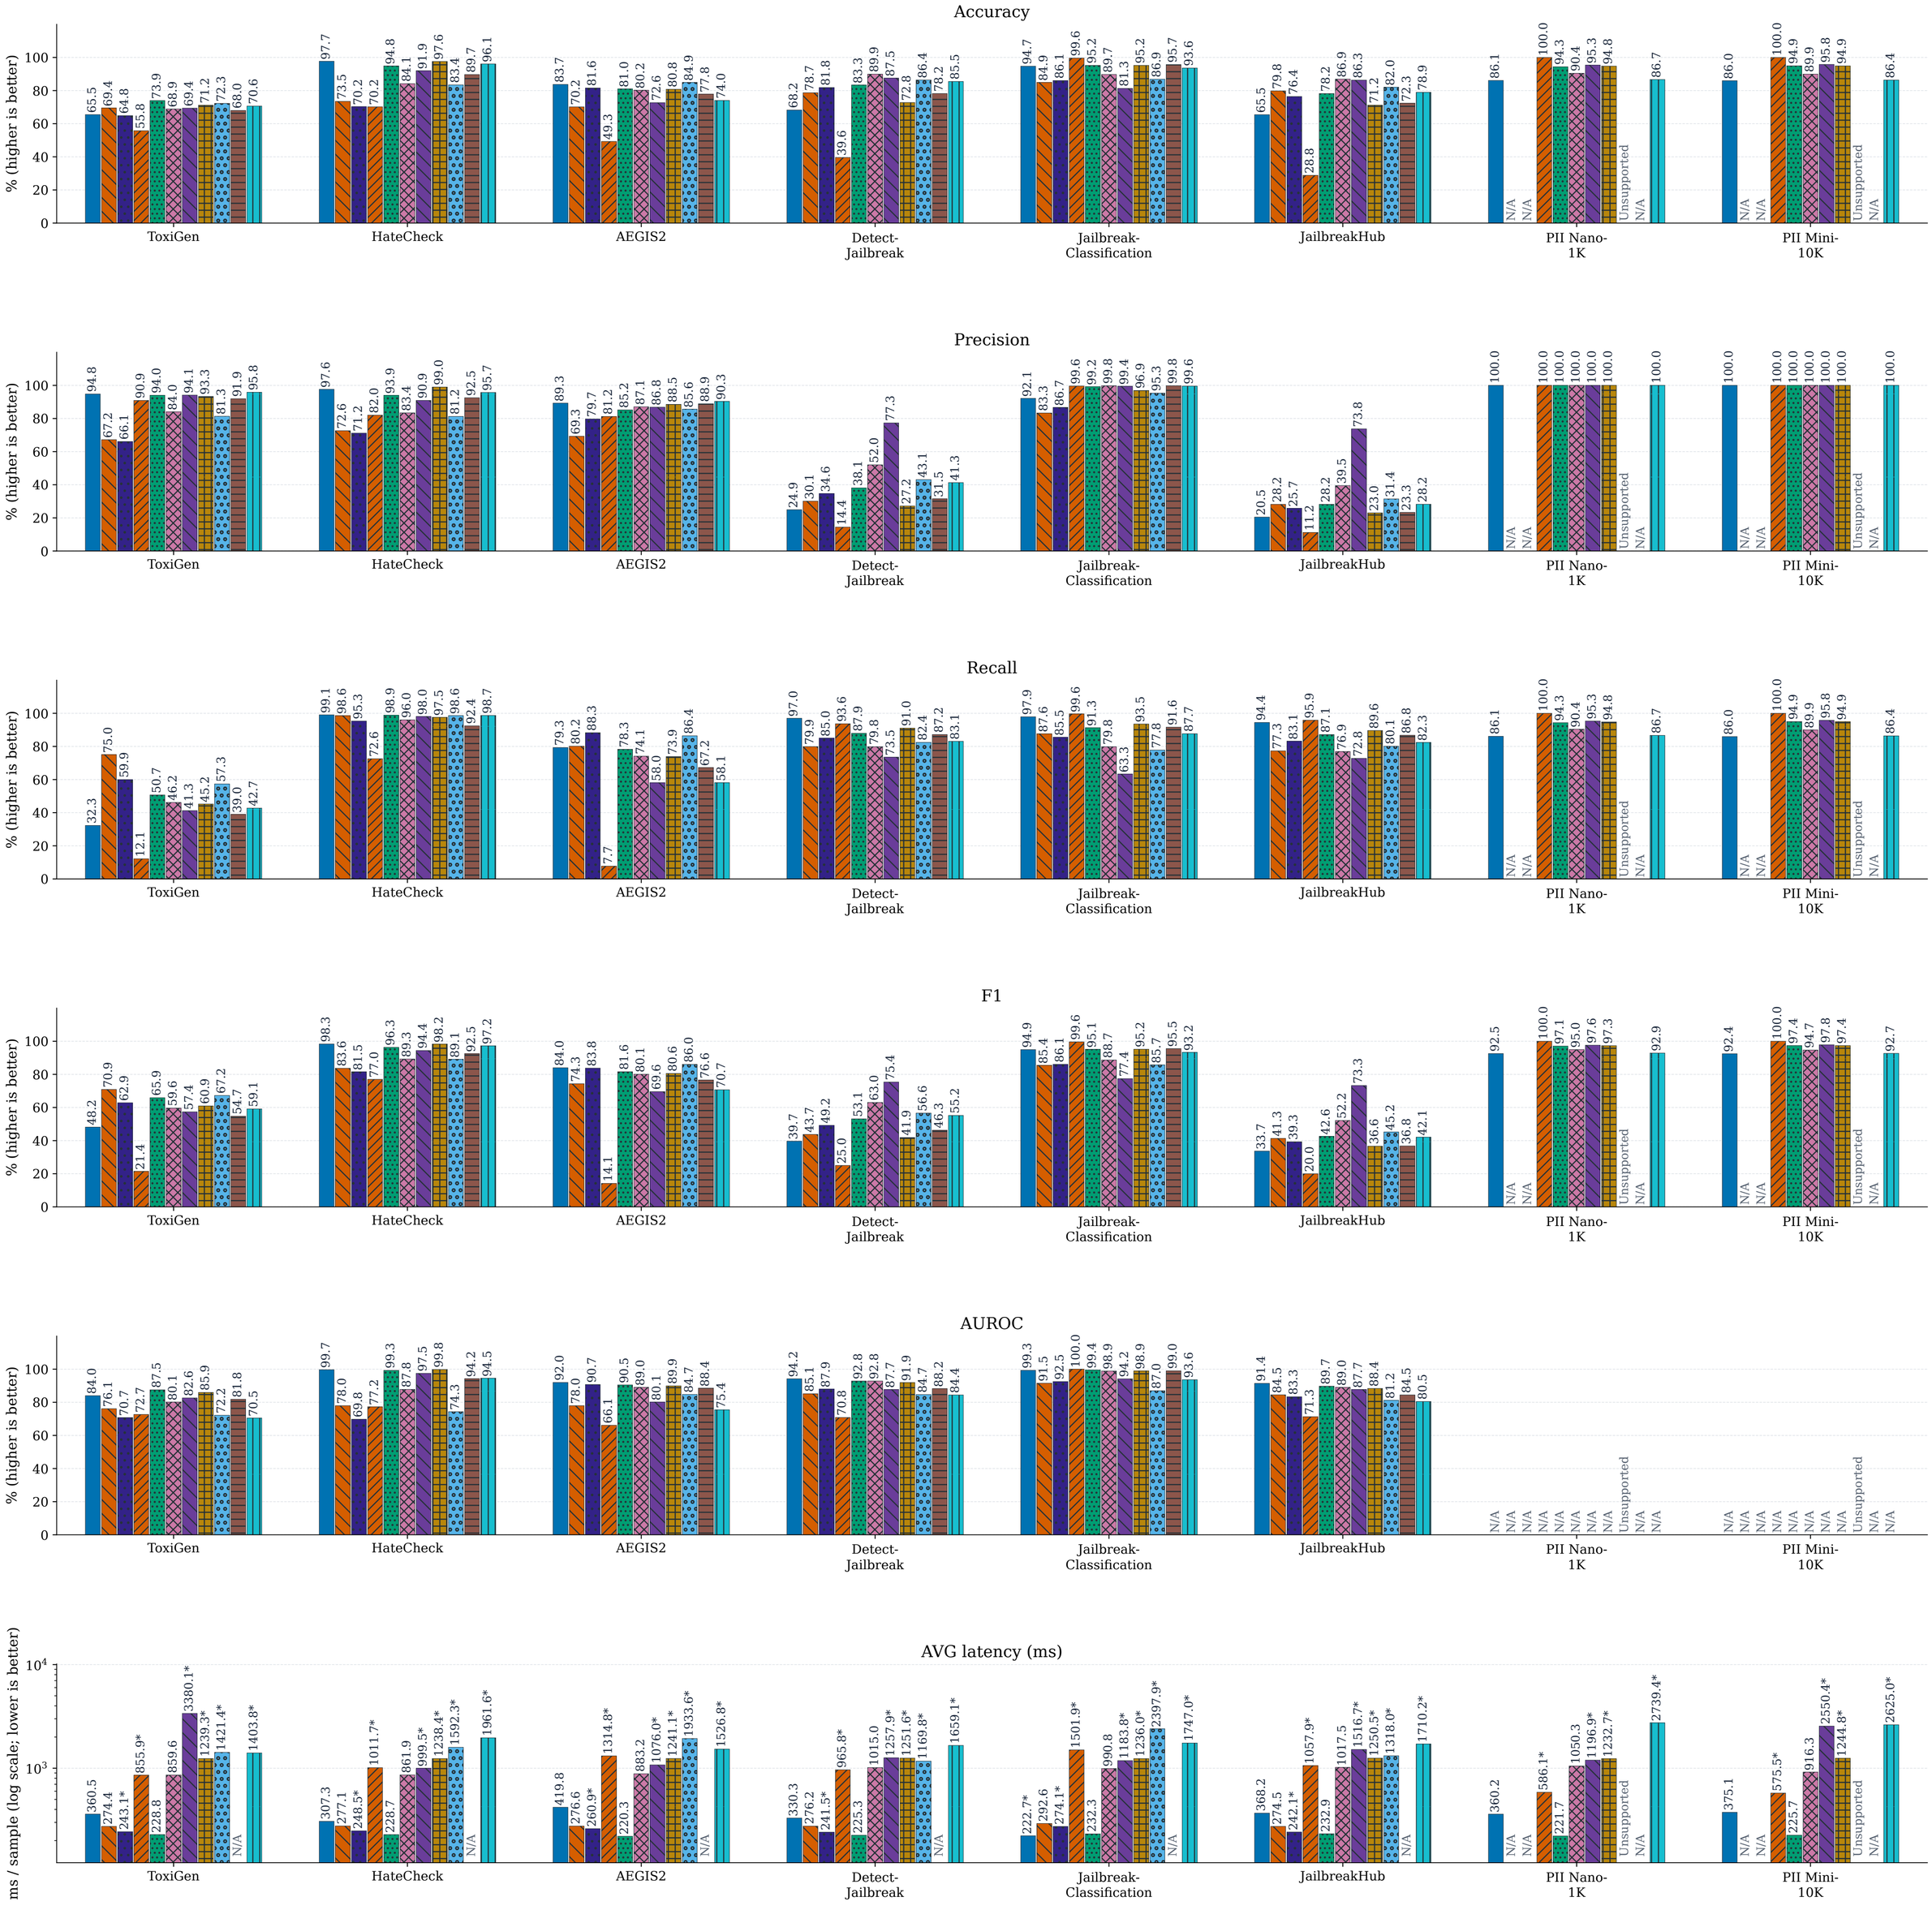
PII = ordinary binary classification (present vs absent), NOT masking quality.

Harmful content: ToxiGen, HateCheck, AEGIS2. Jailbreak: Detect-Jailbreak, Jailbreak-Classification, JailbreakHub.
PII: Nano-1K, Mini-10K; ordinary binary classification (present vs absent), not extraction or masking quality.
Current PII datasets contain positives only; AUROC is undefined.
Each metric is the arithmetic mean of per-dataset metrics, not pooled samples or F1 recomputed from means.
Missing required data or undefined metrics = N/A; category scopes are never reduced to available datasets.
Latency averages the same reference timings as the per-dataset chart; protocols differ and include failure waits.
Own successful samples / supplied aggregates; descriptive only, not a fair ranking.
Qwen3Guard-Stream-8B: binary safety labels; PII unsupported (historical proxies are audit-only).
Qwen mean latency: log elapsed/processed (failures included); CSV P95: successful requests. No ranking.
Luna: original task-specific presets, no custom prompts (prompt N/A); binary threshold 0.5.
Other models use dataset-specific v1/v2 prompts. Task-definition differences may affect results.
Luna three-dataset reference latency: 181.97 ms (user-measured; timing/weights unspecified).
Per-dataset latency and binary PII: N/A; PII micro-F1 is separate. Scorer versions unverified.
Straiker: supplied v1 aggregates; threshold and model/version N/A (not confirmed as 0.5).
Single-run log elapsed/processed includes failures; JailbreakHub: 1 error. PII: binary classification.
Raw files unverified; own valid samples only, not a sample-aligned ranking.

Per-Dataset Multi-Model Reference Comparison (v1 / Threshold 0.5) | Luna: native presets (prompt N/A) | Straiker: threshold N/A

threshold = 0.5 | own successful samples | internal models: aggregates, NOT sample-paired

- Fastino | request: fastino/GLiDE resolved: glide_v2
- Fastino GLiNER2.5-Decide model ID: fastino/gliner2.5-decide response model: fastino/gliner2.5-decide v1 / 0.5; PII not tested
- Fastino GLiGuard-LLMGuardrails-300M model ID: fastino/gliguard-LLMGuardrails-300M v1 / 0.5; PII not tested
- Laya | request: typed-decisions serving: laya-rl-agent routed: typed-decisions
- TypeSafe | request: jev-latest resolved: unconfirmed (API response required)
- OpenJev 4B | supplied internal summary model / version: not supplied (not sample-paired)
- OpenJev 32B | supplied internal summary single-run timing (not sample-paired)
- Microsoft Decision-1 | supplied internal summary single-run timing (not sample-paired)
- Qwen3Guard-Stream-8B | supplied summary binary safety labels; PII unsupported
- Luna | original task-specific presets prompt: N/A; binary threshold: 0.5 PII: separate multilabel evaluation
- Straiker | supplied aggregate prompt: v1; threshold: N/A single-run log elapsed/processed



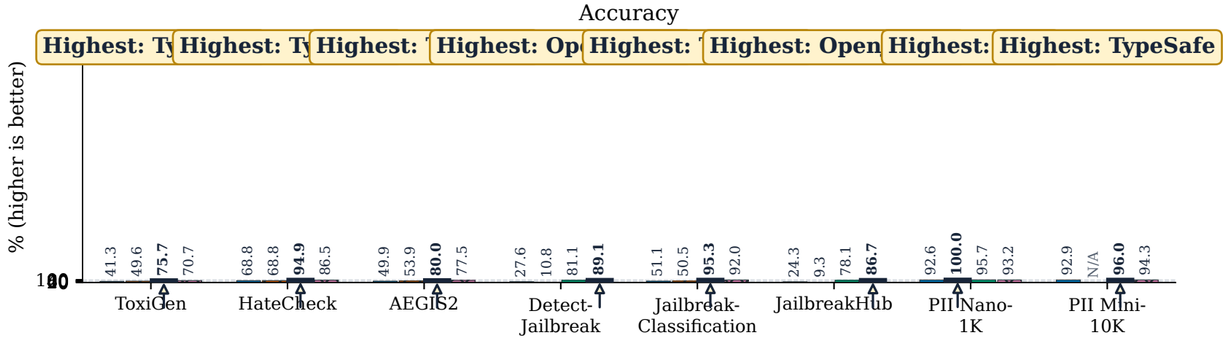
PII = ordinary binary classification (present vs absent), NOT masking quality.

OpenJev 4B: four-run summary timing; local models prefer four-run log averages. 4B valid/error counts are unknown.
* = single-run v1 / 0.5 timing; local models require complete, matched logs; internal models use supplied summaries.
OpenJev 32B / Microsoft Decision-1: supplied single-run timing, including failures; protocols differ.
32B has 27,460 errors / 56,823 rows; quality excludes errors.
N/A = missing quality, undefined metric, or incomplete timing; missing values are not zero.
Descriptive only, not a fair ranking. PII is ordinary binary classification (present vs absent), not masking quality.
Qwen3Guard-Stream-8B: binary safety labels; PII unsupported (historical proxies are audit-only).
Qwen mean latency: log elapsed/processed (failures included); CSV P95: successful requests. No ranking.
Luna: original task-specific presets, no custom prompts (prompt N/A); binary threshold 0.5.
Other models use dataset-specific v1/v2 prompts. Task-definition differences may affect results.
Luna three-dataset reference latency: 181.97 ms (user-measured; timing/weights unspecified).
Per-dataset latency and binary PII: N/A; PII micro-F1 is separate. Scorer versions unverified.
Straiker: supplied v1 aggregates; threshold and model/version N/A (not confirmed as 0.5).
Single-run log elapsed/processed includes failures; JailbreakHub: 1 error. PII: binary classification.
Raw files unverified; own valid samples only, not a sample-aligned ranking.

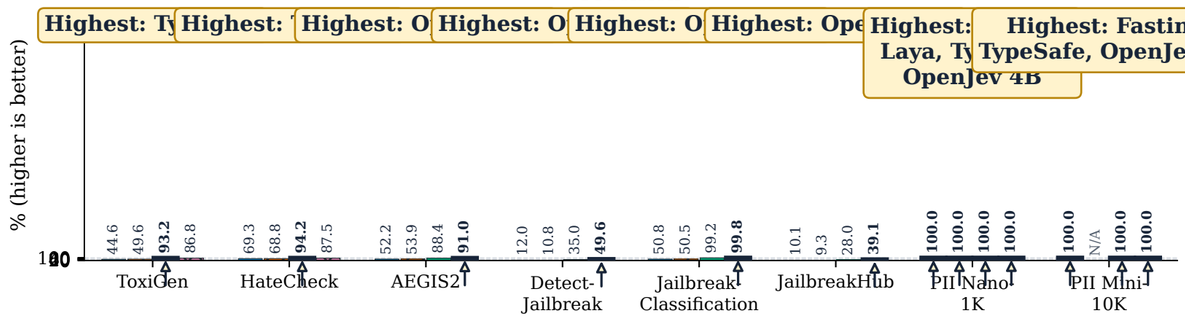
Per-Dataset Reference Comparison: v2 / Threshold 0.5 (Fastino, Laya, TypeSafe, OpenJev 4B)

threshold = 0.5 | own successful samples | internal models: aggregates, NOT sample-paired

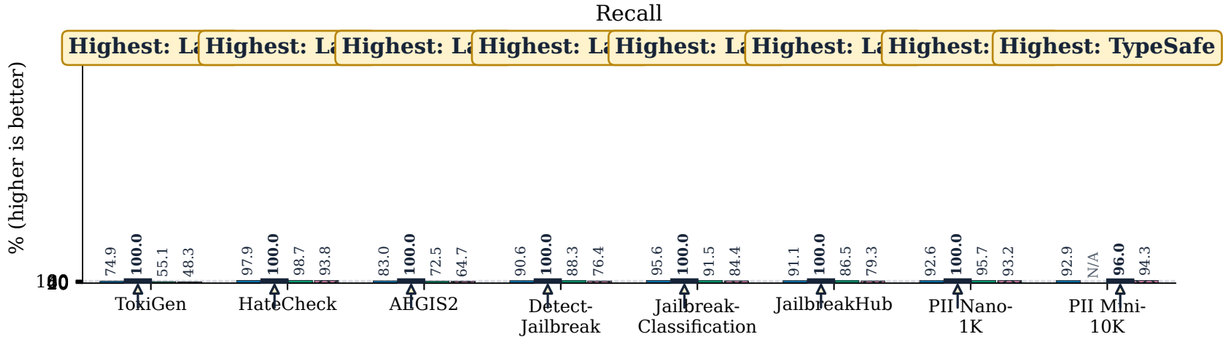
Accuracy



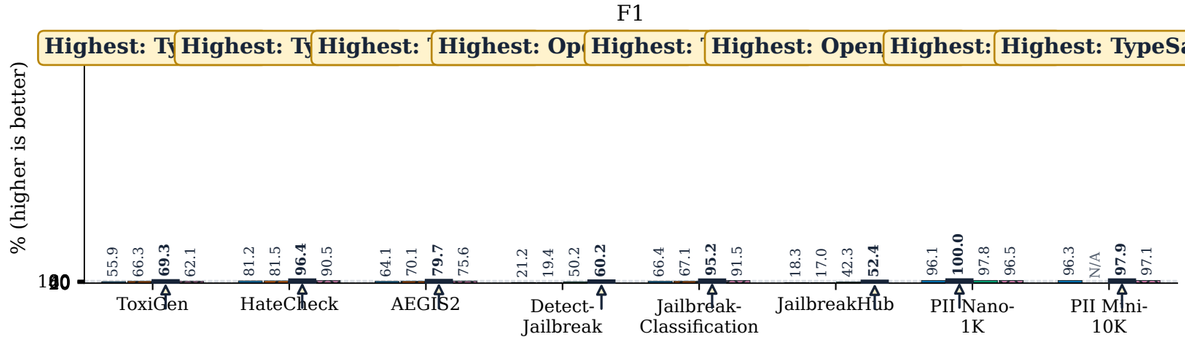
Precision



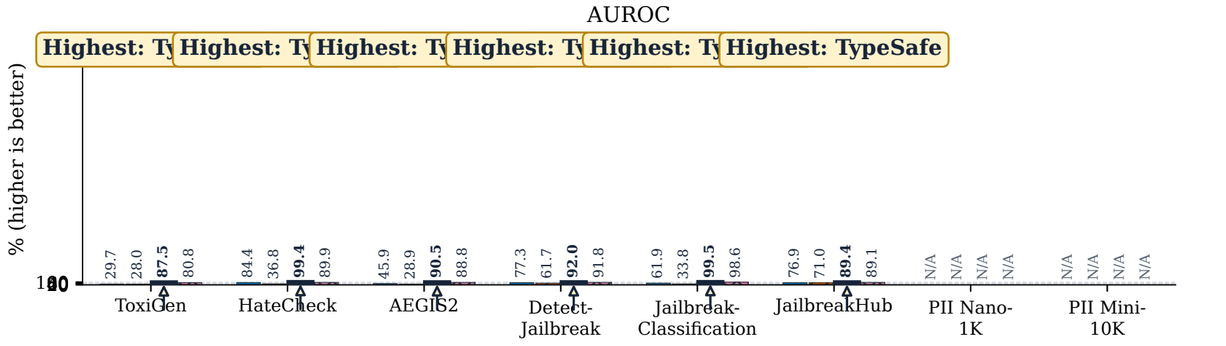
Recall



F1

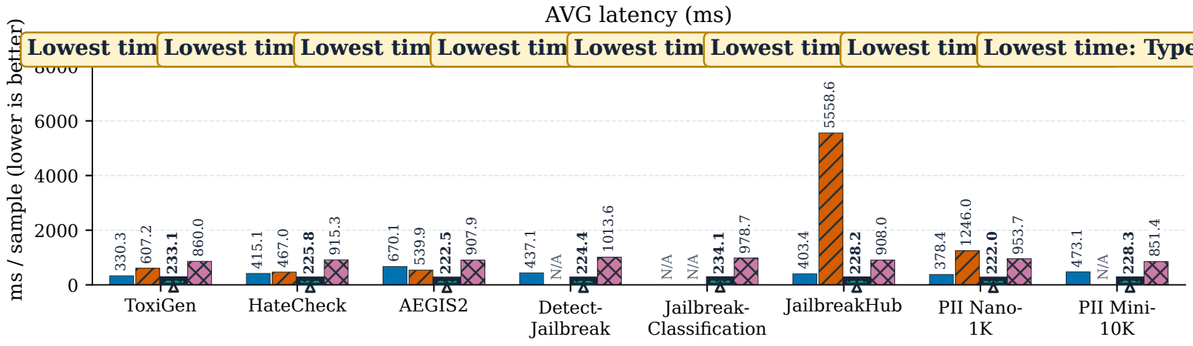


AUROC



PII = ordinary binary classification (present vs absent), NOT masking quality.

AVG latency (ms)



OpenJev 4B: four-run summary timing; local models prefer four-run log averages. 4B valid/error counts are unknown.
N/A = missing quality, undefined metric, or incomplete timing; missing values are not zero.
Descriptive only, not a fair ranking. PII is ordinary binary classification (present vs absent), not masking quality.

Fastino | request: fastino/GLiDE
resolved: glide_v2

Laya | request: typed-decisions
serving: laya-rl-agent
routed: typed-decisions

TypeSafe | request: jev-latest
resolved: unconfirmed (API response required)

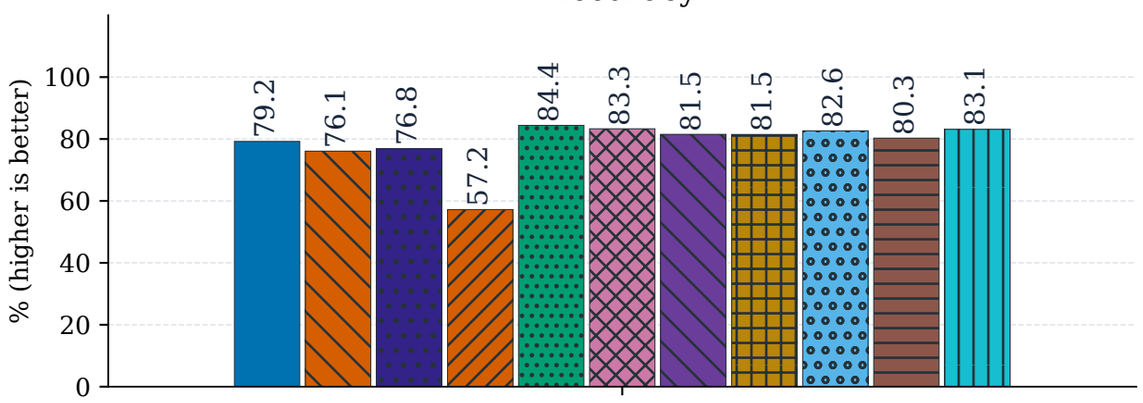
OpenJev 4B | supplied internal summary
model / version: not supplied (not sample-paired)

Multi-Model Reference Averages (v1 / Threshold 0.5) | Luna: native presets (prompt N/A) | Straiker: threshold N/A

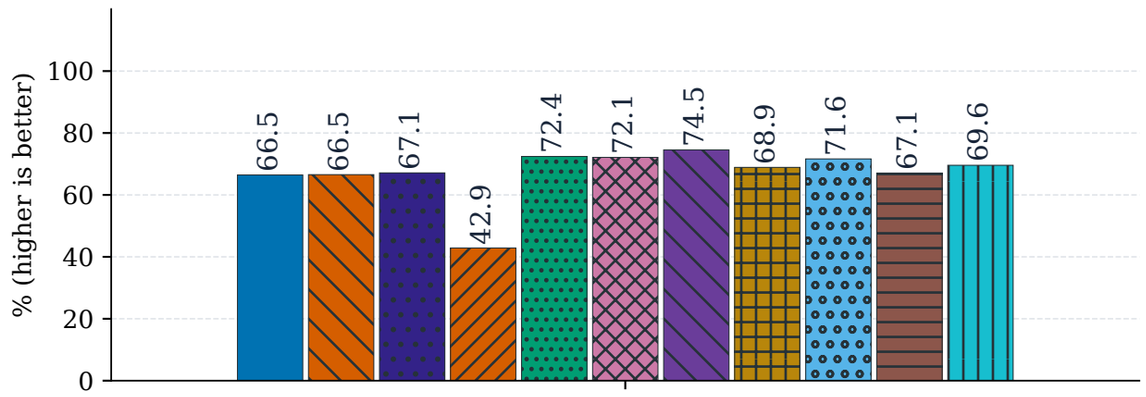
v1 | threshold = 0.5 | shared non-PII dataset names | equal dataset weights | PII reported separately

- Fastino | request: fastino/GLiDE
  resolved: glide_v2
- Fastino GLiNER2.5-Decide
  model ID: fastino/gliner2.5-decide
  response model: fastino/gliner2.5-decide
  v1 / 0.5; PII not tested
- Fastino GLiGuard-LLMGuardrails-300M
  model ID: fastino/gliguard-LLMGuardrails-300M
  v1 / 0.5; PII not tested
- Laya | request: typed-decisions
  serving: laya-rl-agent
  routed: typed-decisions
- TypeSafe | request: jev-latest
  resolved: unconfirmed (API response required)
- OpenJev 4B | supplied internal summary
  model / version: not supplied (not sample-paired)
- OpenJev 32B | supplied internal summary
  single-run timing (not sample-paired)
- Microsoft Decision-1 | supplied internal summary
  single-run timing (not sample-paired)
- Qwen3Guard-Stream-8B | supplied summary
  binary safety labels; PII unsupported
- Luna | original task-specific presets
  prompt: N/A; binary threshold: 0.5
  PII: separate multilabel evaluation
- Straiker | supplied aggregate
  prompt: v1; threshold: N/A
  single-run log elapsed/processed

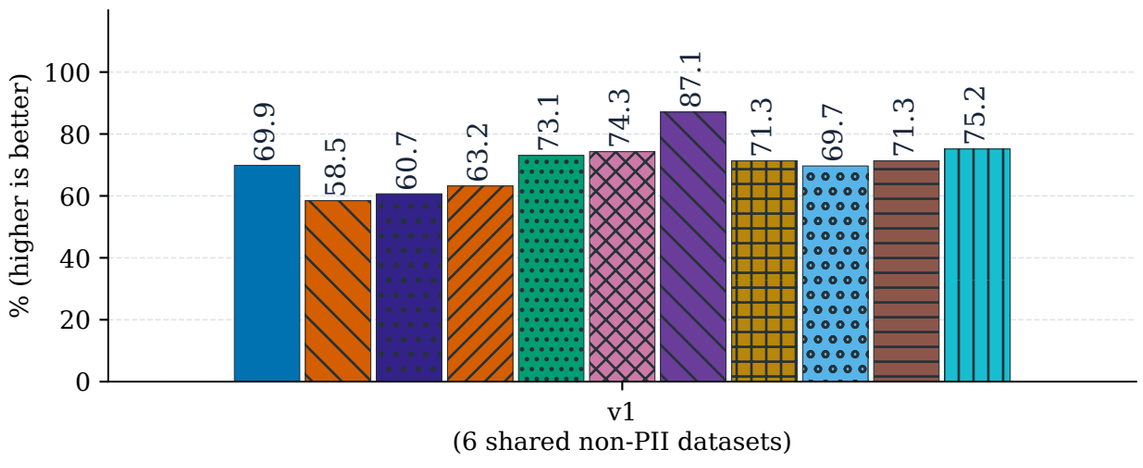
Accuracy



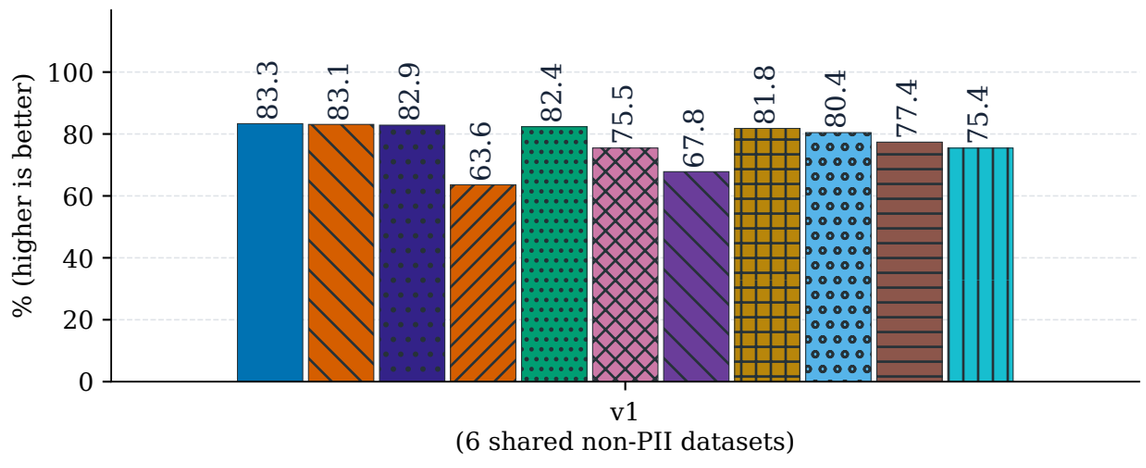
F1



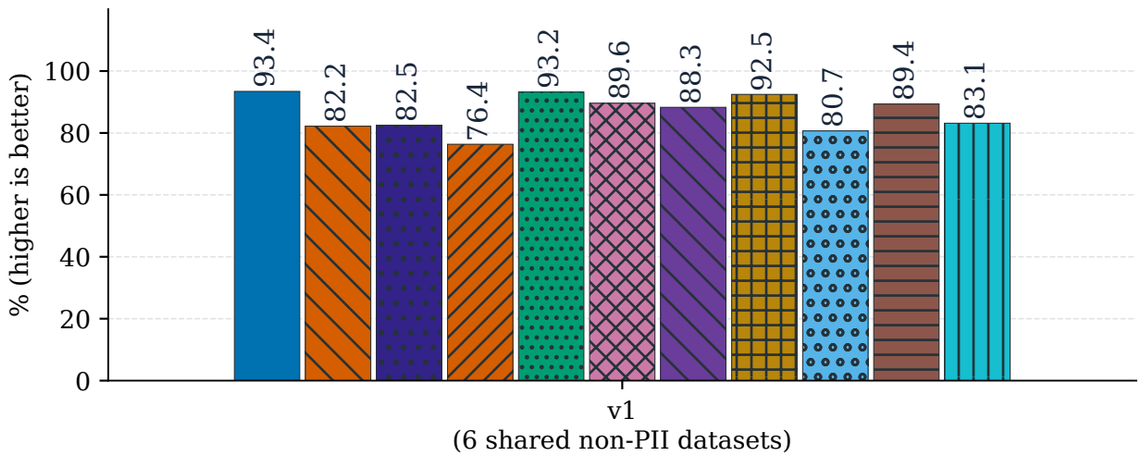
Precision



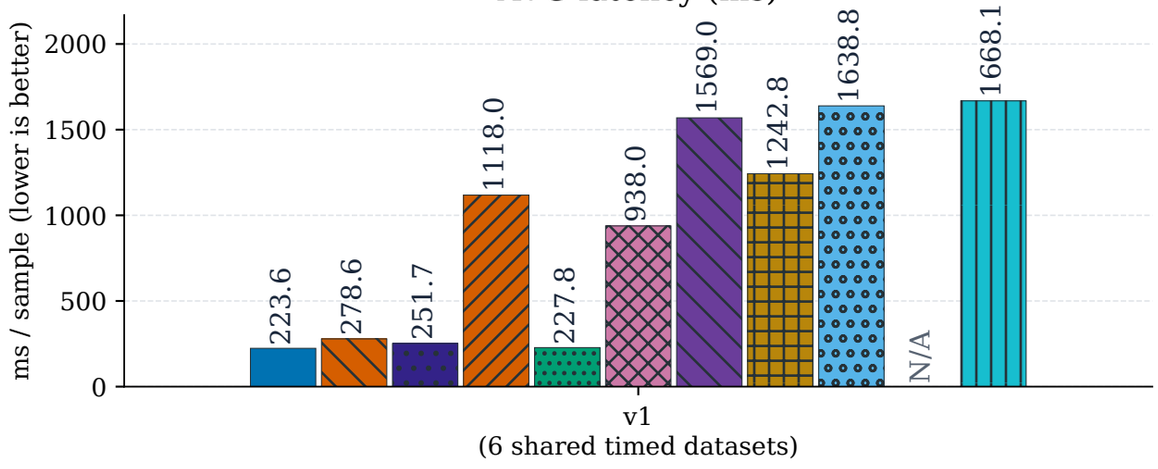
Recall



AUROC



AVG latency (ms)

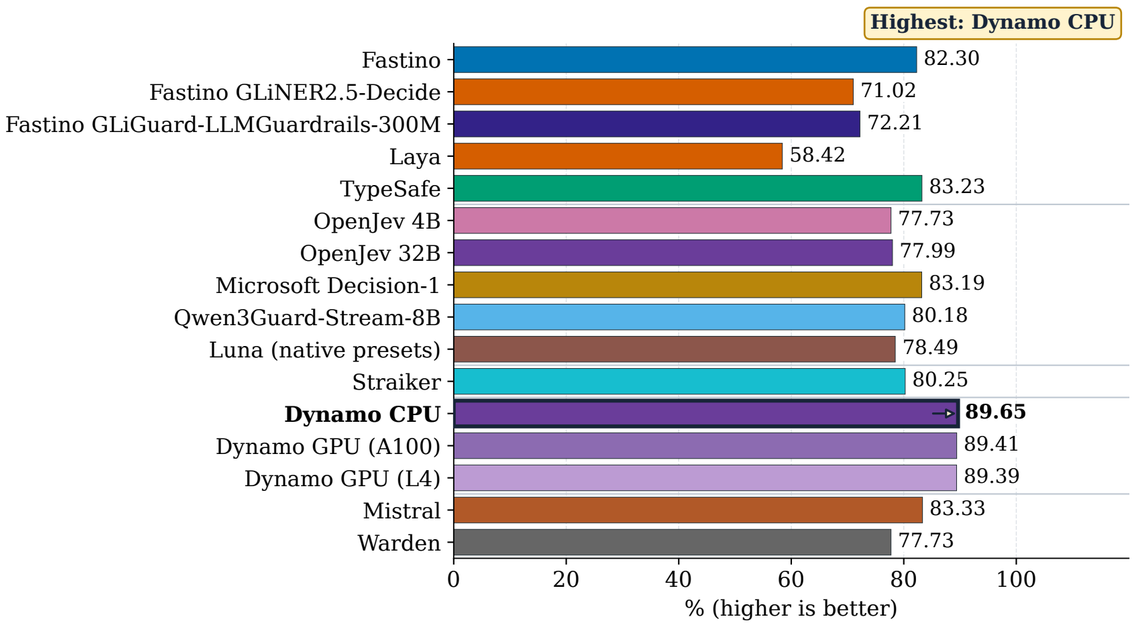


Each model uses its own successful samples; aggregate models have no paired sample IDs. Not a fair ranking.
Latency: original models' shared fully timed v1 / 0.5 subset; Qwen follows the same dataset scope:
ToxiGen, HateCheck, AEGIS2, Detect-Jailbreak, Jailbreak-Classification, JailbreakHub. Quality and latency dataset scopes can differ.
APIs: v1 / 0.5 single-run timing. OpenJev 4B: supplied summary; 32B / Decision-1: supplied single-run timing.
PII is ordinary binary classification (present vs absent), not masking quality; positives-only results are separate.
Missing quality stays N/A; no shared timing subset = N/A latency. Timing protocols differ; not inference time.
Qwen3Guard-Stream-8B: binary safety labels; PII unsupported (historical proxies are audit-only).
Qwen mean latency: log elapsed/processed (failures included); CSV P95: successful requests. No ranking.
Luna: original task-specific presets, no custom prompts (prompt N/A); binary threshold 0.5.
Other models use dataset-specific v1/v2 prompts. Task-definition differences may affect results.
Luna three-dataset reference latency: 181.97 ms (user-measured; timing/weights unspecified).
Per-dataset latency and binary PII: N/A; PII micro-F1 is separate. Scorer versions unverified.
Straiker: supplied v1 aggregates; threshold and model/version N/A (not confirmed as 0.5).
Single-run log elapsed/processed includes failures; JailbreakHub: 1 error. PII: binary classification.
Raw files unverified; own valid samples only, not a sample-aligned ranking.

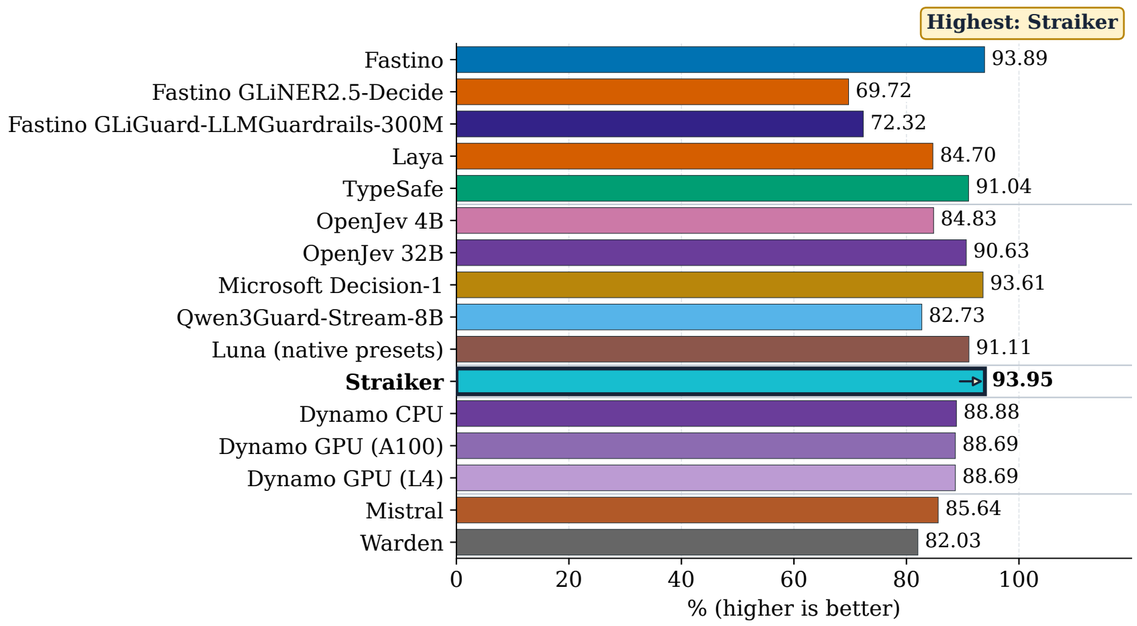
ToxiGen, HateCheck, AEGIS2: Multi-model Reference Averages | Luna: native presets (prompt N/A) | Straiker: threshold N/A

Computed averages: ToxiGen / HateCheck / AEGIS2 | supplied-model dataset scopes unverified

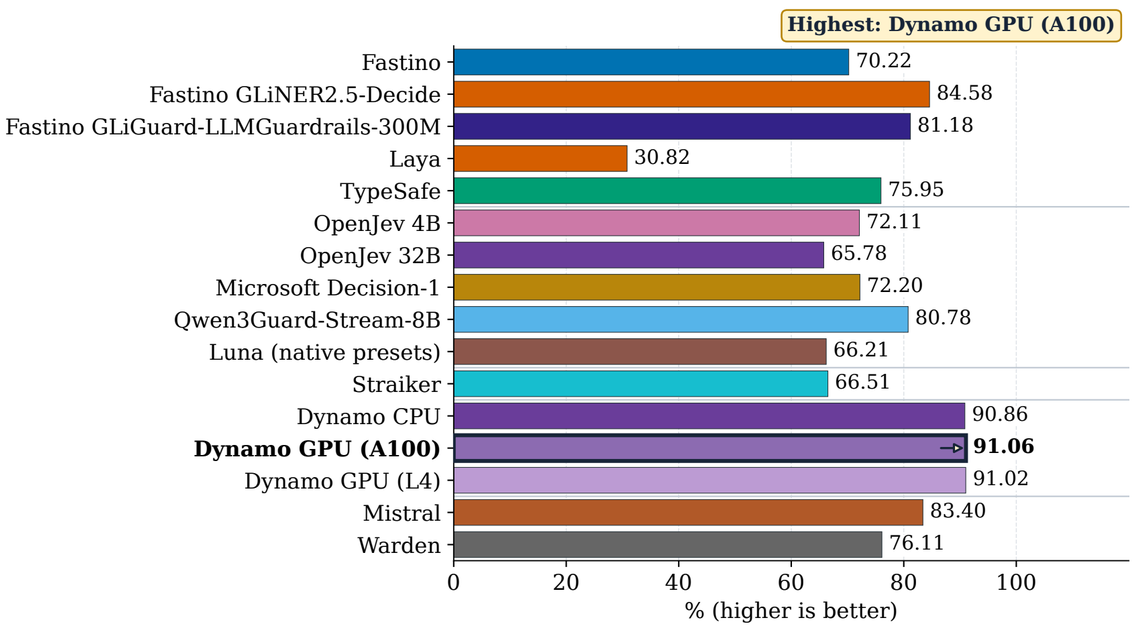
Accuracy



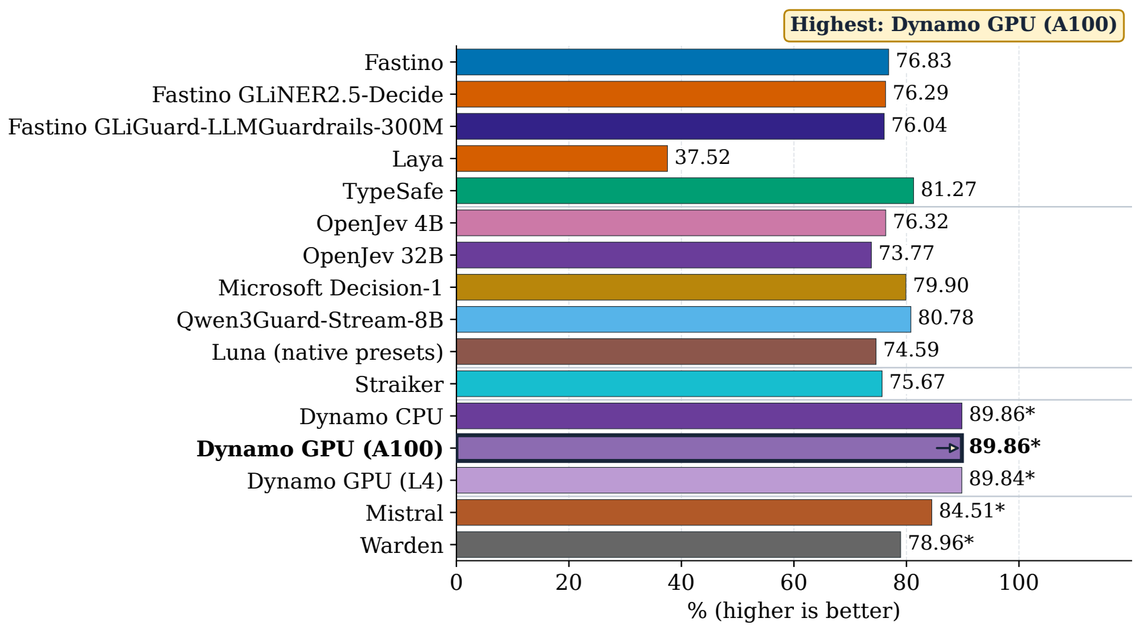
Precision



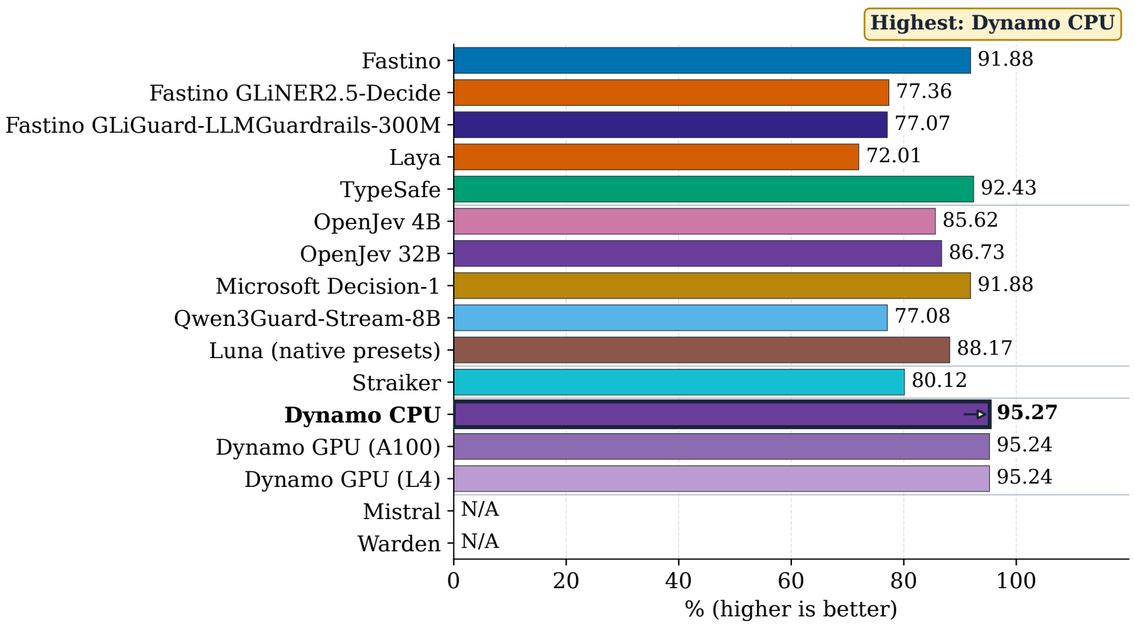
Recall



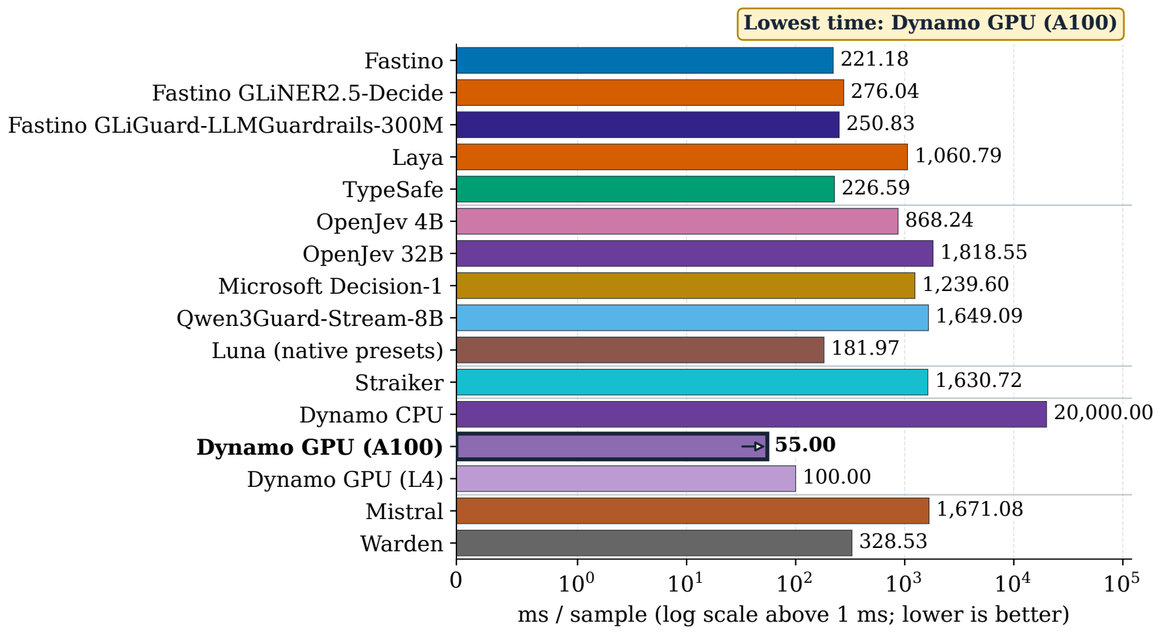
F1



AUROC



AVG latency (ms)



Computed models: no internal dataset was available; v1 / 0.5 means use only the three shared datasets.
Local latency uses v1 / 0.5 runs; OpenJev 4B / 32B and Decision-1 use their supplied summary latency.
Dynamo / Mistral / Warden: supplied averages; their evaluations may also include an internal dataset.
Its inclusion in their averages is unconfirmed; prompt, threshold and timing protocols were not supplied.
* F1 = 2PR / (P + R) from supplied averages; not a mean of per-dataset F1. N/A = unavailable.
Successful samples and timing protocols are not aligned; descriptive reference only, not a fair ranking.
Qwen3Guard-Stream-8B: binary safety labels; PII unsupported (historical proxies are audit-only).
Qwen mean latency: log elapsed/processed (failures included); CSV P95: successful requests. No ranking.
Luna: original task-specific presets, no custom prompts (prompt N/A); binary threshold 0.5.
Other models use dataset-specific v1/v2 prompts. Task-definition differences may affect results.
Luna three-dataset reference latency: 181.97 ms (user-measured; timing/weights unspecified).
Per-dataset latency and binary PII: N/A; PII micro-F1 is separate. Scorer versions unverified.
Straiker: supplied v1 aggregates; threshold and model/version N/A (not confirmed as 0.5).
Single-run log elapsed/processed includes failures; JailbreakHub: 1 error. PII: binary classification.
Raw files unverified; own valid samples only, not a sample-aligned ranking.