Baseline quality, errors and elapsed time

v1 | threshold = 0.5 | 4 shared datasets | dataset-macro quality

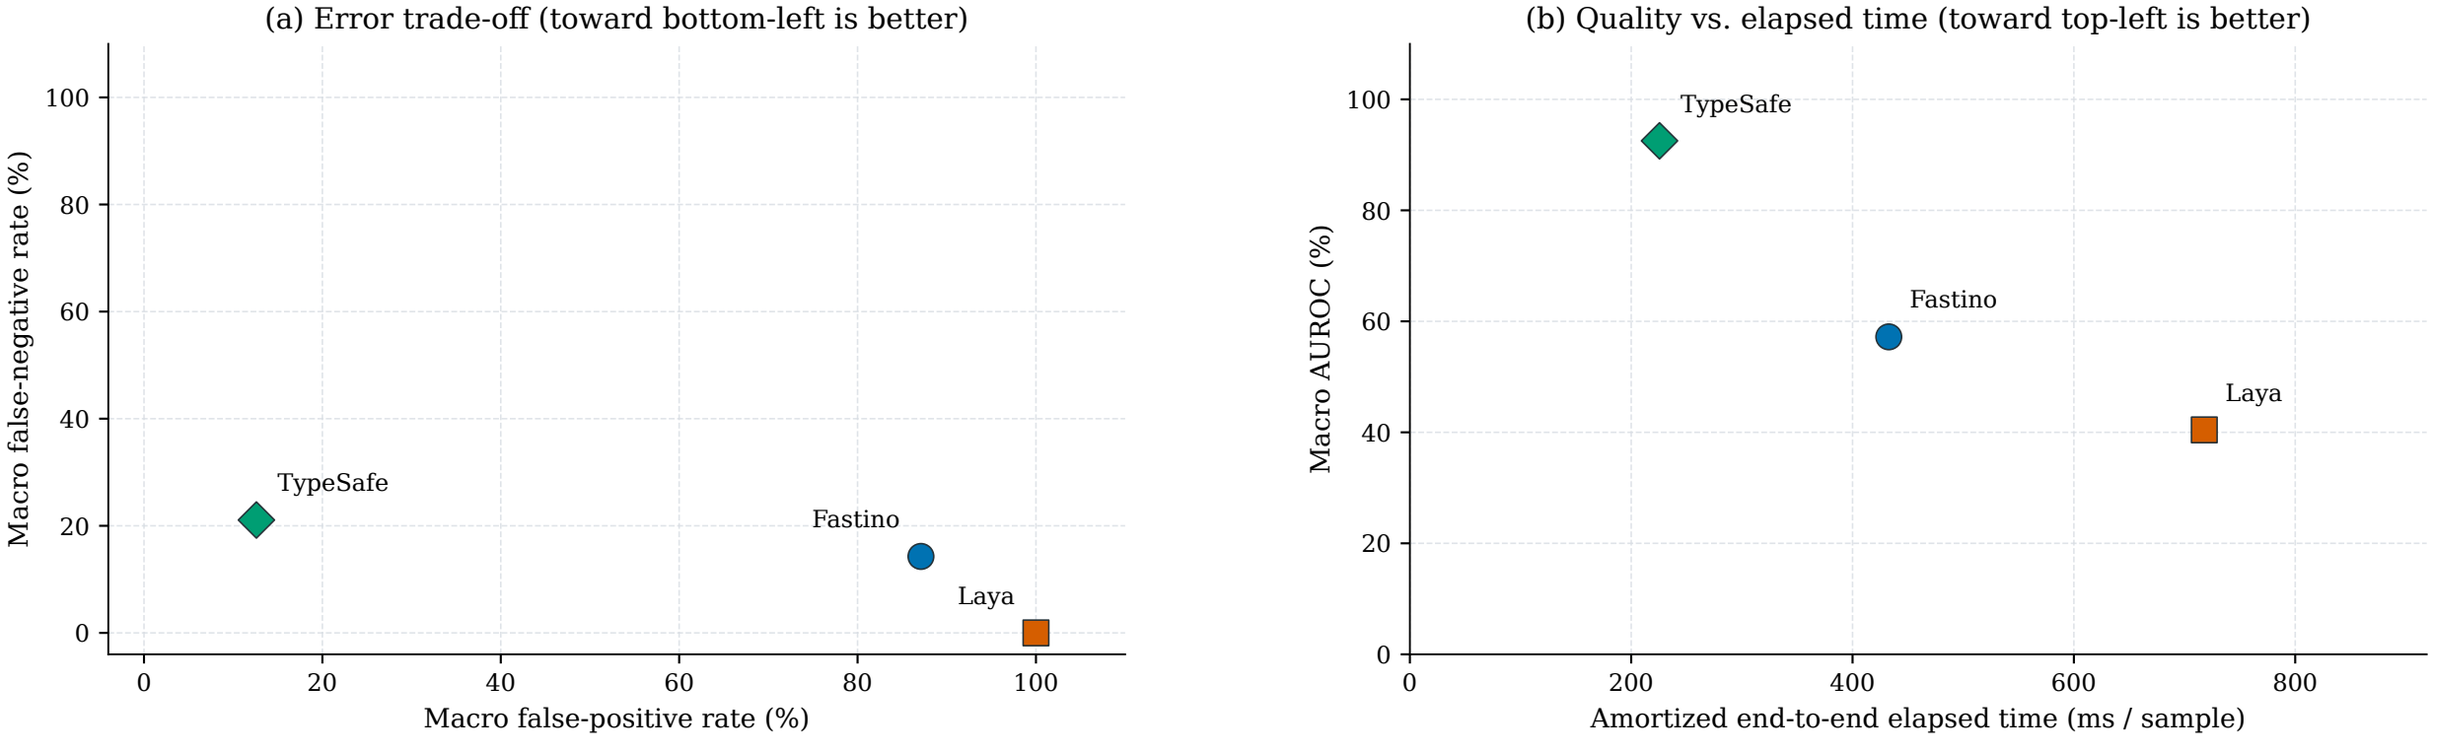

Elapsed time = total logged elapsed / all processed baseline rows, including failed requests.
Runs occurred at different times; this is not pure inference latency or a request-level percentile.

Fastino | request: fastino/GLiDE
resolved: glide_v2

Laya | request: typed-decisions
serving: laya-rl-agent
routed: typed-decisions

TypeSafe | request: jev-latest
resolved: unconfirmed (API response required)