Three-dataset model reference averages

Computed averages: ToxiGen / HateCheck / AEGIS2 | supplied-model dataset scopes unverified

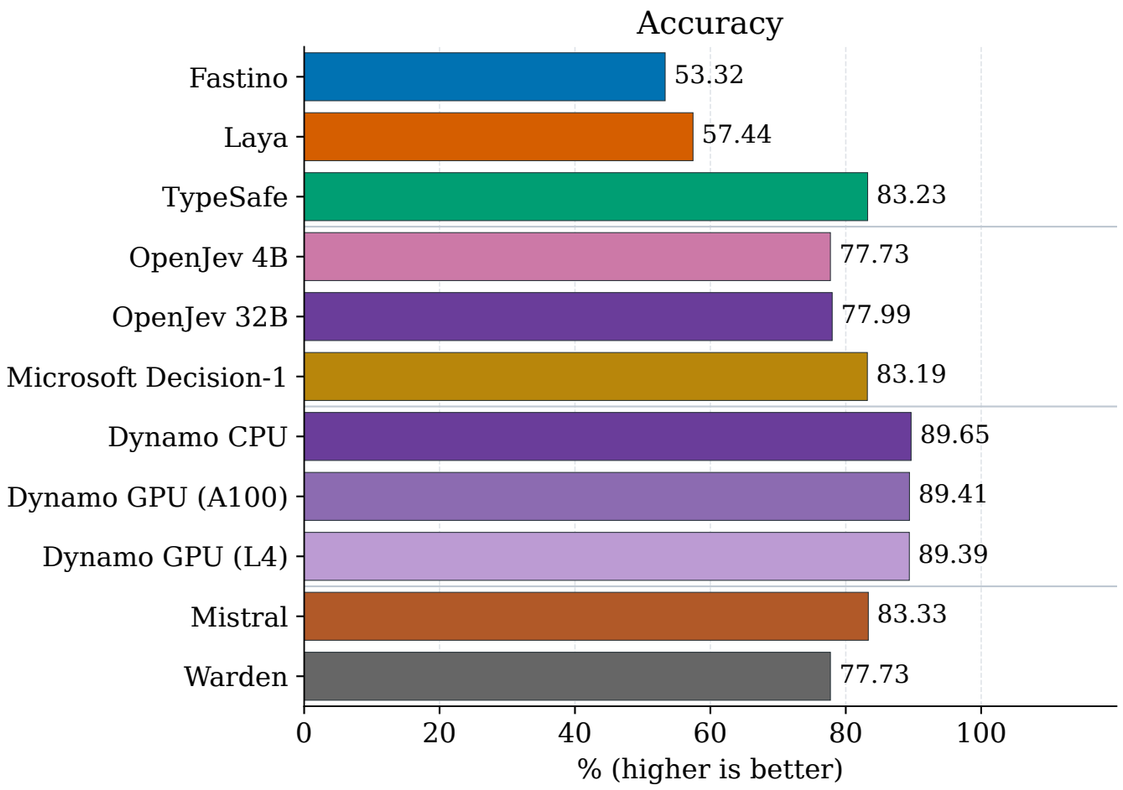

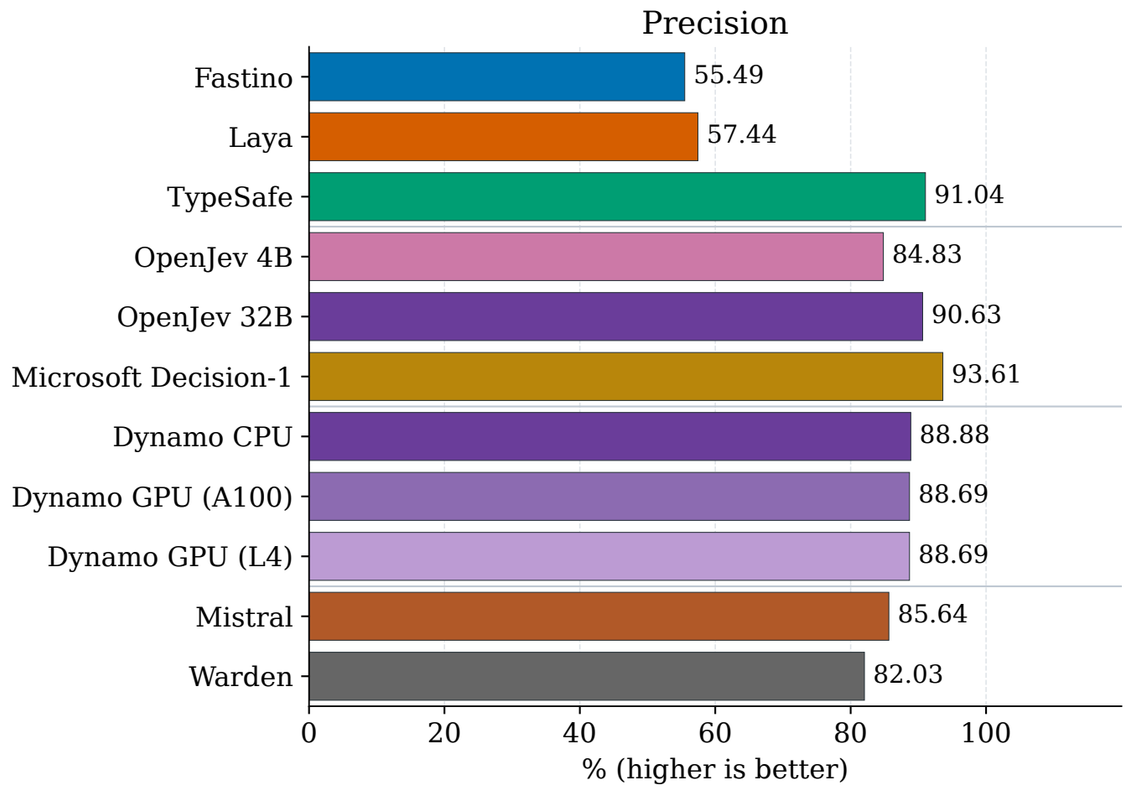

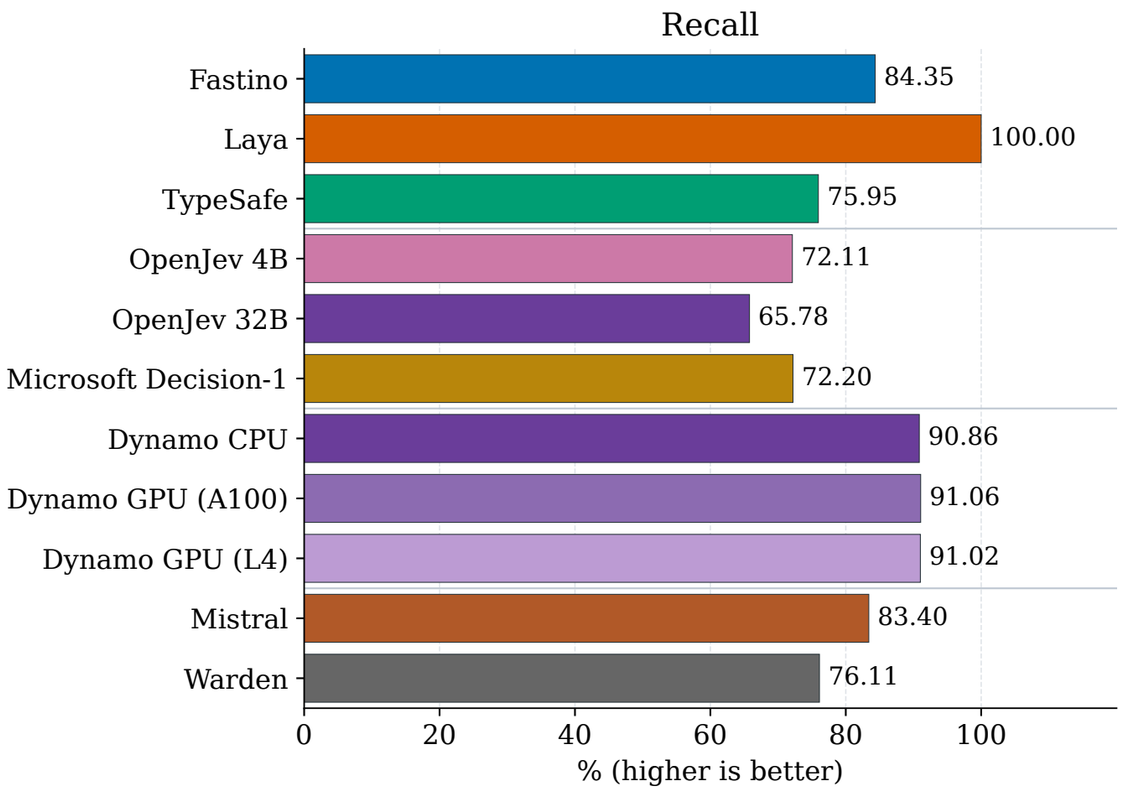

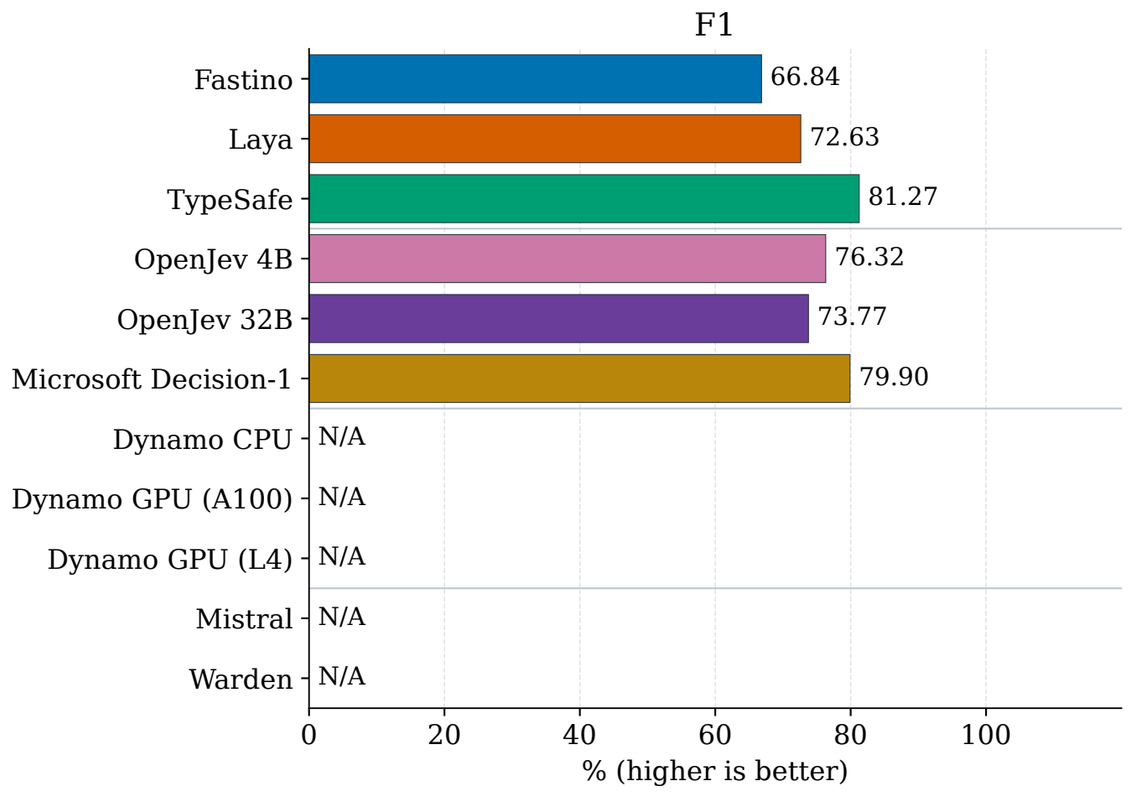

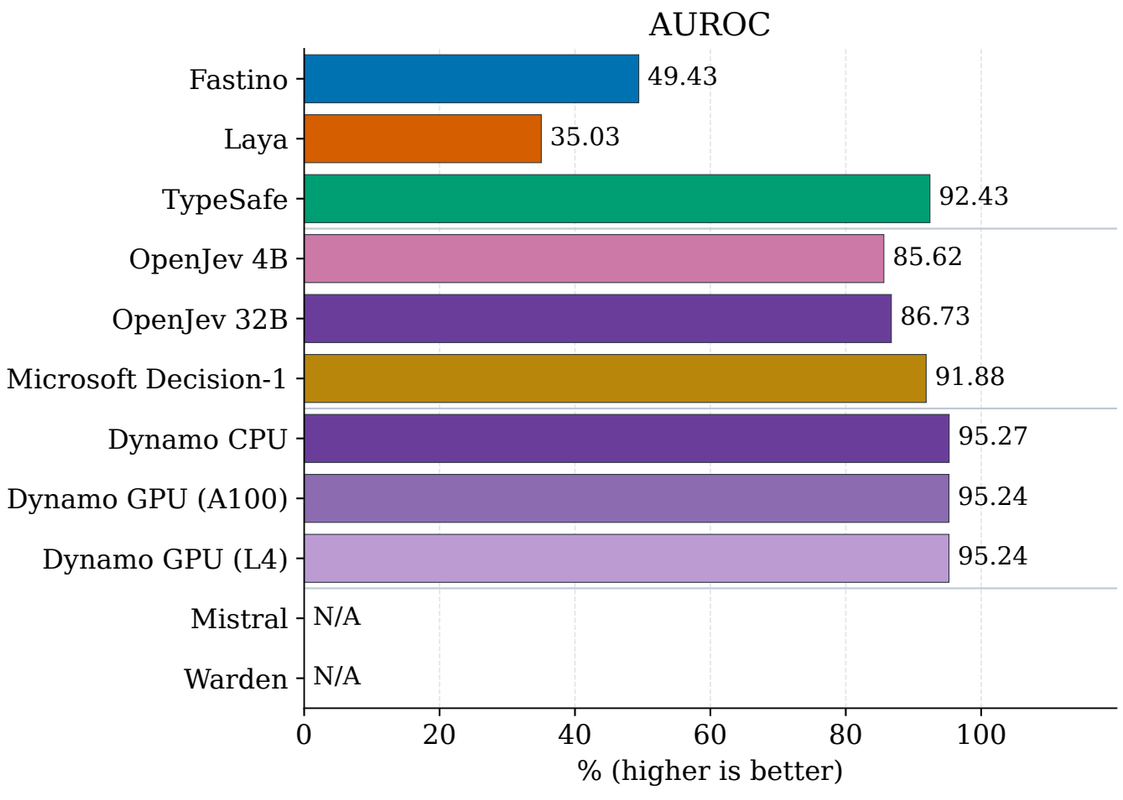

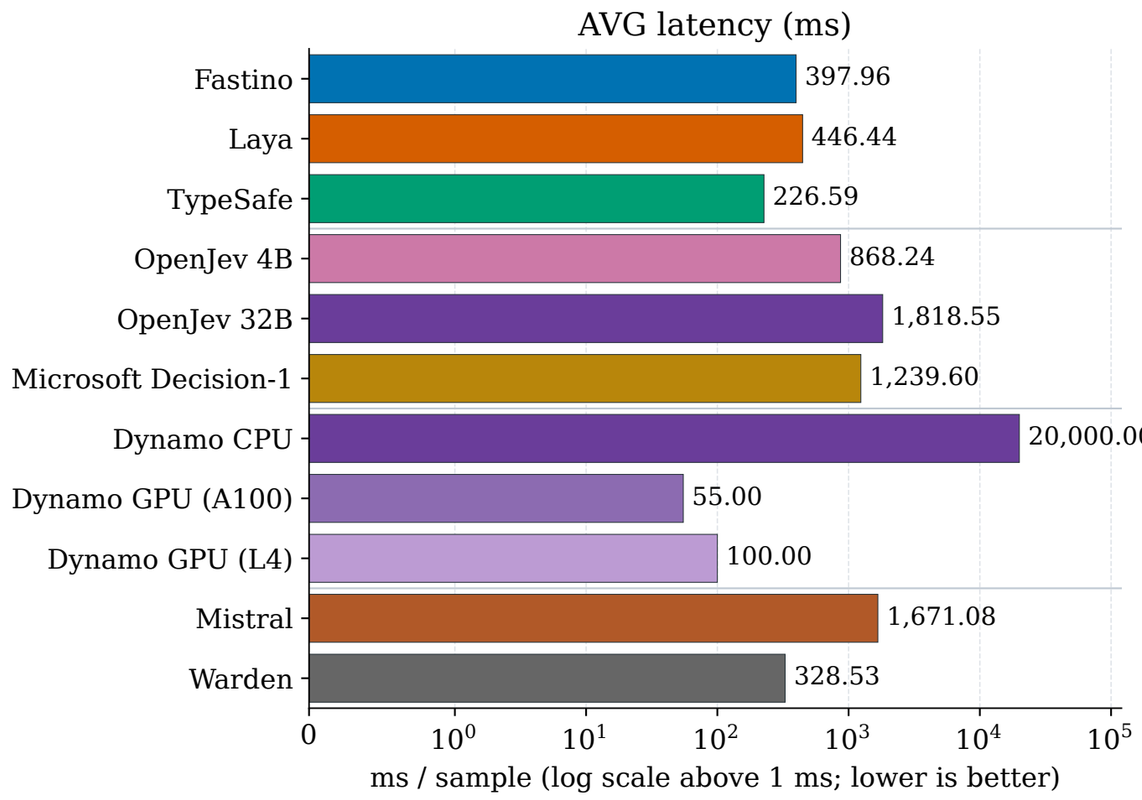

Computed models: no internal dataset was available; v1 / 0.5 means use only the three shared datasets.
Local latency uses v1 / 0.5 runs; OpenJev 4B / 32B and Decision-1 use their supplied summary latency.
Dynamo / Mistral / Warden: supplied averages; their evaluations may also include an internal dataset.
Its inclusion in their averages is unconfirmed; prompt, threshold and timing protocols were not supplied.
N/A = not supplied or required source data missing. F1 is not inferred from average precision / recall.
Successful samples and timing protocols are not aligned; descriptive reference only, not a fair ranking.